


Lecture 3: Processes

Hossein Asadi (asadi@sharif.edu)

Rasool Jalili (jalili@sharif.edu)

Fall 2024





Lecture 3: Processes

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- Examples of IPC Systems
- Communication in Client-Server Systems





Objectives

- To Introduce **Notion** of a **Process**
- To Describe Various **Features** of Processes
 - **Scheduling, Creation / Termination, & Communication**
- To Explore Interprocess Communication
 - Shared Memory
 - Message Passing





Reminder

■ Early Computers

- Only one program in execution at a time
- Program had complete control of system

■ Current Computers Allow

- Multiple programs to be loaded into memory
- Multiple programs executed concurrently

■ ➔ Notion of **Process**





Process Concept

■ **Process** – a Program in Execution

- Process execution must progress in sequential fashion

■ A System Consists of a Set of Processes

- OS processes executing **system code**
- User processes executing **user code**

■ Textbook Uses terms **Job** and **Process** almost Interchangeably

- Batch system: **jobs**
- Time-shared systems: **user programs** or **tasks**





Process Concept (Cont.)

■ Multiple Parts

- Program code, also called **text section**
- Current activity including **program counter**, processor **registers**
- **Stack** containing temporary data
 - ▶ Function parameters, return addresses, local variables
- **Data section** containing global variables
- **Heap** containing memory dynamically allocated during run time






Process Concept (Cont.)

- Program is **Passive** Entity Stored on Disk (**executable file**), process is **active**
 - Program becomes process when executable file loaded into memory
- Execution of Program Starts via:
 - GUI mouse clicks or
 - Command line entry of its name
- One Program can be Several Processes
 - Consider multiple users executing same program





Process in Memory





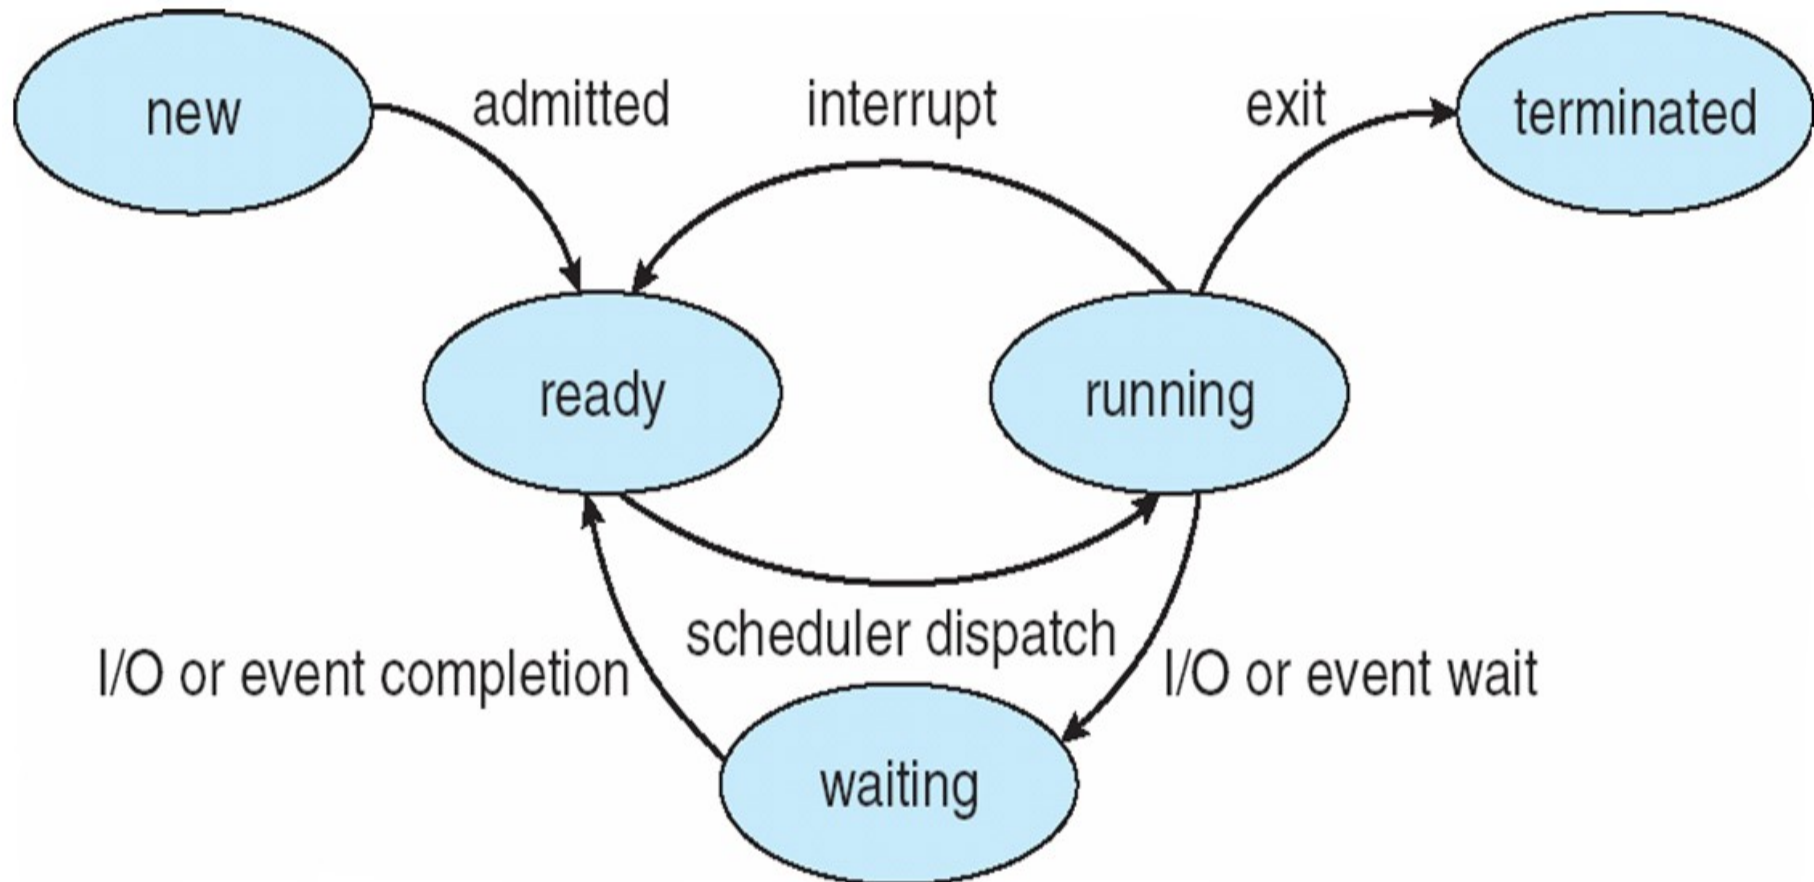
Process State

- As a Process Executes, it Changes **State**
 - **New**: Process is being created
 - **Running**: Instructions are being executed
 - **Waiting**: Process is waiting for some event to occur
 - **Ready**: Process is waiting to be assigned to a processor
 - **Terminated**: Process has finished execution





Diagram of Process State





Process Control Block (PCB)

■ Info Associate with each Process

- Aka, task control block

■ Process State

- Running, waiting, ready, new, & termin.

■ Program Counter

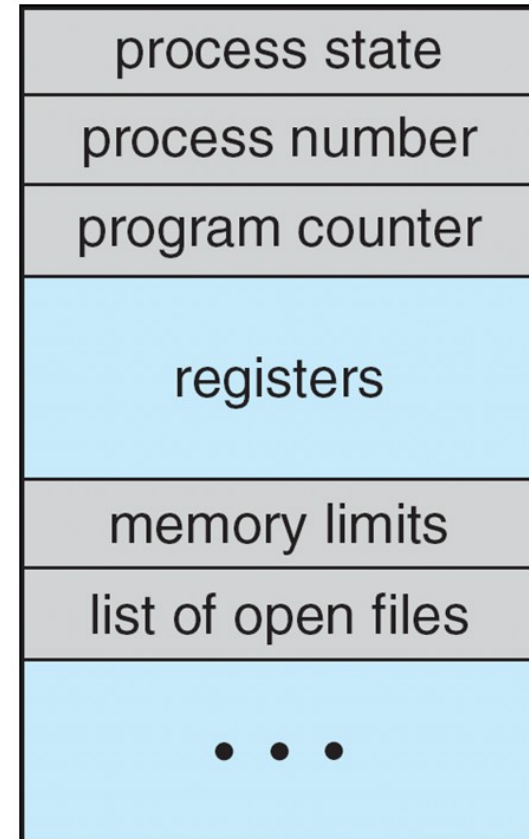
- Location of instruction to next execute

■ CPU Registers

- Contents of all process-centric registers: GP registers, stack register

■ CPU Scheduling Info

- Priorities, scheduling queue pointers





Process Control Block (PCB) (cont.)

■ Memory-Management Info

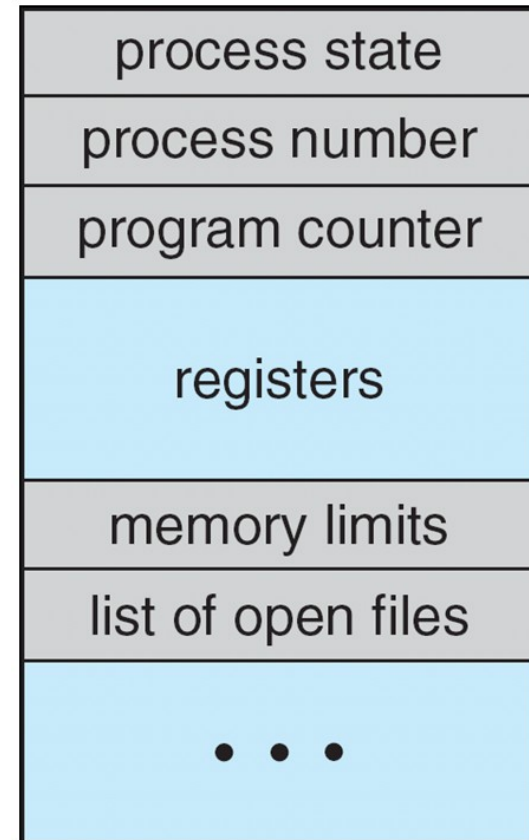
- Memory allocated to process
- Page or segment tables

■ Accounting Info

- CPU used
- Clock time elapsed since start
- Time limits

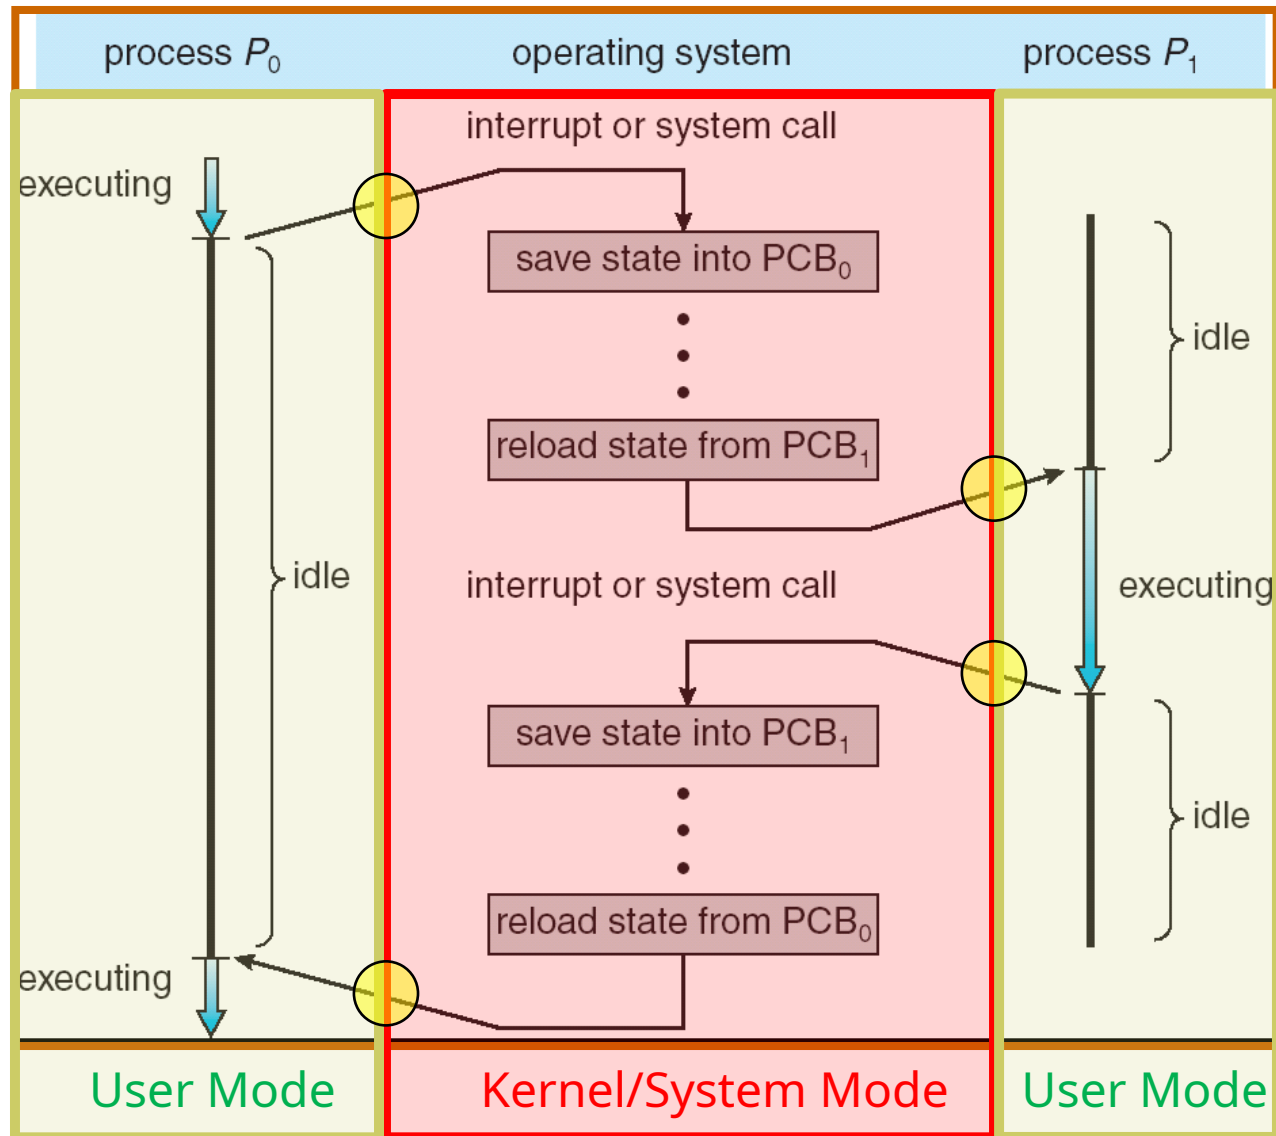
■ I/O Status Info

- I/O devices allocated to process
- List of open files





CPU Switch From Process to Process





Threads

- Originally, Process has a **Single Thread** of Execution
 - E.g, : cannot do different searches with **Mozilla** at **the same time** (or typing & spell check)
- Consider Having **Multiple** Program Counters per Process
 - Multiple locations can execute at once
 - Multiple threads of control → **threads**
- Must then Have Storage for Thread Details
 - **Multiple program counters** in PCB
 - Will be covered in detail next lecture



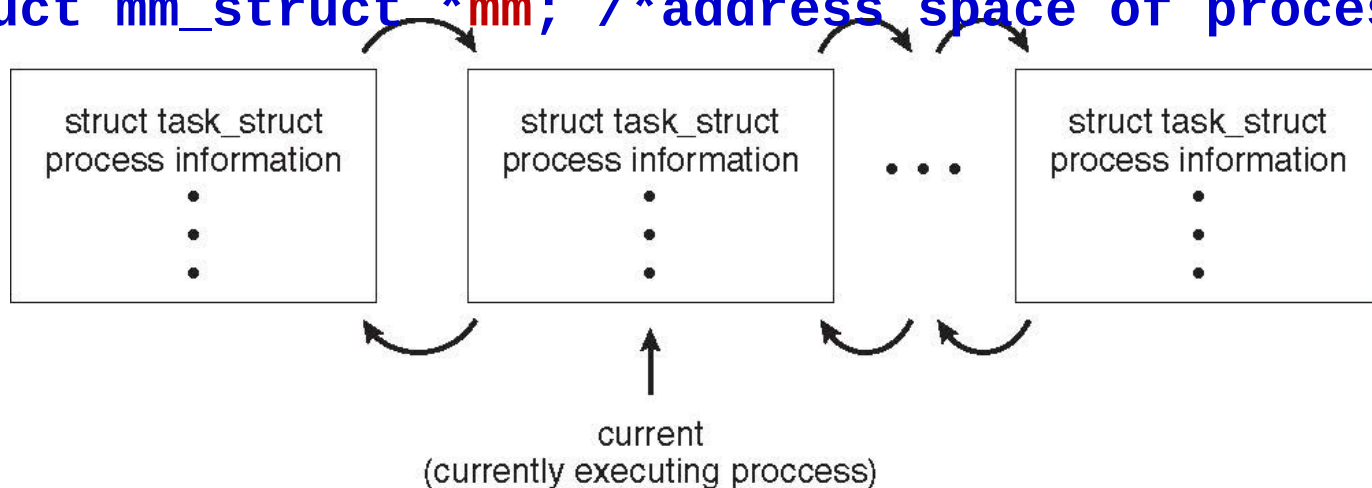


Process Representation in Linux

■ Doubly Linked List of Task_Struct

● Kernel Maintains a **Pointer** to Current Process

```
pid t_pid; /*process identifier*/  
long state; /*state of the process*/  
unsigned int time_slice /* scheduling information*/  
struct task_struct *parent; /*this process's parent */  
struct list_head children; /*this process's children*/  
struct files_struct *files; /*list of open files*/  
struct mm_struct *mm; /*address space of process*/
```





Process Scheduling

■ Process Scheduler

- Selects among available processes for next execution on CPU

■ Possible Goals of Process Scheduling

- To maximize CPU utilization
 - ▶ Quickly switch processes onto CPU for time sharing
- To meet deadline of each process
 - ▶ In real-time applications
- To meet fairness among processes





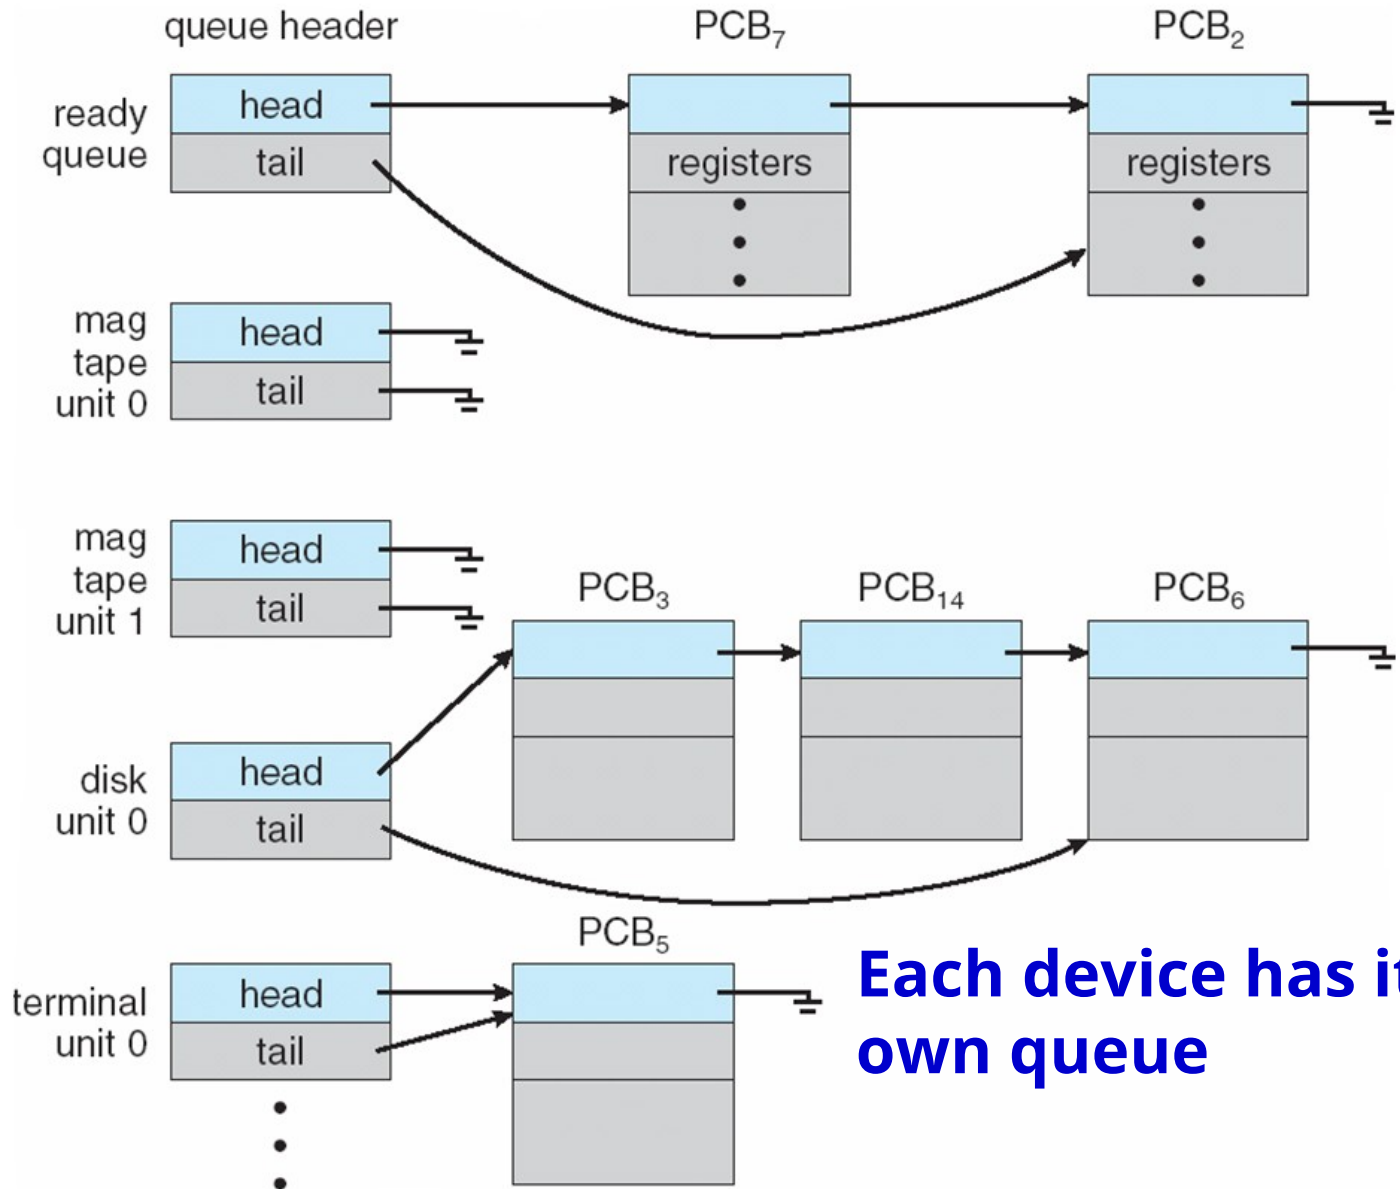
Process Scheduling (cont.)

- Maintains **Scheduling Queues** of Processes
 - **Job queue** – set of all processes in system
 - **Ready queue** – set of all processes residing in main memory, ready, and waiting to execute
 - ▶ Pointers to the 1st and final PCB in the list
 - **Device queues** – set of processes waiting for an I/O device (**one queue** for **each device**)
- Processes **Migrate** among Various Queues
 - Job queue → ready queue
 - Device queue → ready queue





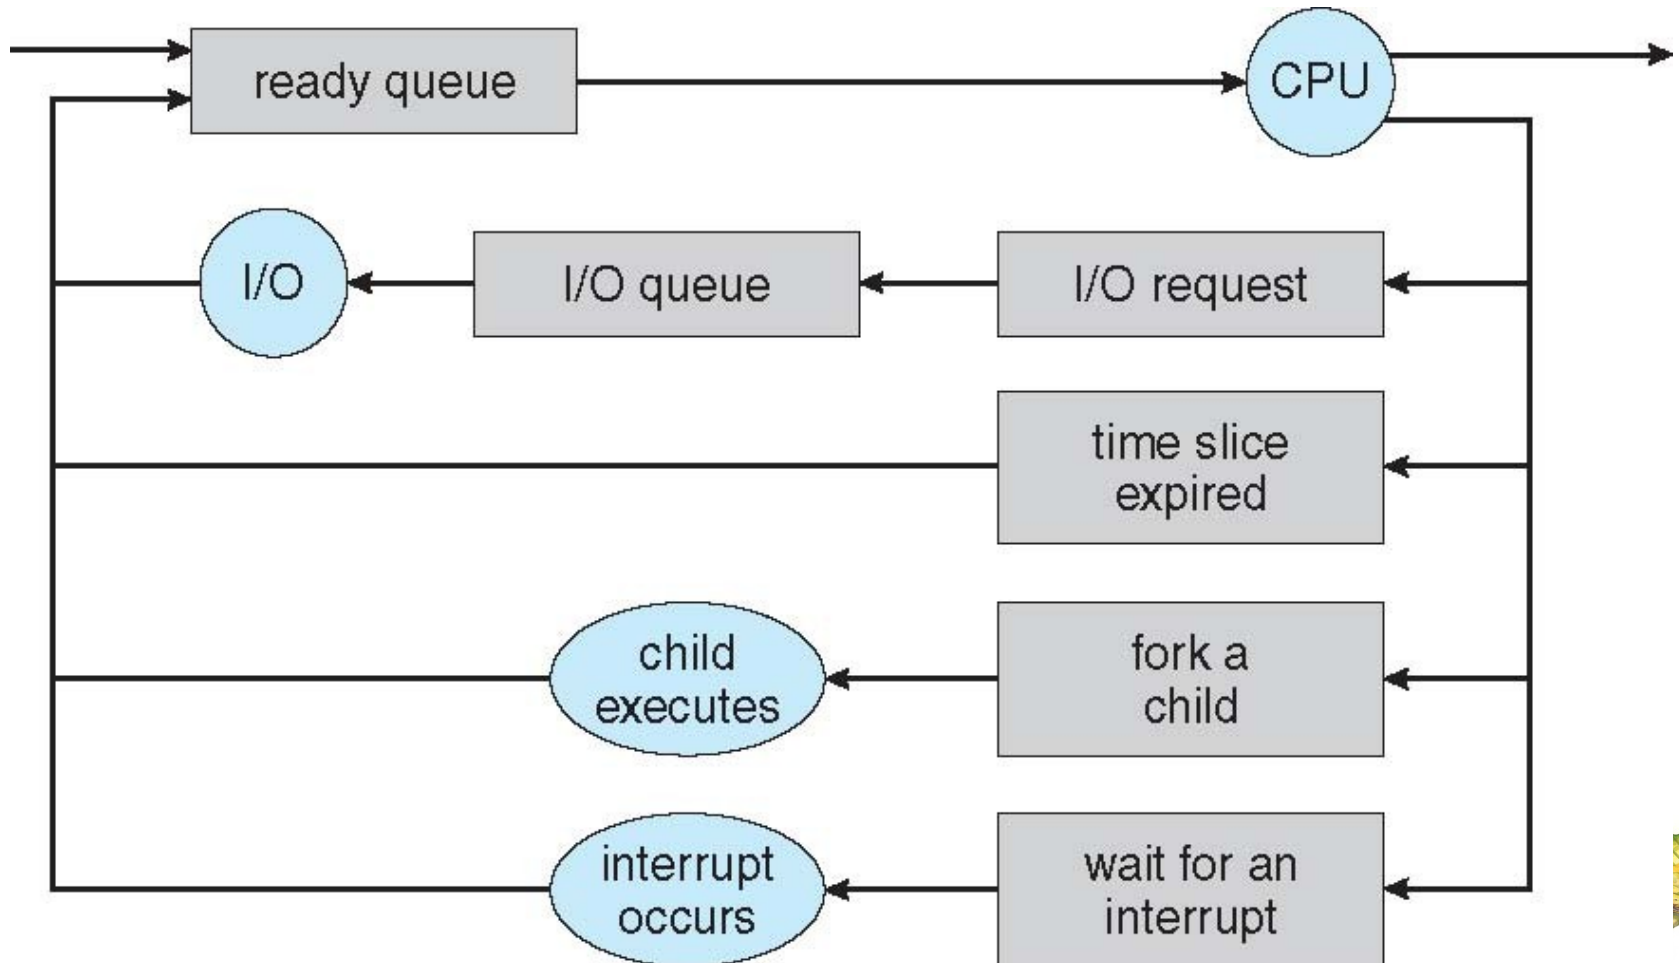
Ready Queue and Various I/O Device Queues





Representation of Process Scheduling


- **Queuing Diagram** represents queues, resources, and flows





Schedulers

- A Process **Migrates** Among various Scheduling Queues throughout its Lifetime
 - OS selects (schedules) processes in queues
- **Short-Term (ST) Scheduler** (CPU Scheduler)
 - Selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Is invoked frequently (milli-sec)
 - ▶ Must be fast
 - ▶ Example: scheduler takes 10ms to run and it is invoked every 100ms → ~9% OS overhead on CPU





Schedulers (cont.)

■ Long-Term (LT) scheduler (job scheduler)

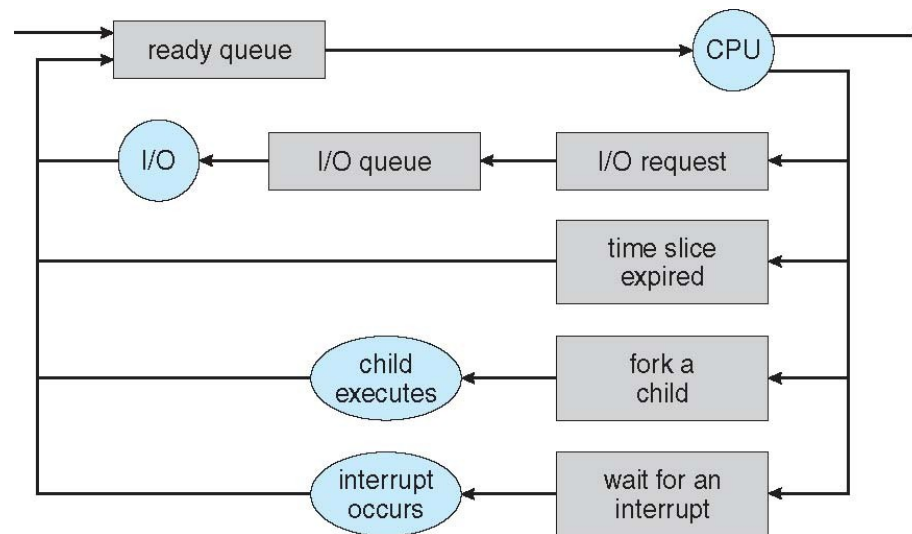
- Selects which processes should be brought into ready queue
- Invoked infrequently (sec~min) \Rightarrow (may be slow)
- Controls **degree of multiprogramming**
 - i.e., determines **no of processes in memory**
- **Steady-State:** long-term scheduler should be invoked only a process leaves the system
 - Assuming degree of multiprogramming is stable
- Mostly used in **HPC & clustered Systems**





Schedulers (cont.)

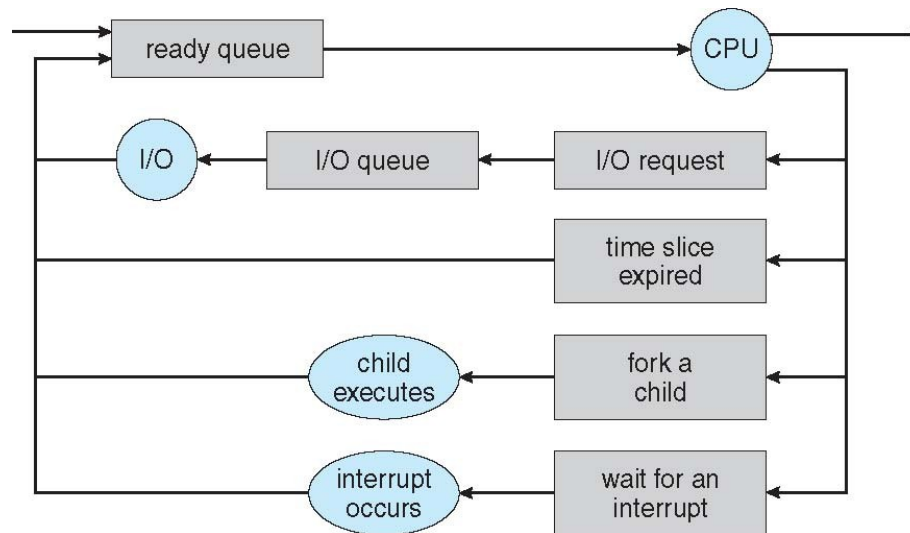
- Processes can be Described as Either:
 - I/O-bound process** – spends more time doing I/O than computations
 - CPU-bound process** – spends more time doing computations





Schedulers (cont.)

- LT Scheduler Strives for Good *process mix*
 - All processes I/O bound → ready queue will almost always be empty → ST scheduler idle
 - All processes CPU bound → I/O waiting queue will almost always be empty → devices idle





Schedulers (cont.)

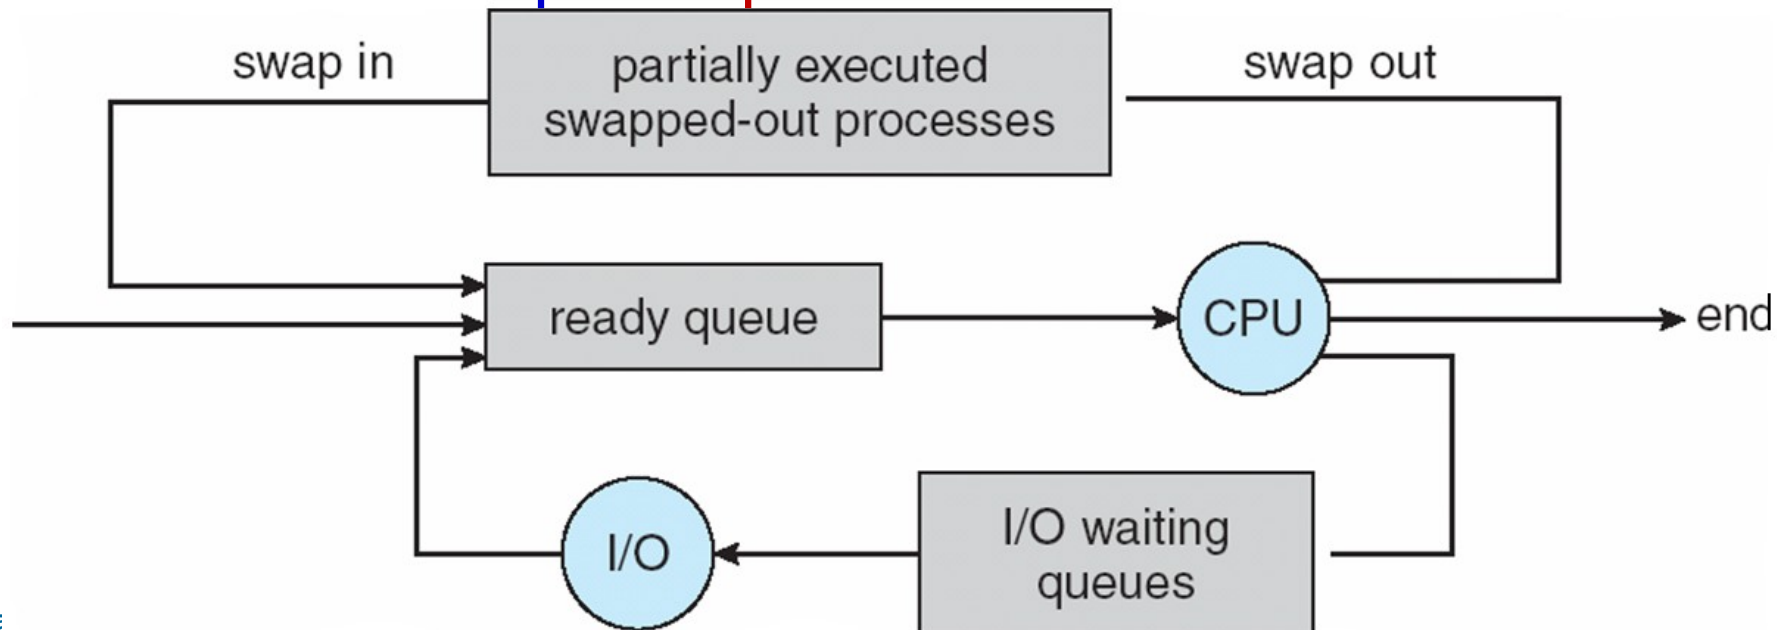
- LT Scheduler May be Absent or Minimal
 - E.g., UNIX and Windows have no LT scheduler and put every new process in memory
 - → Can adversely affect performance
 - → So, users may quit and their processes are terminated





Addition of Medium Term Scheduling


- **Medium-Term Scheduler** Used if Degree of Multi-Programming Needs to Decrease
 - **Swapping:** remove process from memory, store on disk (**swapped out**), bring back in from disk to continue execution (**swapped in**)
 - Used to improve **process mix**





Context Switch

- When CPU Switches to another Process
 - → System must **save state** of old process
 - → OS, then, loads **saved state** for new process via a **context switch**
- **Context** of a Process Represented in PCB
- Context-Switch Time is Overhead
 - System does **no useful work** while switching
 - More complex OS and PCB → longer context switch
 - More context switching time when using VMs





Context Switch (cont.)

- Context-Switch Time Depends on HW Support
 - Memory speed
 - # of registers in RF
 - Special instruction to copy RF
 - ▶ Also, HW (micro-architecture) support
 - Some HW provides multiple sets of registers per CPU → multiple contexts loaded at once
 - ▶ Context switch: changing a pointer to a target RF





Operations on Processes

- System must Provide Mechanisms for
 - Process creation
 - Process termination
 - and so on as detailed next





Process Creation

- **Parent** Process Creates **Children** processes
 - Which, in turn creates other processes, forming a **tree** of processes
- Process Identified and Managed via a **Process Identifier (pid)**
- Root Process in Linux: “init”
- Root Process in Solaris: “sched”
 - Children processes: “Init”, “pageout”, “fsflush”
 - ▶ “Init” root of user processes

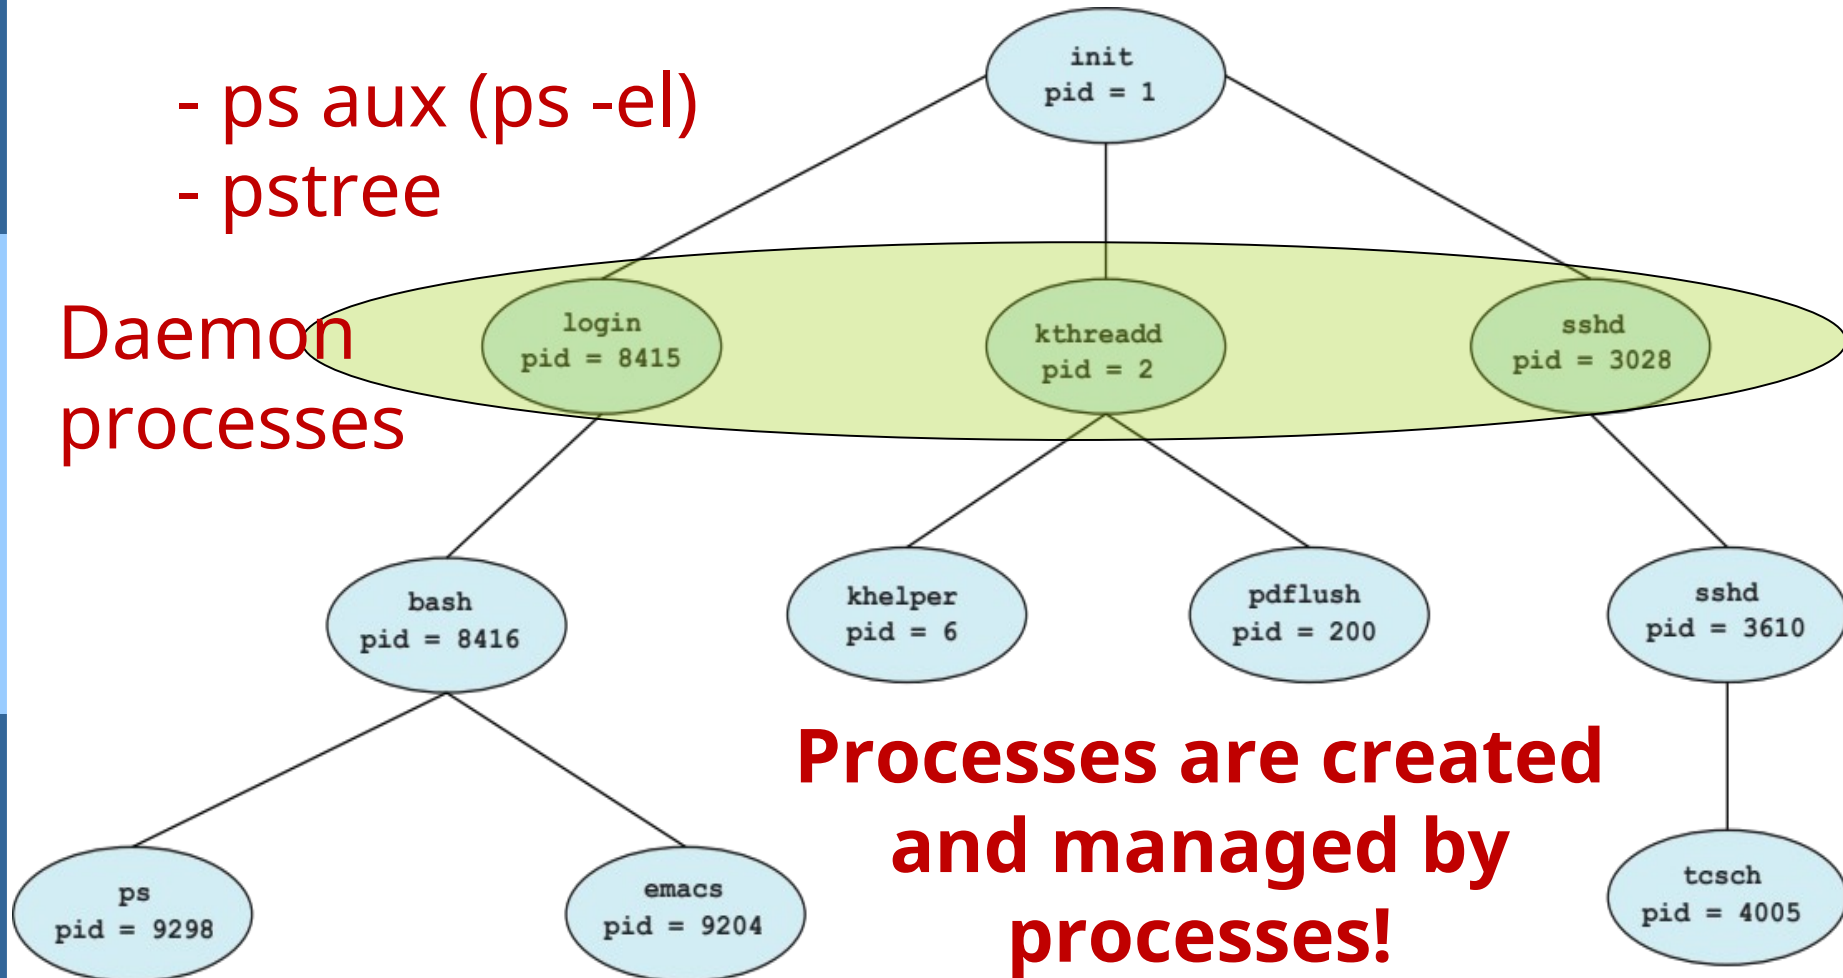




A Tree of Processes in Linux

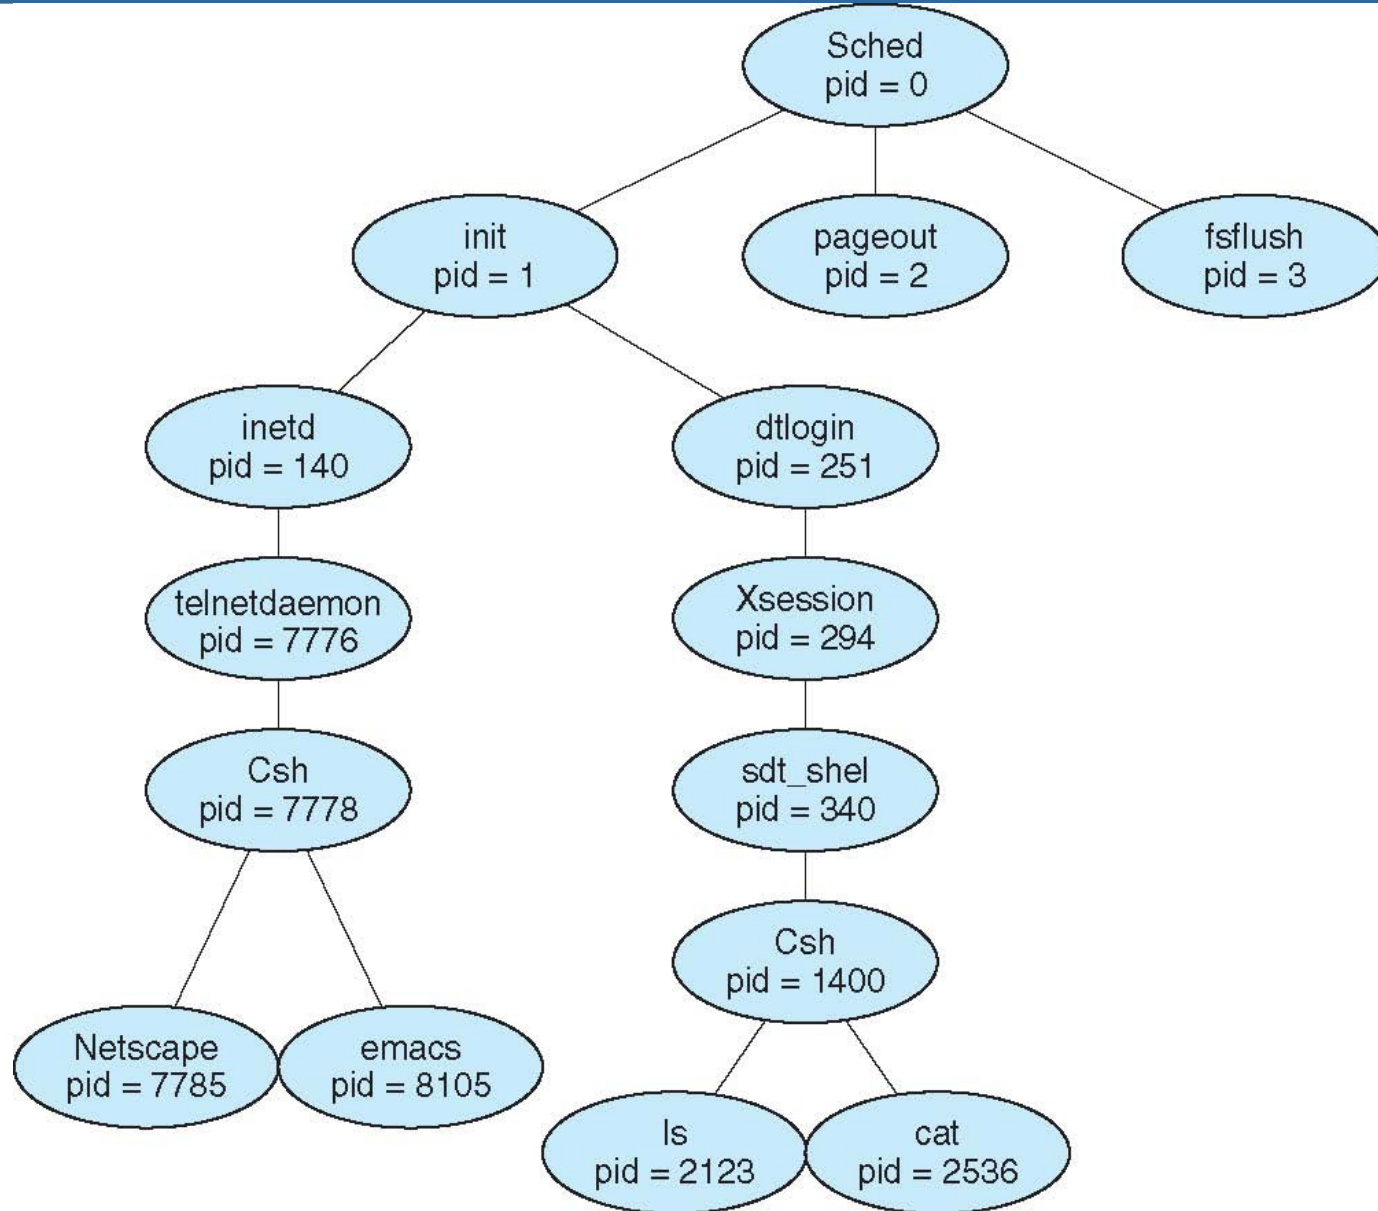
- ps aux (ps -el)
- pstree

Daemon
processes





A Tree of Processes in Solaris





Process Creation (cont.)

■ Resource **Sharing Options**

- Parent and children **share all resources**
- Children **share subset** of parent's resources
 - ▶ Prevents any process from overloading system by creating too many sub-processes
- Parent and child **share no resources**

■ **Execution Options**

- Parent and children execute concurrently
- Parent waits until children terminate





Process Creation (Cont.)

■ Address Space

- Child duplicate of parent
- Child has a new program loaded into it

■ UNIX Examples

- **fork()** system call creates new process
 - ▶ New process consists of a copy of address space of original process
 - ▶ Convenient communication between two processes
- **exec()** system call used after a **fork()** to replace process' memory space with a new program

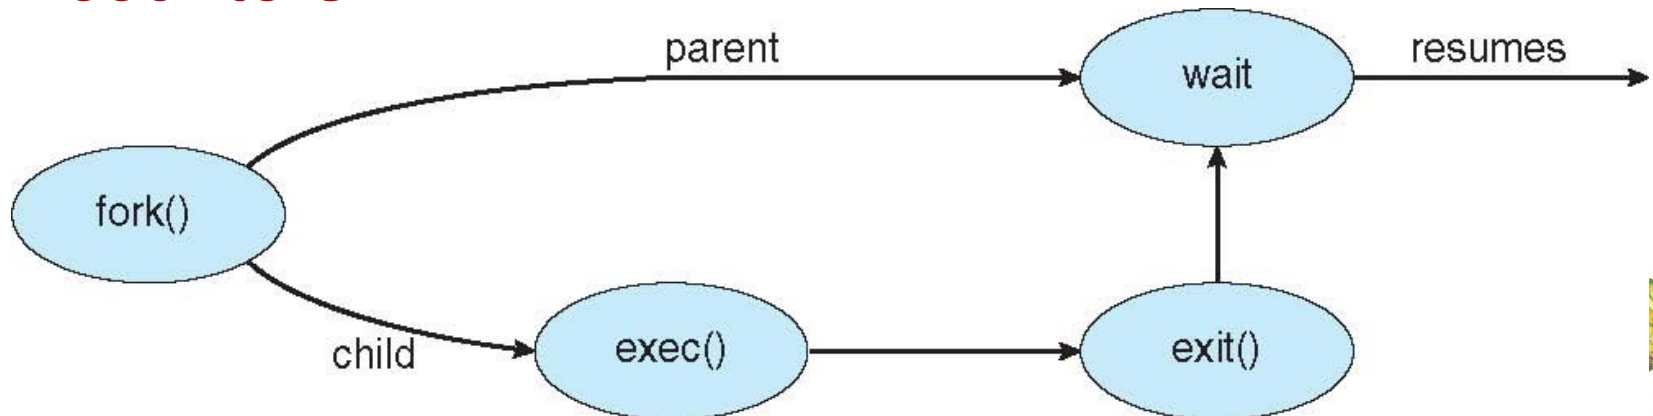




Process Creation (Cont.)

■ What Happens after “fork ()”?

- Both processes continue execution at the instruction after “fork”
- Child gets unique process ID
- Child's PPID = parent's PID
- Reset child's **resource utilization** and **CPU time counters**





Process Creation (Cont.)

■ Sample Code 1

■ Process status "Sleep" → "Running"

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main() {
    printf("Hello \n");
    fork(); //make a child process of same type
    printf("pid = %d \n", getpid());
    sleep(10);
    int i=0;
    while (1) i+=1;
    printf("Sample code for fork \n");
    return 0;
}
```





Process Creation (Cont.)

- `pid_t fork()` – copy current process
 - New process has different pid
 - New process contains a single thread
- Return value from **fork()**: pid (like an integer)
 - When > 0 :
 - Running in (original) **Parent** process
 - return value is **pid** of new child
 - When $= 0$:
 - Running in new **Child** process
 - When < 0 :
 - Error! Must handle somehow
 - Running in original process





Process Creation (cont.)

Sample Code 2: Checking process ID

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
int main(int argc, char *argv[]) {
    pid_t cpid, mypid;
    pid_t pid = getpid(); /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) { /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) { /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
    }
}
```

Operating System Concepts – 9th Edition

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
int main(int argc, char *argv[]) {
    pid_t cpid, mypid;
    pid_t pid = getpid(); /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) { /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) { /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
    }
}
```

Operating System Concepts – 9th Edition





Process Creation (Cont.)

- What Does not Child Inherits from Parent?
 - Parents memory locks
 - Parent's timers
 - Semaphore's adjustments and pending signals
- What Child Process Inherits from Parent?
 - Privileges and scheduling attributes
 - Certain resources such as open files
 - Address space





Process Creation (Cont.)


■ Sample Code 3: Possible Race?

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
int main(int argc, char *argv[]) {
    int i;
    pid_t cpid = fork();
    if (cpid > 0) {
        for (i = 0; i < 10; i++) {
            printf("Parent: %d\n", i);
            //sleep(1);
        }
    } else if (cpid == 0) {
        for (i = 0; i > -10; i--) {
            printf("Child: %d\n", i);
            //sleep(1);
        }
    }
}
```

- What if we change 10 to 1000?

- Would adding the calls to `sleep()` matter?





C Program Forking Separate Process

Sample Code 4: Running Separate Process

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
```

```
#include <sys/wait.h>int main(int argc, char *argv[]) {
```

```
int main() {
```

```
    pid_t pid;
```

```
    long unsigned i=0;
```

```
    pid = fork();
```

```
    if (pid<0){
```

```
        fprintf(stderr, "Fork Failed"); return 1;
```

```
    }
```

```
    else if (pid==0) { /* child process */
```

```
        printf("Running child process:\n");
```

```
        execlp("/bin/ls", "ls", NULL, \
```

```
    }
```

```
    else /*parent process*/
```

```
    {
```

```
        printf("waiting for the c
```

```
        wait (NULL);
```

```
        printf("Running parent pr
```

```
        printf("Child completed \
```

```
    }
```

```
    return 0;
```

```
        if (pid<0){
```

```
            fprintf(stderr, "Fork Failed"); return 1;
```

```
        }
```

```
        else if (pid==0) { /* child process */
```

```
            printf("Running child process:\n");
```

```
            execlp("/bin/ls", "ls", NULL);
```

```
        }
```

```
        else /*parent process*/
```

```
        {
```

```
            printf("waiting for the child process:\n
```


```
            wait (NULL);
```

```
            printf("Running parent process:\n");
```

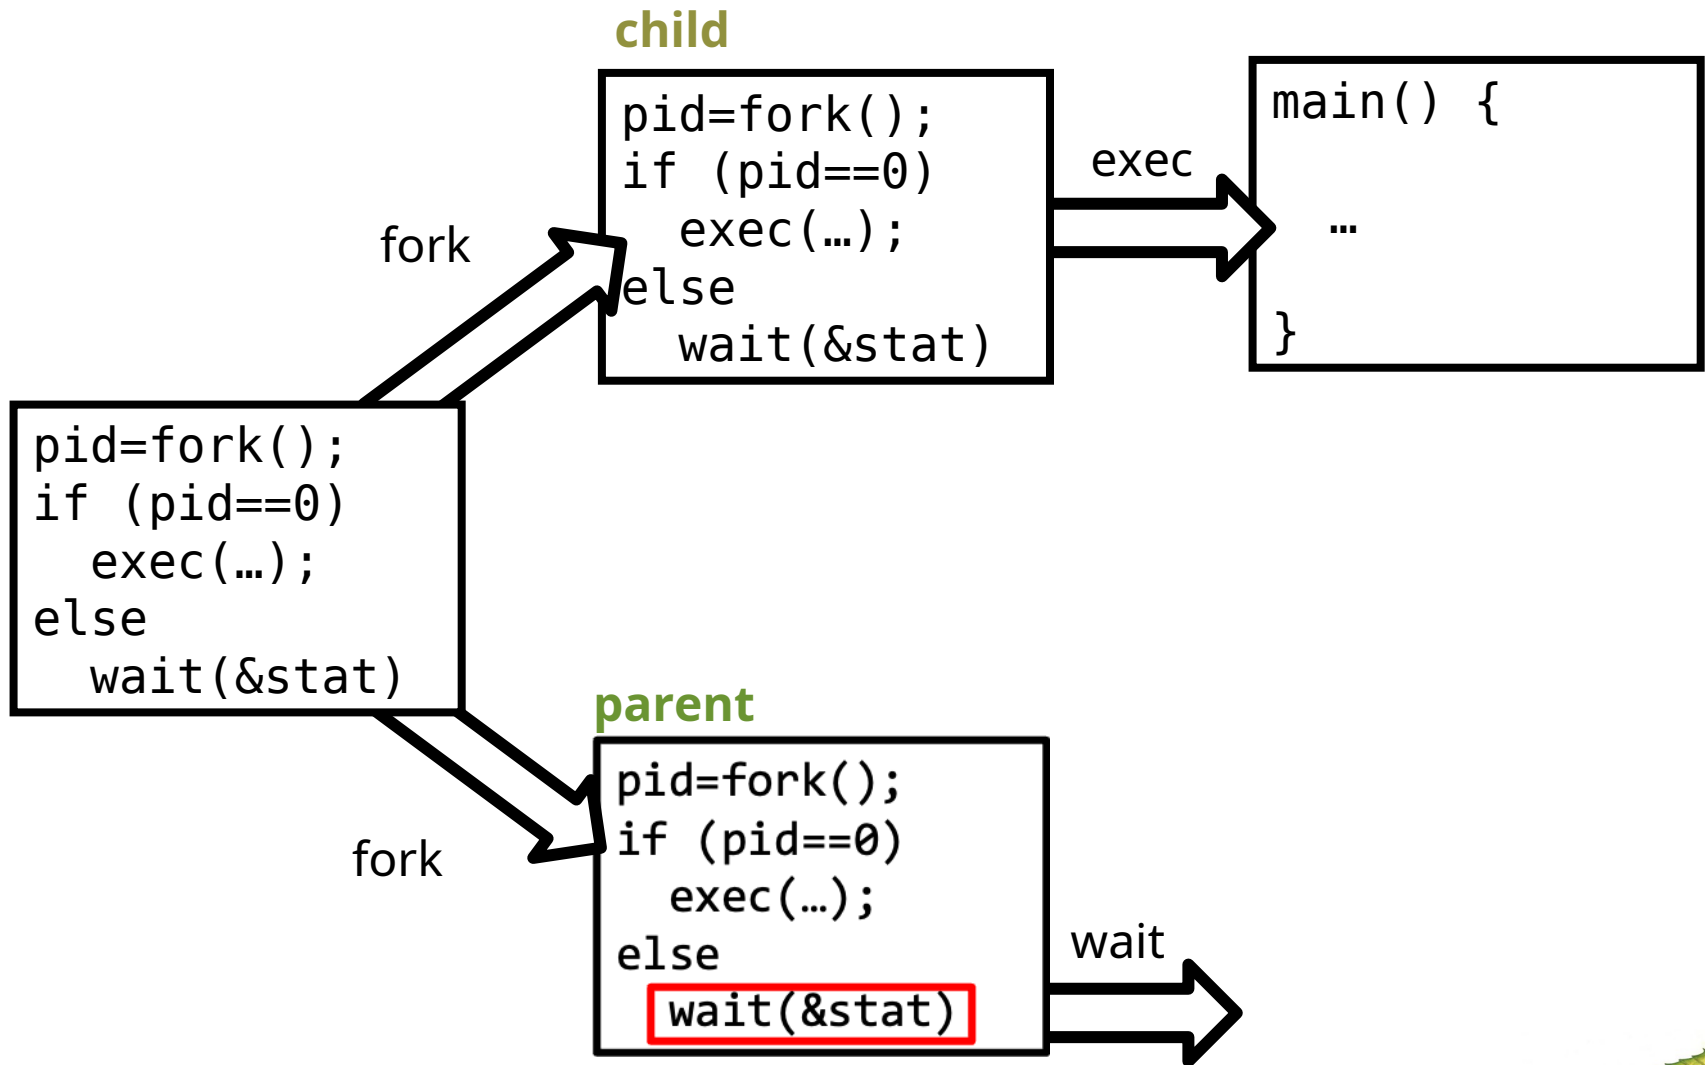
```
            printf("Child completed \n");
```

```
        }
```

```
    return 0;
```

C Program Forking Separate Process





Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```






Process Termination

■ Process Executes Last Statement

- Then asks OS to delete it
 - ▶ Using **exit()** system call
- Returns status data from child to parent via **wait()**
- Process' resources are de-allocated by OS

■ Parent may Terminate Execution of Children Processes Using **abort()** System Call

- Or TerminateProcess() in Win32





Process Termination (cont.)

■ Reasons to Terminate Child Processes

- Child has exceeded allocated resources
- Task assigned to child is no longer required
- Parent is exiting and OS does not allow a child to continue if its parent terminates
- Or ...





Process Termination (cont.)

- Some OSes do not Allow Child to Exist
 - If its parent has terminated
 - If a process terminates, then all its children must also be terminated
 - **Cascading termination**
 - ▶ All children, grandchildren, etc. are terminated
 - Termination is initiated by OS
- Parent process may Wait for Termination of a Child Process by Using **wait()** System Call





Process Termination (cont.)

- Call Returns Status Info and pid of Terminated Process

`pid = wait(&status);`

- Zombie Process

- If no parent waiting (did not invoke **`wait()`**)
- Process **has completed execution** but still has an **entry in process table**
- Entry needed for possible reading of exit status by its parent





Process Termination (cont.)

■ Orphan Process

- Parent terminated without Invoking **wait**
- In UNIX, any orphan process is immediately **adopted** by “**init**” process
 - ▶ Called “**re-parenting**”
- A process can become orphan intentionally or unintentionally
 - ▶ Process crash
 - ▶ To run a process indefinitely (in the background)





Process Termination (cont.)

■ Sample Code 5:

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main() {
    printf("Hello \n");
    int pid=fork();
    printf("pid = %d \n", getpid());
    sleep(10);
    printf("Sample cod for fork \n");
    return 0;
```

■ Two Scenarios }

- Case 1: Terminate parent process → reparenting
- Case 2: Terminate child process → Zombie process





Process Termination (cont.)

■ Example Code 6:

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main() {
    printf("Hello \n");
    int pid=fork();
    printf("pid = %d \n", getpid());
    sleep(20);
    printf("Sample cod for fork \n");
    wait(NULL);
    sleep(10);
    return 0;
}
```

■ Test Scenario

- Terminate child process → Zombie process → removed after wait call by parent





Process Termination (cont.)

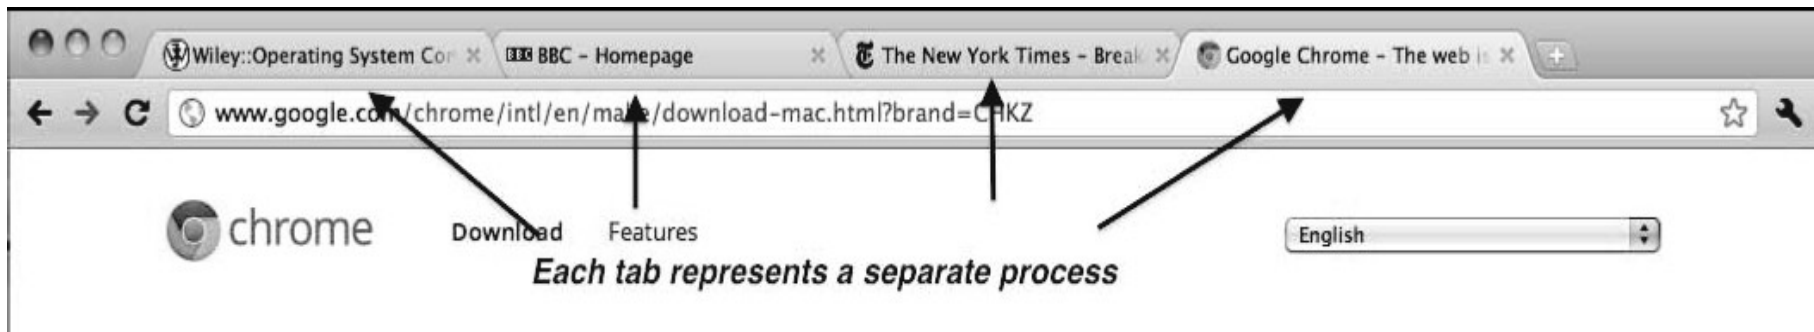
- Other Variations of Wait
 - waitid()
 - waitpid()
- Waits for specific child process





Multiprocess Architecture – Chrome Browser

- Many Web Browsers Ran as **Single Process** (Some Still Do)
 - If one web-site causes trouble, entire browser can hang or crash
- Google Chrome Browser is **Multi-Process** with Three Different Types of Processes





Multiprocess Architecture – Chrome Browser

- Google Chrome Browser is Multi-Process with Three Different Types of Processes
 - **Browser** process manages user interface, disk and network I/O
 - **Renderer** process renders web pages, deals with HTML, Javascript.
 - A new renderer created for each website opened
 - ▶ Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in** process for each type of plug-in





Interprocess Communication

■ Processes: *Independent* or *Cooperating*

■ *Independent* Process

- Cannot affect or be affected by execution of another Process
- Does not share data with any other processes

■ *Cooperating* Process can Affect or be Affected by Execution of another Process

- Including shared data






Cooperating Processes

■ Motivations of Process Cooperation

- Information sharing
 - ▶ E.g., shared file
- Computation speed-up
 - ▶ Breaking into subtasks & executing on multiple cores
 - ▶ E.g., computing Pi
- Modularity
 - ▶ Constructing a system in a modular fashion
 - ▶ Dividing system functions into separate processes
- Convenience
 - ▶ Several tasks of a single user (editing, printing, ...)





Interprocess Communication (cont.)

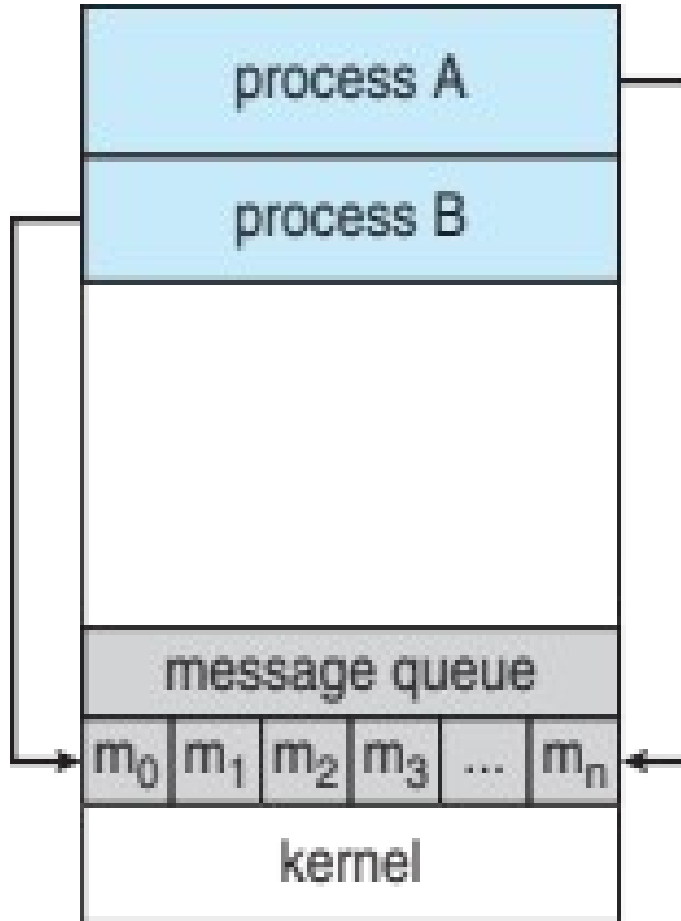
- Cooperating Processes need **Interprocess Communication (IPC)**
- Two Models of IPC
 - **Shared memory**
 - **Message passing**
- Most OSes Implement Both Models






Communications Models

(a) **Message passing**



(a)

(b) **Shared memory**



(b)





Interprocess Communication (cont.)

■ Shared memory

- Maximum speed
- Convenience of communication
- System calls used only to establish shared-memory regions
 - ▶ Thereafter, all accesses are treated as normal memory accesses

■ Message passing

- Useful for exchanging smaller amounts of data
- No conflicts need to be avoided
- Easier to be implemented for inter-computer communication
- Implemented by system calls → time-consuming





Shared-Memory Solution: Producer-Consumer Problem

■ Paradigm for Cooperating Processes

- *Producer* process produces info that is consumed by a *consumer* process
- **Unbounded-buffer** places no practical limit on size of buffer
- **Bounded-buffer** assumes that there is a fixed buffer size





Bounded-Buffer: Shared-Memory Solution

■ Shared data

```
#define BUFFER_SIZE 10  
typedef struct {  
    . . .  
} item;  
  
item buffer[BUFFER_SIZE];  
int in = 0;  
int out = 0;
```

■ Solution is correct, but can only use BUFFER_SIZE-1 elements





Bounded-Buffer: Producer

```
item next_produced;  
while (true) {  
    /* produce an item in next produced */  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
}
```





Bounded Buffer: Consumer

```
item next_consumed;  
while (true) {  
    while (in == out)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    /*consume item in next consumed*/  
}
```





Interprocess Communication: Shared Memory

■ Main Idea

- An area of memory shared among processes that wish to communicate

■ Communication is under Control of Users Processes not OS

■ Major Issue

- To provide mechanism that will allow user processes to **synchronize** their actions when they access shared memory
- Synchronization will be discussed in details in next lectures





Interprocess Communication: Message Passing

- Mechanism for Processes to Communicate and to Synchronize their Actions
- Processes Communicate with each other
 - Without resorting to shared variables
 - Without having a shared address space
- Significantly useful in a Distributed Systems
 - E.g., a chat program
- Two Major Operations
 - **send**(*message*)
 - **receive**(*message*)





Message Passing (Cont.)

- If Processes P and Q Wish to Communicate, they need to:
 - Establish a **communication link** between them
 - Exchange messages via send/receive
- *Message* size is either **Fixed** or **Variable**
- **Fixed** Message Size
 - Straightforward implementation 😊
 - Difficulty for programmer ☹️





Message Passing (Cont.)

■ Implementation Issues:

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is capacity of a link?
- Is size of a message that link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?





Message Passing: Communication Link

- Communication Link can be Viewed at either
 - Physical or Logical
- Physical
 - Shared memory
 - Hardware bus
 - Network
- Logical
 - Direct or indirect
 - Synchronous or asynchronous
 - Automatic or explicit buffering





Direct Communication

- Processes must Name each other Explicitly
 - **send** (P, msg): send a message to process P
 - **receive**(Q, msg): receive a message from Q
- Properties of Direct Communication Link
 - Links are **established automatically**
 - A **link** is associated with **exactly one pair** of communicating processes
 - Between **each pair** there exists **exactly one link**
 - Link may be **unidirectional**, but is usually **bidir**
 - ▶ **Bidirectional**: symmetric (example as above)
 - ▶ **Directional**: Asymmetric





Indirect Communication

- Main Issue in Direct Communication
 - Identifiers must **explicitly stated** (e.g, process P or Q)
- Messages are Directed and Received from Mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of an **Indirect** Communication Link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bidirectional





Indirect Communication

■ Operations

- Create a new mailbox (port)
- Send and receive messages through mailbox
- Destroy a mailbox

■ Primitives are Defined as:

send(*A, message*) – send a message to mailbox A

receive(*A, message*) – receive a message from mailbox A





Indirect Communication (cont.)

■ Mailbox Sharing

- P_1 , P_2 , and P_3 share mailbox A
- P_1 sends; P_2 and P_3 receive
- Who gets message?

■ Solutions

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow system to select arbitrarily receiver
- Sender is notified who receiver was





Synchronization

- Message Passing may be either Blocking or Non-Blocking
- **Blocking** considered **Synchronous**
 - **Blocking send** -- sender is blocked until message is received
 - **Blocking receive** -- receiver is blocked until a message is available





Synchronization (cont.)

■ Non-Blocking Considered **Asynchronous**

- **Non-blocking send** -- sender sends message and continue
- **Non-blocking receive** -- receiver receives:
 - A valid message, or
 - Null message

■ Different Combinations Possible

- If both send and receive are blocking → we have a **rendezvous**





Synchronization (Cont.)

■ Producer-Consumer becomes Trivial

```
message next_produced;  
while (true) {  
    /*produce an item in next produced*/  
    send(next_produced);  
}
```

```
message next_consumed;  
while (true) {  
    receive(next_consumed);  
    /*consume the item in next consumed*/  
}
```





Buffering

■ Implemented in one of Three Ways

1. **Zero capacity** – no messages are queued on a link
Sender must wait for receiver (rendezvous)
2. **Bounded capacity** – finite length of n messages
Sender must wait if link full
3. **Unbounded capacity** – infinite length
Sender never waits



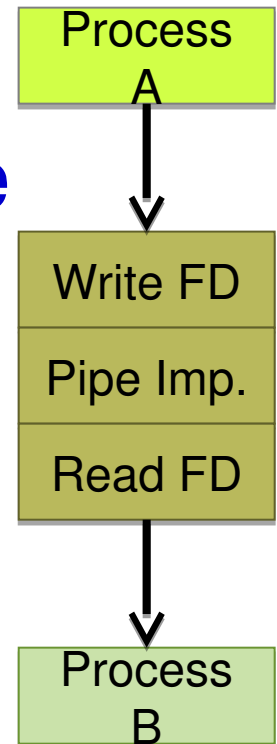


Linux Pipes

Linux `pipe()` system call:

- Allows *anonymous* (un-named) communication
- Unidirectional
- Produces / consumes data through file syscalls `write()` and `read()`
- Data stored in kernel via *pipefs* in virtual filesystem

`/fs/pipe.c` for implementation





Named Pipes or FIFO

- Created with `mkfifo()` library function
- Handle to FIFO exists as a regular file
- Read and written like regular file
- Allows non-related processes to communicate
- Must be open at both ends before reading or writing
- Even though FIFOs have a handle in regular file system, they are not files!





Limits and Best Practice

■ Pipe: used between parent and child processes

■ FIFO: used across unrelated processes

■ **Limits:**

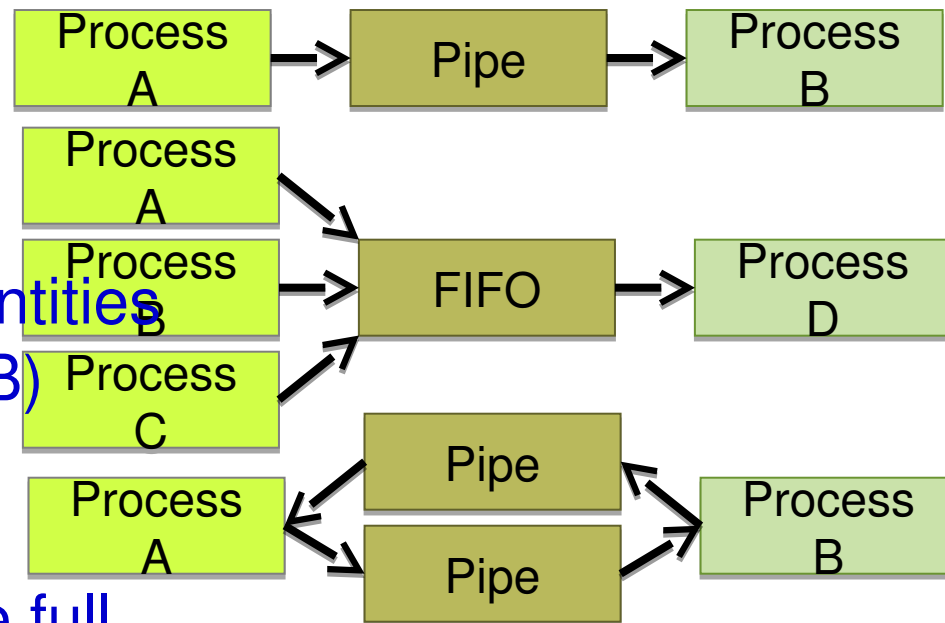
■ Atomicity

- I/O is atomic for data quantities less than PIPE_BUF (4KB)

■ Write capacity: 64K

- Writers block or fail if pipe full

■ Polling vs. Blocking: Readers may block or fail based on flags set during pipe creation





Pipe Sample Code

```
#include<stdio.h>
#include<unistd.h>
int main() {
    int pipefds[2];
    int returnstatus;
    int pid;
    char writemessages[2][20]={"Hi", "Hello"};
    char readmessage[20];
    returnstatus = pipe(pipefds);
    if (returnstatus == -1) {
        printf("Unable to create pipe\n");
        return 1;
    }
    pid = fork();
    // Child process
```

```
    if (pid == 0) {
        read(pipefds[0], readmessage,
            sizeof(readmessage));
        printf("Child Process - Reading from pipe –
            Message 1 is %s\n", readmessage);
        read(pipefds[0], readmessage,
            sizeof(readmessage));
        printf("Child Process - Reading from pipe –
            Message 2 is %s\n", readmessage);
    } else { //Parent process
        printf("Parent Process - Writing to pipe -
            Message 1 is %s\n", writemessages[0]);
        write(pipefds[1], writemessages[0],
            sizeof(writemessages[0]));
        printf("Parent Process - Writing to pipe -
            Message 2 is %s\n", writemessages[1]);
        write(pipefds[1], writemessages[1],
            sizeof(writemessages[1]));
    }
    return 0;
}
```





Pipe: Two-Way Sample Code

```
#include<stdio.h>
#include<unistd.h>
int main() {
    int pipefds1[2], pipefds2[2];
    int returnstatus1, returnstatus2;
    int pid;
    char pipe1writemessage[20] = "Hi";
    char pipe2writemessage[20] = "Hello";
    char readmessage[20];
    returnstatus1 = pipe(pipefds1);
    if (returnstatus1 == -1) {
        printf("Unable to create pipe 1 \n");
        return 1;
    }
    returnstatus2 = pipe(pipefds2);
    if (returnstatus2 == -1) {
        printf("Unable to create pipe 2 \n");
        return 1;
    }
    pid = fork();
```

```
    if (pid != 0) // Parent process {
        close(pipefds1[0]); // Close the unwanted pipe1 read side
        close(pipefds2[1]); // Close the unwanted pipe2 write side
        printf("In Parent: Writing to pipe 1 – Message is %s\n",
            pipe1writemessage);
        write(pipefds1[1], pipe1writemessage,
            sizeof(pipe1writemessage));
        read(pipefds2[0], readmessage, sizeof(readmessage));
        printf("In Parent: Reading from pipe 2 – Message is %s\n",
            readmessage);
    } else { //child process
        close(pipefds1[1]); // Close the unwanted pipe1 write side
        close(pipefds2[0]); // Close the unwanted pipe2 read side
        read(pipefds1[0], readmessage, sizeof(readmessage));
        printf("In Child: Reading from pipe 1 – Message is %s\n",
            readmessage);
        printf("In Child: Writing to pipe 2 – Message is %s\n",
            pipe2writemessage);
        write(pipefds2[1], pipe2writemessage,
            sizeof(pipe2writemessage));
    }
}
```

```
return 0;
```





Examples of IPC Systems

■ Reading Assignments

- POSIX shared memory
- Message passing in Mach OS
- Winx XP (message passing/shared memory)
- Sockets
- Remote Procedure Calls
- Pipes
- Remote Method Invocation (Java)






Examples of IPC Systems - POSIX

■ POSIX Shared Memory

- Process first creates shared memory segment
`shm_fd = shm_open(name, 0_CREAT | 0_RDWR, 0666);`
 - Also used to open existing segment to share it
 - Processes wish to access shared memory must attach it to their address space using `mmap` or
`shared_mem = (char*) shmat(id, NULL, 0);`
 - Now process could write to shared memory
`sprintf(shared memory, "Writing to shared memory");`
- Finally, to detach: `shmdt(shared_mem)`





IPC POSIX Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0);
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);

    return 0;
}
```





IPC POSIX Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```





Examples of IPC Systems - Mach

- Mach Communication is Message Based
 - Even system calls are messages
 - Each task gets two mailboxes at creation- Kernel & Notify
 - Only three system calls needed for message transfer
msg_send(), msg_receive(), msg_rpc()
 - Mailboxes needed for communication, created via
port_allocate()
 - Send and receive are flexible, e.g. 4 options if mailbox full:
 - ▶ Wait indefinitely
 - ▶ Wait at most *n* milliseconds
 - ▶ Return immediately
 - ▶ Temporarily cache a message





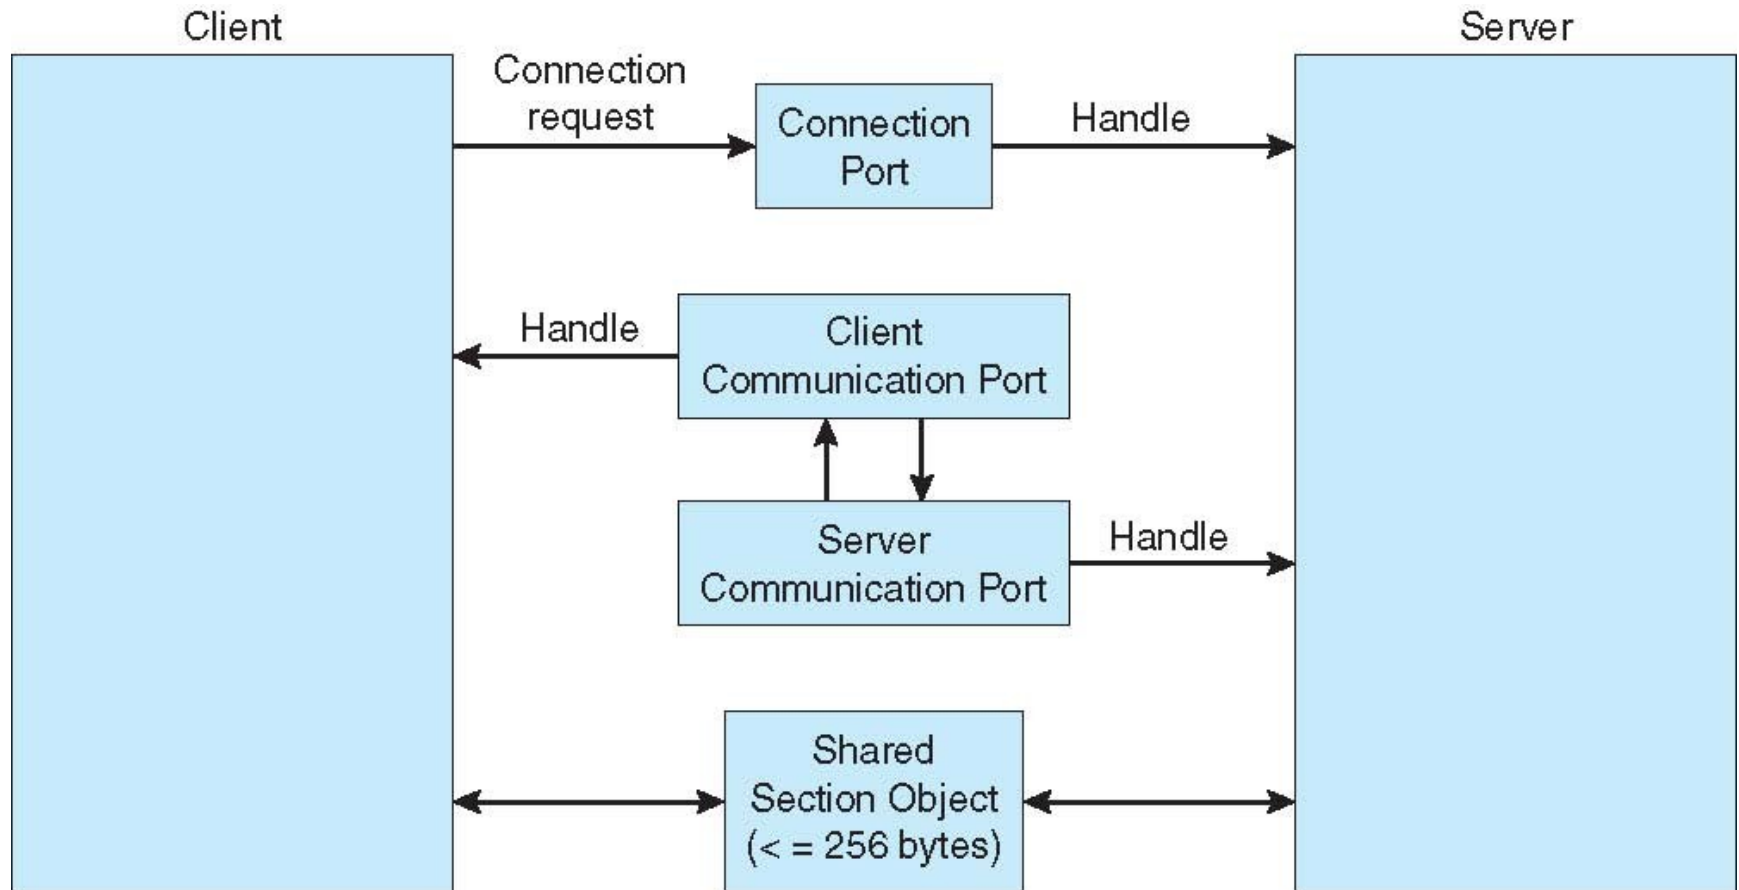
Examples of IPC Systems – Windows

- Message-Passing Centric via **Advanced Local Procedure Call (LPC)** facility
 - Only works between processes on the same system
 - Uses ports (like mailboxes) to establish and maintain communication channels
 - Communication works as follows:
 - ▶ Client opens a handle to the subsystem's **connection port** object
 - ▶ Client sends a connection request
 - ▶ Server creates two private **communication ports** and returns the handle to one of them to the client
 - ▶ Client and server use the corresponding port handle to send messages or callbacks and to listen for replies





Local Procedure Calls in Windows





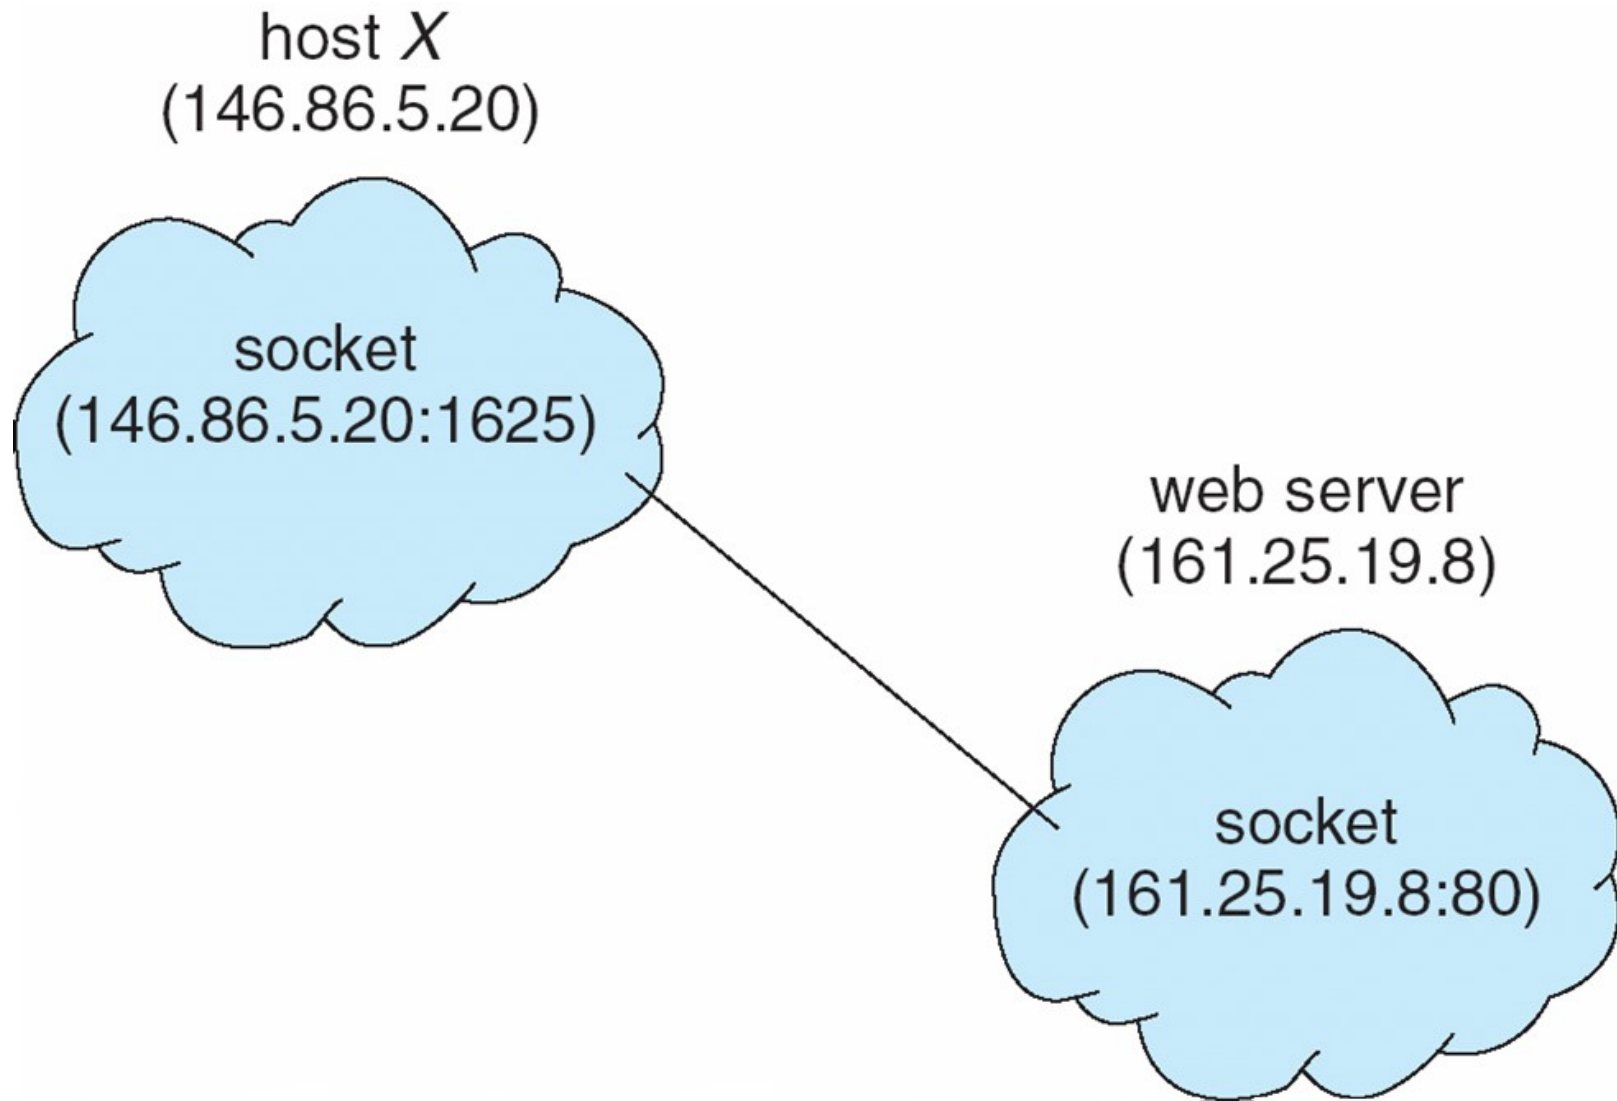
Sockets

- **Socket:** an Endpoint for Communication
- Concatenation of IP address and **Port**
 - A number included at start of message packet to differentiate network services on a host
 - Socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All Ports below 1024 are **Well Known**
 - Used for standard services
- Special IP address 127.0.0.1 (**loopback**)
 - Refers to system on which process is running





Socket Communication





Sockets in Java

- 3 Types of Sockets
 - **Connection-oriented (TCP)**
 - **Connectionless (UDP)**
 - **MulticastSocket** class— data can be sent to multiple recipients
- Consider this “Date” server

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```





Remote Procedure Calls

■ Remote Procedure Call (RPC)

- Abstracts procedure calls between processes on networked systems
- Again uses ports for service differentiation
- Stubs – client-side proxy for actual procedure on server
- Client-side stub locates server and marshalls parameters
- Server-side stub receives this message, unpacks marshalled parameters, and performs procedure on server





Remote Procedure Calls (Cont.)

■ On Windows

- Stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**





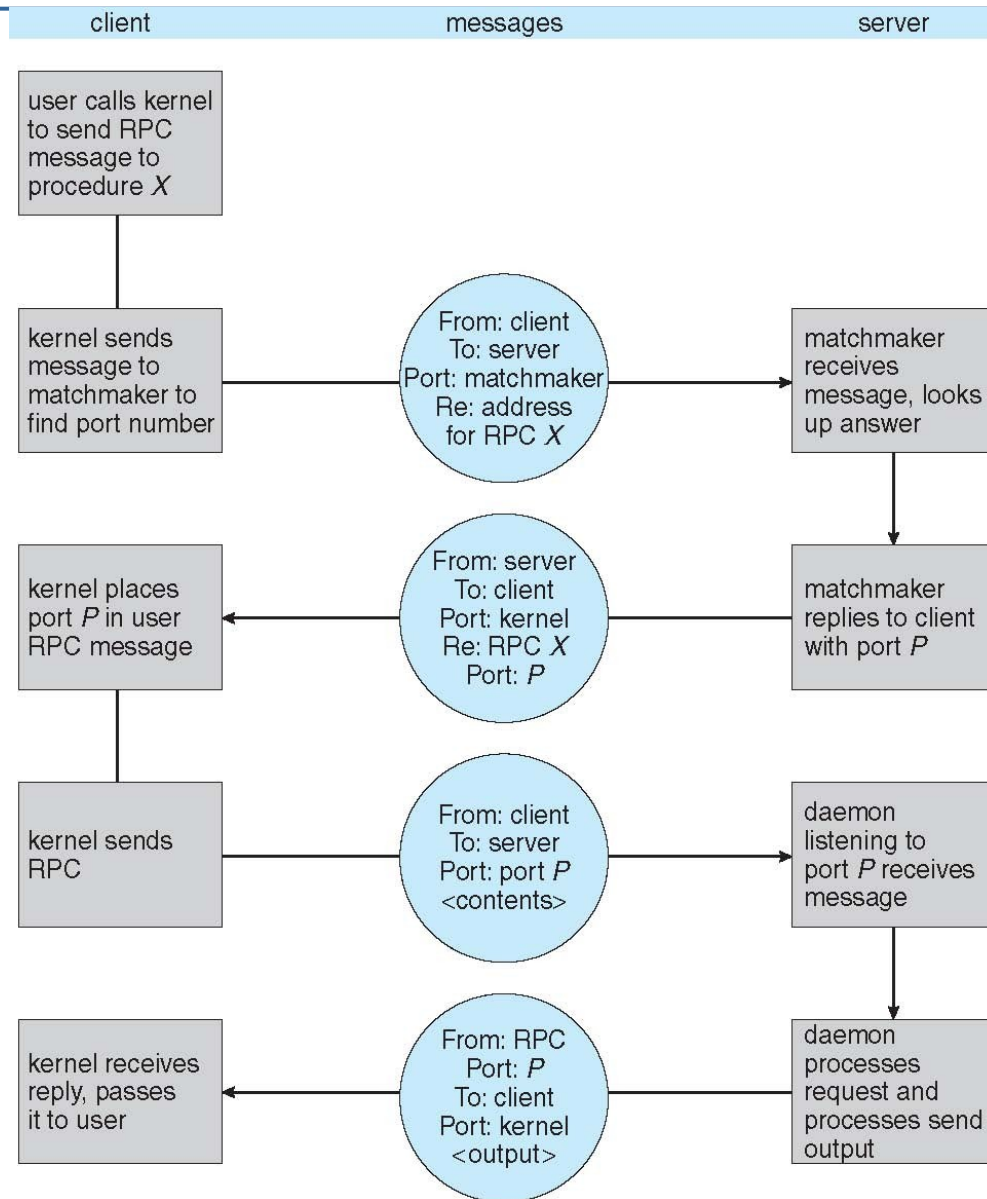
Remote Procedure Calls (Cont.)

- Data Representation Handled via **External Data Representation (XDL)** Format
 - To account for different architectures
 - ▶ **Big-endian** and **little-endian**
- Remote Communication Has more Failure Scenarios than Local
 - Messages can be delivered ***exactly once*** rather than ***at most once***
- OS Typically Provides a Rendezvous (or **matchmaker**) Service to Connect Client and Server





Execution of RPC





Pipes

- Acts as a Conduit allowing Two Processes to Communicate
 - 1st IPC mechanism used in early UNIX
- Issues:
 - Is communication unidirectional or bidirectional?
 - In case of two-way communication, is it half or full-duplex?
 - Must there exist a relationship (i.e., ***parent-child***) between the communicating processes?
 - Can pipes be used over a network?





Pipes (cont.)

■ Ordinary Pipes

- Cannot be accessed from outside process that created it.
- Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.

■ Named Pipes


- Can be accessed without a parent-child relationship





Ordinary Pipes


- Ordinary Pipes allow Communication in Standard Producer-Consumer Style
- Producer Writes to One End (**write-end** of Pipe)
- Consumer Reads from other end (**read-end** of Pipe)





Ordinary Pipes (Cont.)

- Ordinary Pipes are therefore Unidirectional
- Require Parent-Child Relationship between Communicating Processes
- Windows Calls these **Anonymous Pipes**
- See Unix and Windows Code Samples in Textbook





Named Pipes

- Named Pipes More powerful than Ordinary Pipes
- Communication is Bidirectional
- No Parent-Child Relationship is Necessary between Communicating Processes
- Several Processes can Use Named Pipe for Communication
- Provided on both UNIX and Windows Systems



End of Lecture 3

