


Lecture 9: Deadlocks

Hossein Asadi (asadi@sharif.edu)

Rasool Jalili (jalili@sharif.edu)





Lecture 9: Deadlocks

- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock





Lecture Objectives

- To Develop a **Description of Deadlocks**, which Prevent sets of Concurrent Processes from Completing their Tasks
- To Present a number of Different Methods for **Preventing** or **Avoiding** Deadlocks in a Computer System





System Model

■ System Consists of Resources

- To be distributed among a number of competing processes

■ Resource Types R_1, R_2, \dots, R_m


E.g.,: CPU cycles, memory space, I/O devices

■ Each Resource Type R_i has W_i instances

■ Each Process Utilizes a Resource as follows

● Request, Use, Release

- ▶ Accomplished using system calls such as request() and release() device or open() and close() file, OR allocate() and free() memory





System Model (cont.)

■ Physical Resources

- Printers, tape drives, memory space, or CPU cores/cycles

■ Logical Resources

- Files, semaphores, and monitors

■ Example: Consider a system with 3 CD RW drives and 3 running processes

- Each process holds one CD RW drive and now requests another drive → Deadlock

■ Multithreaded Programs:

- Good candidate for deadlock
- Multiple threads compete for shared resources





Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously

- **Mutual Exclusion:** only one process at a time can use a resource
- **Hold and Wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- **No Preemption:** a resource can be released only voluntarily by process holding it, after that process has completed its task





Deadlock Characterization (cont.)

- **Circular Wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .





Resource-Allocation Graph

A set of vertices V and a set of edges E .

■ V is partitioned into two types:

- $P = \{P_1, P_2, \dots, P_n\}$, set consisting of all **processes** in the system
- $R = \{R_1, R_2, \dots, R_m\}$, set consisting of all **resource types** in the system

■ **Request Edge** – directed edge $P_i \rightarrow R_j$


■ **Assignment Edge** – directed edge $R_j \rightarrow P_i$





Resource-Allocation Graph (cont.)


■ Process




■ Resource Type with 4 instances



■ P_i requests instance of R_j




■ P_i is holding an instance of R_j





Example of a Resource Allocation Graph



■ $P = \{P_1, P_2, P_3\}$

■ $R = \{R_1, R_2, R_3, R_4\}$


- R_1 : one instance, R_2 two instances, R_3 one instance, and R_4 three instances

■ $E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_3, \dots, R_3 \rightarrow P_3\}$






Resource Allocation Graph With a Deadlock





Graph With A Cycle But No Deadlock





Basic Facts

- If Graph Contains **no Cycles** \Rightarrow **No Deadlock**
- If Graph Contains a Cycle \Rightarrow
 - Only one instance per resource type \rightarrow **deadlock**
 - Cycle involves only a set of resource types, each of which has only a single instance \rightarrow **deadlock**
 - ▶ In the above two cases, cycle is necessary & sufficient condition for existence of deadlock
 - Several instances per resource type \rightarrow **possibility of deadlock**
 - ▶ In this case, cycle is necessary but not sufficient





Methods for Handling Deadlocks

- Ensure that System will **Never** Enter a Deadlock State
 - Deadlock **prevention**: try to **violate** one of necessary **conditions** for deadlock
 - Deadlock **avoidance**: try to **regulate** how/when requests can be made to acquire resources
 - ▶ More **conservative** approach than deadlock prevention
- Allow System to **enter a Deadlock State** and then **recover**





Methods for Handling Deadlocks (cont.)

- Ignore Problem and Pretend that Deadlocks Never occur in system
 - Used by most OSes, including UNIX
 - Up to application developer to detect and handle deadlocks
- What if Deadlocks are not Resolved?
 - Deterioration of system performance
 - ▶ Eventually need a manual restart
 - Deadlock occur very infrequent → cheaper approach in mainstream applications
 - ▶ Instead of employing prevention, avoidance, or detection and recovery methods





Deadlock Prevention

■ Mutual Exclusion

- Not required for sharable resources
- A process never needs to wait for a sharable resources
- Must hold for non-sharable resources
- Example
 - ▶ Read-only files





Deadlock Prevention (cont.)

- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
 - **Solution 1**: Require process to request and be allocated **all its resources** before it **begins** execution
 - **Solution 2**: Or allow process to request resources only when process **has none** allocated to it
 - **Cons**
 - ▶ Low resource utilization ☹️
 - ▶ Starvation possible ☹️





Deadlock Prevention (cont.)

■ No Preemption


- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then **all resources currently** being held are **released**
- **Preempted resources** are **added** to **list** of resources for which the process is **waiting**
- Process will be **restarted only** when it can **regain** its **old** resources, **as well as new ones**
- This protocol **applicable only** to resources whose state can be easily saved and restored later
 - ▶ CPU registers and memory space: applicable 😊
 - ▶ Printers and tape drives: not (easily) applicable ☹️





Deadlock Prevention (cont.)

- **Circular Wait** – impose a total ordering of all resource types, and require that each process requests resources in an **increasing order** of enumeration
 - A process which holds $R(i)$, can request instance of $R(j)$ if $F(Rj) > F(Ri)$
 - Ensuring order by application developer
 - ▶ Can use **lock-order verifier** (e.g., **witness** in FreeBSD)
- **Example**
 - $F(\text{tape})=1$, $F(\text{disk drive})=5$, and $F(\text{printer})=12$





Deadlock Example

```
/* thread one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /** * Do some work */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);
    pthread_exit(0);
}
/*****/
/* thread two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /** * Do some work */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);
    pthread_exit(0);
}
```






Deadlock Example with Lock Ordering

■ Lock Ordering does not Guarantee Deadlock Prevention if Locks can be acquired Dynamically

- Ordering is broken with **unordered arguments**

```
void transaction(Account from, Account to, double amount)
{
    mutex lock1, lock2;
    lock1 = get_lock(from);
    lock2 = get_lock(to);
    acquire(lock1);
    acquire(lock2);
    withdraw(from, amount);
    deposit(to, amount);
    release(lock2);
    release(lock1);
}
```

- Transactions 1 and 2 execute concurrently.
- Transaction 1 transfers \$25 from account A to account B, &
Transaction 2 transfers \$50 from account B to account A





Deadlock Avoidance

Requires that system has some additional ***a priori*** information available

- Simplest and most useful model requires that each process **declare *maximum number*** of resources of each type that it may need
- Deadlock-avoidance algorithm dynamically examines **resource-allocation state** to ensure that there can never be a **circular-wait** condition
- “Resource-allocation *state*” is defined by number of **available** and **allocated** resources, and **maximum demands** of processes





Safe State

- When a process requests an available resource, system must decide if immediate allocation **leaves** system in a **safe state**
- System is in **safe state** if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ (aka, **safe sequence**) of ALL processes in systems such that for each P_i , resources that P_i can still request can be satisfied by **currently available** resources + **resources held** by all P_j , with $j < i$





Safe State (cont.)

■ That is:

- If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
- When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
- When P_i terminates, P_{i+1} can obtain its needed resources, and so on







Basic Facts

- If a System is in Safe State \Rightarrow No Deadlocks
- If a System is in Unsafe State \Rightarrow Possibility of Deadlock
- Avoidance \Rightarrow Ensure that a System will Never enter an Unsafe State





Safe, Unsafe, Deadlock State





Avoidance Algorithms

- **Single Instance** of a Resource Type
 - Use a resource-allocation graph

- **Multiple Instances** of a Resource Type
 - Use banker's algorithm
 - ▶ Reading assignment





Resource-Allocation Graph Scheme

■ Claim edge

- $P_i \rightarrow R_j$ indicates that process P_i may request resource R_j (represented by a dashed line)

■ Claim edge converts to request edge when a process requests a resource

■ Request edge converted to an assignment edge when resource is allocated to process


■ Resource is Released by a process → Assignment Edge reconverts to a claim edge

■ Resources must be claimed *a priori* in system






Resource-Allocation Graph





Unsafe State In Resource-Allocation Graph



Why not assigning R_2 to P_2 after making sure that P_1 doesn't need R_2 in the near future?





Resource-Allocation Graph Algorithm

- Suppose that P_i requests a R_j
- Request can be Granted only if Converting Request Edge to an Assignment Edge does not Result in Formation of a Cycle in Resource Allocation Graph
- Algorithm Complexity $O(n^2)$
- Low Resource Utilization
 - Resources might be available and not be allocated to processes





Deadlock Detection

- Allow System to Enter Deadlock State
- Detection Algorithm
 - Single instance of each resource type
 - Multiple instances of a resource type
 - Reading assignment
- Recovery Scheme






Single Instance of Each Resource Type

- Maintain **wait-for** graph
 - Nodes are processes
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j
- **Periodically** invoke an Algorithm that Searches for a **Cycle** in graph
 - If there is a cycle \rightarrow there exists a deadlock
- An Algorithm to **Detect** a Cycle in a Graph Requires an **Order** of n^2 operations
 - Where n is number of vertices in graph






Resource-Allocation Graph and Wait-for Graph



(a)

**Resource-Allocation
Graph**



(b)

**Corresponding
wait-for Graph**





Detection-Algorithm Usage

- **When, and How often, to Invoke depends on:**
 - How often a deadlock is likely to occur?
 - How many processes will need to be rolled back?
 - ▶ One for each disjoint cycle
- **If Detection Algorithm is Invoked Arbitrarily, there may be many cycles in resource graph and so we would not be able to tell which of many deadlocked processes “caused” deadlock**





Recovery from Deadlock: **Process Termination**

■ **Abort** all Deadlocked Processes

- Significant expense
- Most processes need to be restarted for re-computation

■ Abort one process at a time until Deadlock Cycle is Eliminated

- After each process is aborted, a deadlock-detection algorithm must be invoked
 - ▶ To see if any process is still deadlocked





Recovery from Deadlock: **Process Termination**

■ Issues with Aborting a Process

- Updating a file
- Printing data

■ In which Order should we Choose to Abort?

1. Priority of process
2. How long process has computed, and how much longer to completion
3. Resources process has used
4. Resources process needs to complete
5. How many processes will need to be terminated
6. Is process interactive or batch?





Recovery from Deadlock: **Resource Preemption**

- **Selecting a Victim** – minimize cost
- **Rollback** – return to some safe state, restart process for that state
- **Starvation** – same process may always be picked as victim, include number of rollback in cost factor





Reading Assignment

- Banker's Algorithm
- Deadlock Detection
 - Multiple instances of a resource type





Banker's Algorithm

- Multiple Instances
- Each Process must a Priori Claim Max Use
- When a Process Requests a Resource it may have to Wait
- When a Process Gets all its Resources it must Return them in a Finite Amount of Time





Data Structures for Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- **Available:** Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available
- **Max:** $n \times m$ matrix. If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j
- **Allocation:** $n \times m$ matrix. If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j
- **Need:** $n \times m$ matrix. If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task
 $Need[i, j] = Max[i, j] - Allocation[i, j]$





Safety Algorithm

1. Let **Work** and **Finish** be vectors of length m and n , respectively. Initialize:

Work = **Available**

Finish [i] = **false** for $i = 0, 1, \dots, n-1$

2. Find an i such that both:

(a) **Finish** [i] = **false**

(b) **Need** _{i} ≤ **Work**

If no such i exists, go to step 4

3. **Work** = **Work** + **Allocation** _{i}

Finish [i] = **true**

go to step 2

4. If **Finish** [i] == **true** for all i , then the system is in a safe state





Resource-Request Algorithm for Process P_i

$Request_i$ = request vector for process P_i .

If **$Request_i[j] = k \rightarrow P_i$** wants k instances of R_j

1. If **$Request_i \leq Need_i$** , go to step 2. Otherwise, raise error condition (as process has exceeded its maximum claim)
2. If **$Request_i \leq Available$** , go to step 3. Otherwise P_i must wait, since resources are not available
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$Available = Available - Request_i;$

$Allocation_i = Allocation_i + Request_i;$

$Need_i = Need_i - Request_i;$

- If safe \Rightarrow the resources are allocated to P_i
- If unsafe $\Rightarrow P_i$ must wait, and old resource-allocation state is restored





Example of Banker's Algorithm

■ 5 processes P_0 through P_4 ;

3 resource types:

A (10 instances), B (5 instances), and C (7 instances)

■ Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
P_0	0 1 0	7 5 3	3 3 2
P_1	2 0 0	3 2 2	
P_2	3 0 2	9 0 2	
P_3	2 1 1	2 2 2	
P_4	0 0 2	4 3 3	





Example (cont.)

- Content of ***Need*** is defined to be ***Max – Allocation***

	<u><i>Need</i></u>		
	<i>A</i>	<i>B</i>	<i>C</i>
<i>P</i> ₀	7	4	3
<i>P</i> ₁	1	2	2
<i>P</i> ₂	6	0	0
<i>P</i> ₃	0	1	1
<i>P</i> ₄	4	3	1

- System is in a **Safe** State

- Since sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria





Example: P_1 Request (1,0,2)

- Check that Request \leq Available (i.e., $(1,0,2) \leq (3,3,2) \Rightarrow$ true

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	4	3	2	3	0
P_1	3	0	2	0	2	0			
P_2	3	0	2	6	0	0			
P_3	2	1	1	0	1	1			
P_4	0	0	2	4	3	1			

- Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement
- Can request for (3,3,0) by P_4 be granted?
- Can request for (0,2,0) by P_0 be granted?





Several Instances of a Resource Type

- **Available:** A vector of length m indicates number of available resources of each type
- **Allocation:** An $n \times m$ matrix defines number of resources of each type currently allocated to each process
- **Request:** An $n \times m$ matrix indicates current request of each process
 - If $Request[i][j] = k$, \rightarrow process P_i is requesting k more instances of resource type R_j





Detection Algorithm

1. Let ***Work*** and ***Finish*** be vectors of length ***m*** and ***n***, respectively Initialize:

(a) ***Work = Available***

(b) For $i = 1, 2, \dots, n$, if ***Allocation_i $\neq 0$*** , then
Finish[i] = false; otherwise, ***Finish[i] = true***

2. Find an index ***i*** such that both:

(a) ***Finish[i] == false***

(b) ***Request_i \leq Work***

If no such ***i*** exists, go to step 4





Detection Algorithm (cont.)

3. **$Work = Work + Allocation_i$**
 $Finish[i] = true$
go to step 2

4. If **$Finish[i] == false$** , for some i , $1 \leq i \leq n$,
then the system is in deadlock state.
Moreover, if **$Finish[i] == false$** , then P_i is
deadlocked

**Algorithm requires an order of $O(m \times n^2)$ operations
to detect whether the system is in deadlocked state**





Example of Detection Algorithm

- Five Processes P_0 through P_4 ; three Resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time T_0 :

	<u>Allocation</u>			<u>Request</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	0	0	0	0	0	0
P_1				2	0	0	2	0	2
P_2	3	0	3	0	0	0			
P_3	2	1	1	1	0	0			
P_4	0	0	2	0	0	2			

- Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in ***Finish[i] = true*** for all i






Example (cont.)

- P_2 requests an additional instance of type C

	<u>Request</u>		
	A	B	C
P_0	0	0	0
P_1	2	0	2
P_2	0	0	1
P_3	1	0	0
P_4	0	0	2

- State of system?

- Can reclaim resources held by process P_0 , but insufficient resources to fulfill other processes; requests
- Deadlock exists, consisting of processes P_1 , P_2 , P_3 , and P_4



End of Lecture 9

