



Lecture 2:

Operating-System Structures

Hossein Asadi (asadi@sharif.edu)

Rasool Jalili (jalili@sharif.edu)





Lecture 2: Operating-System Structures

- Operating System Services
- User Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating System Design and Implementation
- Operating System Structure
- Virtual Machines





Objectives

- To Describe **Services** an OS Provides to Users, Processes, and Other Systems
- To Discuss **Various Ways** (or Different Goals) of **Structuring** an OS





Operating System Services

- OSs Provide an Environment for Execution of **Programs** and **Services** to Programs and Users
- Set of OS **Services** Provides Functions that are Helpful to Users for **Convenience** of Programmers
 - User interface
 - Program execution
 - I/O operations
 - File-system manipulation
 - Communications
 - Error detection





Operating System Services (cont.)

■ User Interface

- Almost all OSes have a user interface (**UI**)
 - ▶ **Command-Line (CLI)**
 - ▶ **Graphics User Interface (GUI)**: window/mouse
 - ▶ **Batch interface**: input commands from file

■ Program Execution

- System must be able to **load** a **program** into memory and to **run** that program, **end execution**, either normally or abnormally (indicating error)





Operating System Services (cont.)

■ I/O Operations

- A running program may require I/O, which may involve a file or an I/O device
- Users cannot directly control I/O devices
 - ▶ For sake of protection and efficiency

■ File-System Manipulation

- Programs need to read and write files/dir
- Create and delete files/dir
- Search files/dir
- List file information
- Permission management





Operating System Services (cont.)

■ Communications

- Processes may exchange information
 - ▶ On same computer
 - ▶ Or between computers over a network
- Communications may be via **shared memory** or through **message passing** (packets moved by OS)





Operating System Services (cont.)

■ Error Detection

- OS needs to be constantly aware of possible errors
- May occur in CPU and memory HW, in I/O devices, in user program
- For each type of error, OS should take appropriate action to ensure correct and consistent computing
- Debugging facilities can greatly enhance user's and programmer's abilities to efficiently use system





Operating System Services (cont.)

- Another Set of OS Functions Exist for Ensuring **Efficient Operation** of System itself via **Resource Sharing**
 - Resource allocation
 - Accounting
 - Protection and security





Operating System Services (cont.)

■ Resource Allocation

- When multiple users/jobs running concurrently, resources must be allocated to each of them
- Many types of resources:
 - ▶ CPU cycles
 - ▶ Main memory
 - ▶ File storage
 - ▶ I/O devices





Operating System Services (cont.)

■ Accounting

- To keep track of which users use how much and what kinds of computer resources
- Accounting info can be used for billing or just statistics
- Such statistics info helps researchers/admins to reconfigure system to improve computing efficiency and services





Operating System Services (cont.)

■ Protection and Security

- Owners of information stored in a multiuser or networked computer system may want to control use of that information, **concurrent processes** should **not interfere** with each other

■ Protection

- Involves ensuring that all accesses to system resources are controlled

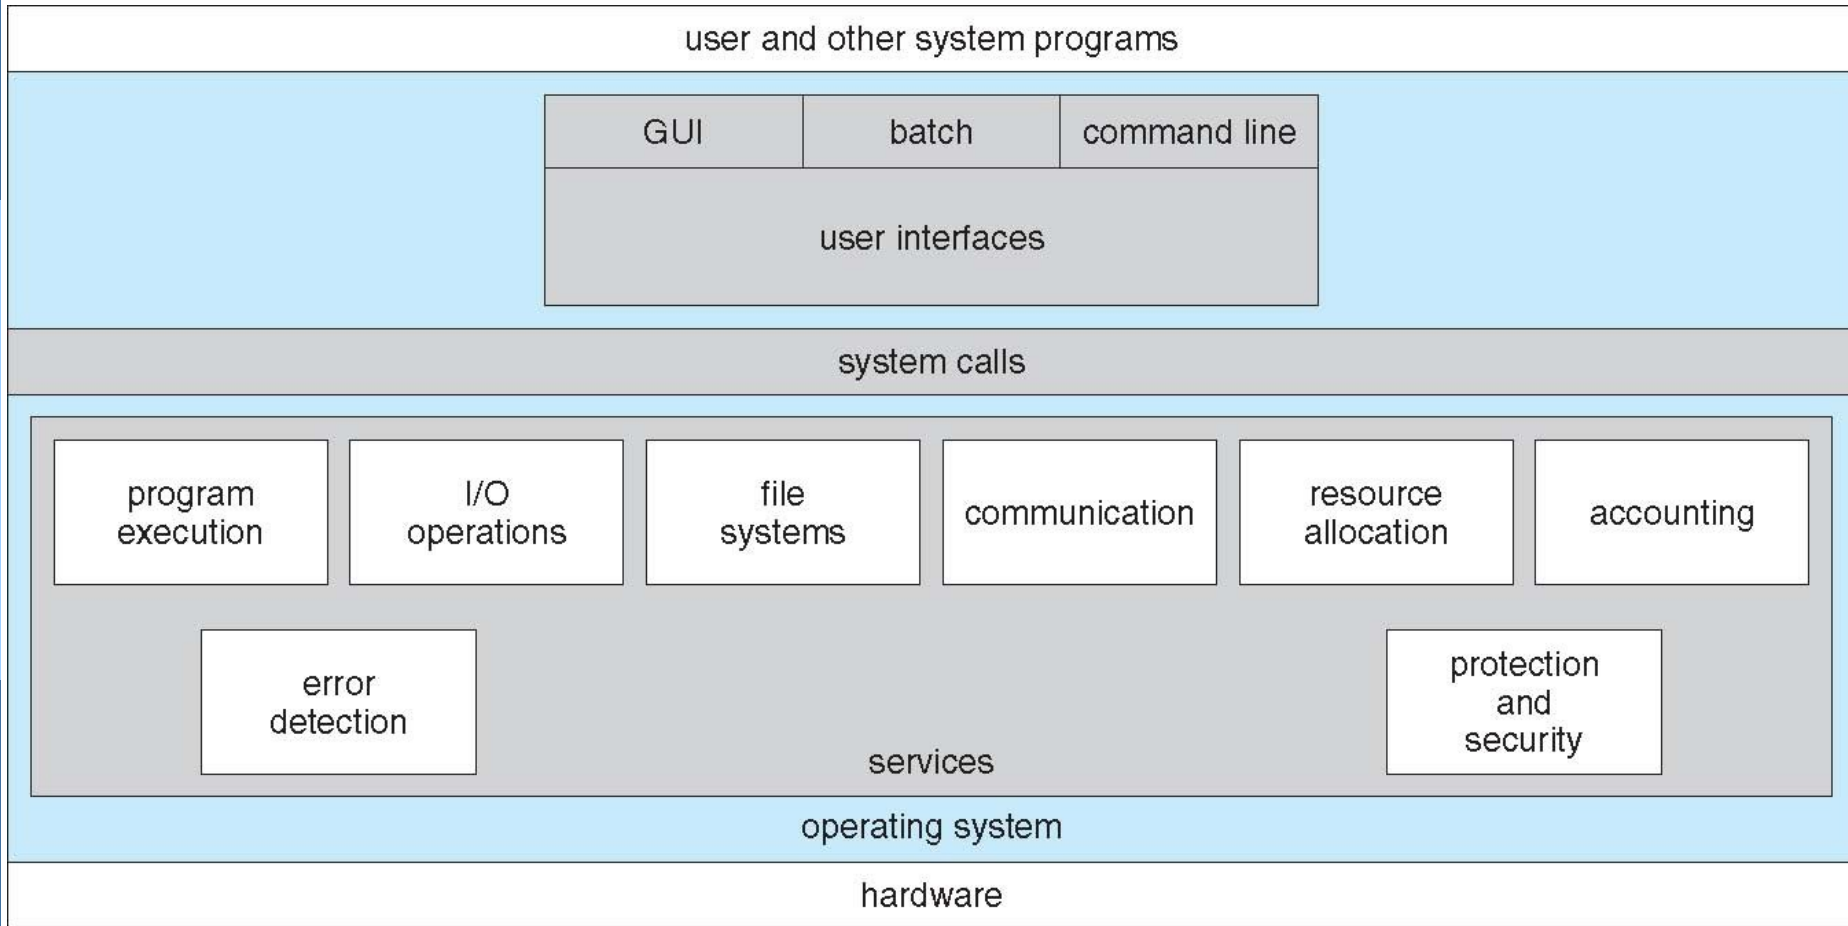
■ Security

- Security of system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts





OS Services: an Overview





User OS Interface - CLI

■ CLI or Command Interpreter

- Allows direct command entry
- Sometimes implemented in **kernel**, sometimes by **systems** program
- Sometimes multiple flavors implemented: **shells**
- Primarily fetches a command from user and executes it





User OS Interface – CLI (cont.)


■ Two Approaches in Implementing Shells

- Commands built-in
 - ▶ For each command, there is a relevant code in Shell
- Just names of programs
 - ▶ Used in UNIX
 - ▶ Does not understand input command
 - Just searches for a file/program with given name of command
 - ▶ E.g.: `rm file.txt`
 - Would search for a file called “rm” → loads file into memory → execute it with the parameter “file.txt”
 - ▶ Adding new features doesn't require shell modification





Bourne Shell Command Interpreter



```
PBG-Mac-Pro:~ pbg$ w
15:24 up 56 mins, 2 users, load averages: 1.51 1.53 1.65
```

```
USER      TTY      FROM            LOGIN@   IDLE   WHAT
pbg       console -                14:34    50    -
pbg       s000    -                15:05    -    w
```

```
PBG-Mac-Pro:~ pbg$ iostat 5

disk0      disk1      disk10      cpu      load average
KB/t tps  MB/s  KB/t tps  MB/s  KB/t tps  MB/s  us sy id  1m  5m  15m
33.75 343 11.30  64.31 14  0.88  39.67 0  0.02 11 5 84  1.51 1.53 1.65
5.27 320  1.65   0.00 0  0.00   0.00 0  0.00  4 2 94  1.39 1.51 1.65
4.28 329  1.37   0.00 0  0.00   0.00 0  0.00  5 3 92  1.44 1.51 1.65
```

```
AC
PBG-Mac-Pro:~ pbg$ ls
Applications
Applications (Parallels)
Desktop
Documents
Downloads
Dropbox
Library
Movies
```

```
Music
Pando Packages
Pictures
Public
Sites
Thumbs.db
Virtual Machines
Volumes
```

```
PBG-Mac-Pro:~ pbg$ pwd
/Users/pbg
PBG-Mac-Pro:~ pbg$ ping 192.168.1.1
PING 192.168.1.1 (192.168.1.1): 56 data bytes
64 bytes from 192.168.1.1: icmp_seq=0 ttl=64 time=2.257 ms
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=1.262 ms
```

```
AC
--- 192.168.1.1 ping statistics ---
2 packets transmitted, 2 packets received, 0.0% packet loss
round-trip min/avg/max/stddev = 1.262/1.760/2.257/0.498 ms
PBG-Mac-Pro:~ pbg$
```

```
WebEx
config.log @hossein-MacBookAir: ~
getsmartdata.txt
imp
top - 11:13 up 11:20, 3 users,
Tasks: log7 total, 2 running, 155 sleeping, 0 stopped, 0 zombie
Cpu(s) 9.0%us, 9.0%sy, 0.0%ni, 70.2%id, 0.9%wa, 0.0%hi, 0.2%st
Mem: 3928500k total, 3647724k used, 138372k free, 174356k buffers
Swap: 3928500k total, 1752k used, 3926748k free, 1024320k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1788	hossein	20	0	1277m	105m	88m	S	16	2.8	5:10.38	compiz
136	root	20	0	237m	73m	17m	S	11	2.0	8:44.68	Xorg
5974	hossein	20	0	2723m	1.0g	1.0g	S	6	29.1	6:17.95	VirtualBox
2005	hossein	20	0	409m	728k	628k	S	1	0.2	0:01.52	zeitgeist-datah
50	hossein	20	0	28300	404k	608k	S	1	0.1	0:08.40	dbus-daemon
1829	hossein	20	0	426m	702k	460k	S	1	0.2	0:07.86	pulseaudio
1846	hossein	20	0	390m	11m	636k	S	1	0.3	0:04.61	bamfdaemon
1888	hossein	20	0	541m	34m	11m	S	1	0.9	0:23.17	unity-panel-ser
3419	hossein	20	0	960m	192m	47m	S	1	5.2	5:13.66	thunderbird
5655	root	20	0	0	0	0	R	1	0.0	0:11.70	kworker/1:2
6193	hossein	20	0	521m	17m	11m	S	1	0.5	0:00.62	gnome-terminal
3	root	20	0	0	0	0	S	0	0.0	0:32.38	ksoftirqd/0
1711	hossein	20	0	391m	10m	7492k	S	0	0.3	0:00.53	gnome-session
1833	hossein	20	0	140m	2512	2016k	S	0	0.1	0:01.04	gvfs-afc-volume
1879	hossein	20	0	328m	13m	9428k	S	0	0.4	0:04.23	gtk-window-deco
1951	hossein	20	0	158m	5388	4284k	S	0	0.1	0:00.34	ubuntu-geoip-pr
2006	hossein	20	0	354m	123m	7432k	S	0	3.4	0:08.07	zeitgeist-fts



User Operating System Interface - GUI

- User-Friendly **Desktop** Metaphor Interface
 - Usually mouse, keyboard, and monitor
 - **Icons** represent files, programs, actions, etc
 - Various **mouse buttons** over objects in interface cause various actions (provide information, options, execute function, open directory (known as a **folder**))
 - Invented at Xerox PARC





User Operating System Interface – GUI (cont.)

- Many Systems now Include both CLI and GUI Interfaces
 - Microsoft **Windows** is GUI with CLI “command” shell
 - Apple **Mac OS X** is “Aqua” GUI interface with UNIX kernel underneath and shells available
 - Unix and Linux have CLI with optional GUI interfaces (CDE, KDE, GNOME)





Touchscreen Interfaces

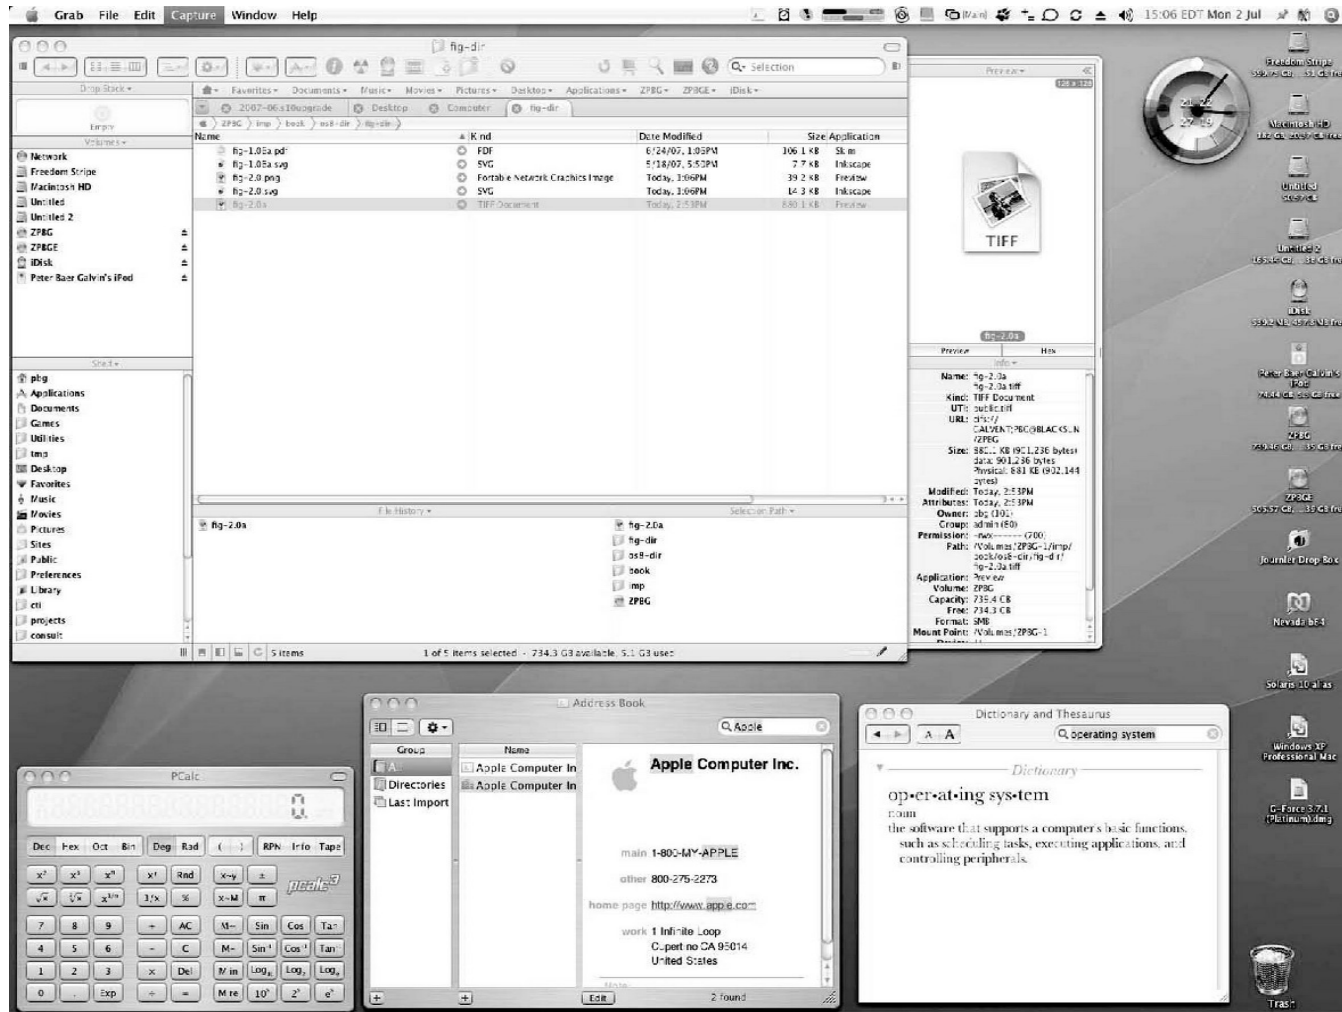
■ Touchscreen Devices Require New Interfaces

- Mouse not possible or not desired
- Actions and selection based on gestures
- Virtual keyboard for text entry
- Voice commands





Mac OS X GUI





System Calls

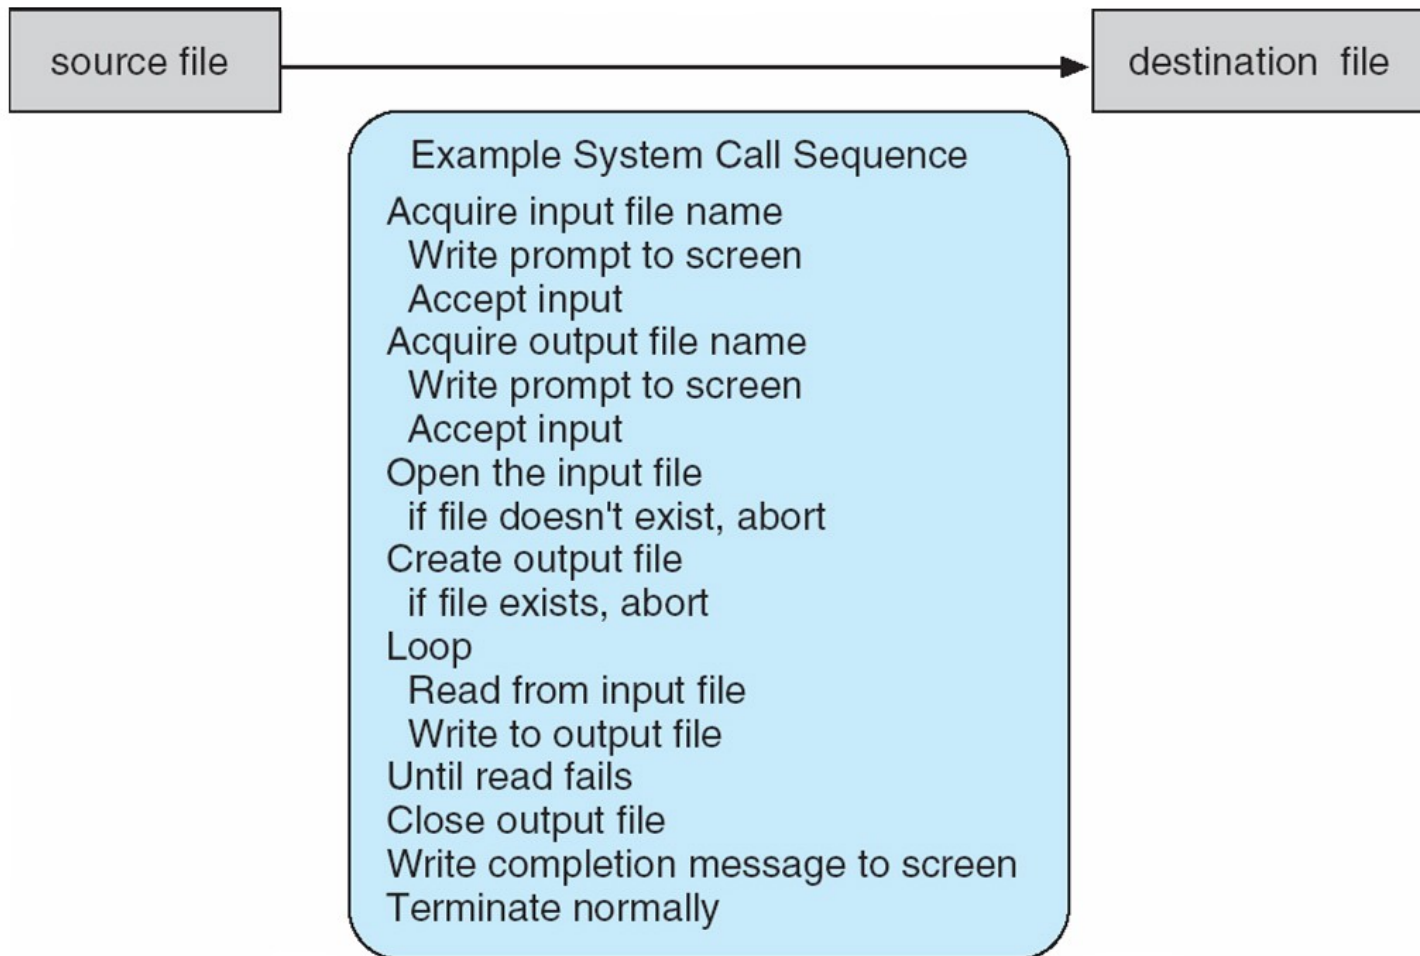
- Programming **Interface to Services** Provided by OS
- Typically Written in a High-Level Language (C/C++)
 - Some **low-level** and **most-frequent** tasks written in **Assembly**
- Examples
 - A system call to control a process
 - ▶ CreateProcess()
 - ▶ ExitProcess()
 - A system call to manipulate a file
 - ▶ ReadFile()
 - ▶ CreateFile()





Example of System Calls

■ System Call Sequence to Copy Contents of one File to another File





System Calls (cont.)

- Mostly Accessed by Programs via **API**
 - Rather than direct system call use
- **Application Programming Interface (API)**
 - Specifies a set of functions that are available to an application programmer
 - ▶ **Parameters** passed to each function
 - ▶ **Return values** programmer can expect
- Three Most Common APIs
 - Win32 API for Windows
 - POSIX API for POSIX-based systems
 - ▶ Including virtually all versions of UNIX, Linux, and Mac OS X
 - Java API for Java Virtual Machine (JVM)





System Calls (cont.)

■ APIs Invoke Actual System Calls

- On behalf of application programmer

■ Examples:

- Win32 function **CreateProcess()**
 - ▶ Calls **NTCreateProcess()** in Windows kernel

■ Why to Use APIs?

- Portability
 - ▶ Programmer can expect his/her program to be compiled on any system that supports the same API
 - ▶ Some system calls are very detailed and hard to work
- More coarse-grained component in programming

■ Mostly a One-to-One Mapping

- Between APIs and system calls





Example of Standard API

EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the `man` page by invoking the command

```
man read
```

on the command line. A description of this API appears below:

```
#include <unistd.h>

ssize_t  read(int fd, void *buf, size_t count)
```

<code>ssize_t</code>	<code>read</code>	<code>(int fd, void *buf, size_t count)</code>
return value	function name	parameters

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns `-1`.





System Call Implementation

■ System-Call Interface

- Invokes intended system call in OS Kernel
- Then returns status of system call and any return values


■ Caller Needs Know Nothing about how system call is Implemented

- Just needs to obey API and understand what OS will do as a result call
- Most details of OS interface hidden from programmer by API





API – System Call – OS Relationship





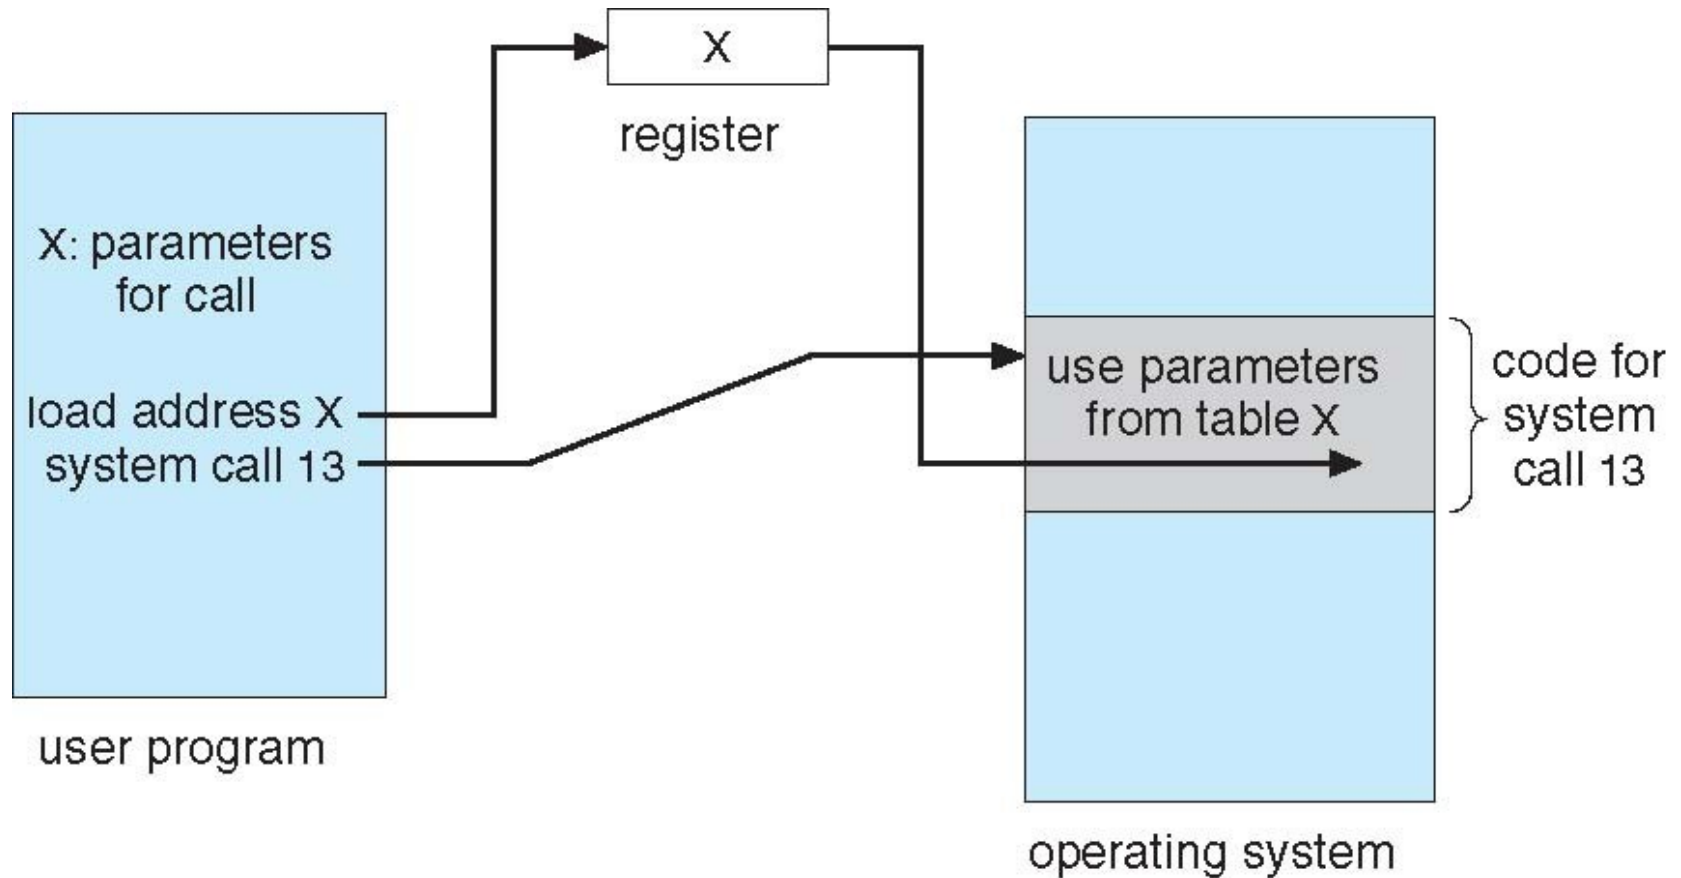
System Call Parameter Passing

- Often, More Information Required than Simply Identity of Desired System Call
 - Exact type and amount of information vary according to OS and call
- Three Methods Used to Pass Parameters to OS
 - **Simplest**: pass parameters in **registers**
 - ▶ In some cases, may be more parameters than registers
 - **Parameters stored in a block, or table, in memory, and address of block passed as a parameter in a register**
 - ▶ This approach taken by Linux and Solaris
 - Parameters placed, or **pushed**, onto **stack** by program and **popped** off the stack by OS
 - Block and stack methods do not limit number or length of parameters being passed





Parameter Passing via Table





Types of System Calls

■ Process Control

- End, abort
- Create process, terminate process, e.g., fork()
- Load, execute
- Get process attributes, set process attributes: e.g., getpid
- Wait for time: e.g. wait()
- Wait event, signal event
- Allocate and free memory
- Dump memory if error
- **Debugger** for determining **bugs**, **single step** execution
- **Locks** for managing access to shared data between processes or **release lock**





Example: MS-DOS

- Single-Tasking
- Shell Invoked when System Booted
- Simple Method to Run Program


- No process created

- Single Memory Space
- Loads Program into memory, overwriting all but kernel
- Program exit → shell reloaded



(a)

At system startup



(b)

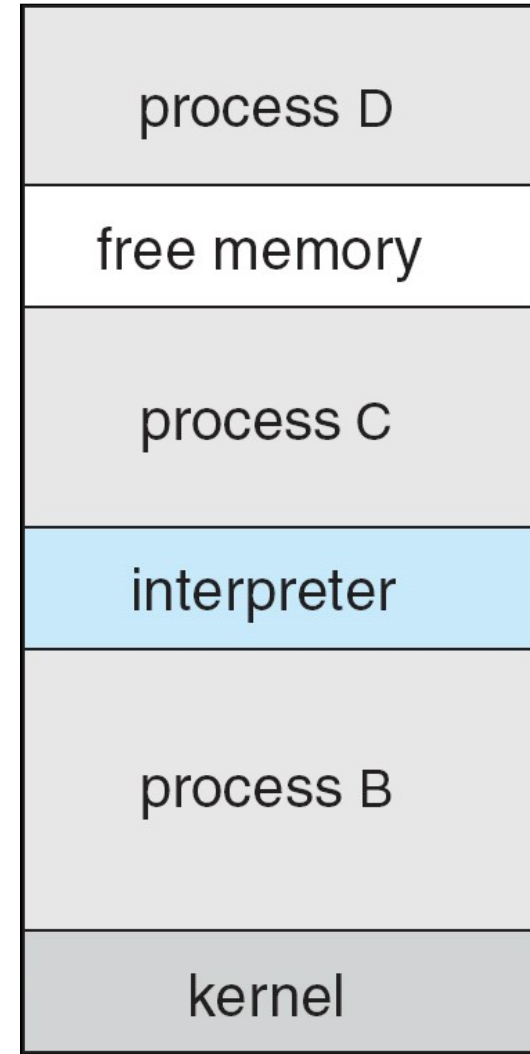
running a program





Example: FreeBSD

- Unix Variant
- Multi-Tasking
- Shell Executes **fork()** System Call to Create Process
 - Executes **exec()** to load program into process
 - Shell waits for process to terminate or continues with user commands
- Process Exits with:
 - (code = 0): no error
 - (code > 0): error code





Sample Code to Call System Call

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main() {
    fork(); //make a child process of same type
    sleep(2);
    printf("Fork testing code\n");
    return 0;
}
```





Types of System Calls (cont.)

■ File Management

- Create file, delete file
 - ▶ `create()`
- Open, close file
 - ▶ `open()`, `close()`
- Read, write, reposition
- Get and set file attributes
 - ▶ File attributes: file name, file type, protection codes, accounting info, etc.
- Copy or move a file
 - ▶ `rename()`, ...





Types of System Calls (cont.)

- Device Management
- Various Resources Controlled by OS can be Thought as **Devices**
 - Physical devices: main memory, disk drives
 - Virtual devices: files
 - Many OSes such as UNIX do not distinguish between physical and virtual devices
 - ▶ → Use a set of similar system calls





Types of System Calls (cont.)

■ Device Management

- Request device, release device
- Read, write, reposition
- Get device attributes, set device attributes
- Logically attach or detach devices





Types of System Calls (cont.)

■ Communications

- Create, delete communication connection
- Send, receive messages if **message passing model** to **host name** or **process name**
 - ▶ From **client** to **server**
 - ▶ **get hostid(), get processid()**
 - ▶ **open() & accept()** to establish a connection
 - ▶ **close()** to terminate the communication
 - ▶ Message passing more suitable for small requests
- **Shared-memory model**
 - ▶ Create and gain access to memory regions
 - ▶ Form of data in shared memory: under process control
 - ▶ Maximum speed and convenience of communication





Types of System Calls (Cont.)

■ Information Maintenance

- Get time or date, set time or date
- Get system data, set system data
- Number of current users
- OS version
- Get and set process, file, or device attributes
- Debugging info
 - ▶ E.g., trace program (where lists each system call as it is executed)
 - ▶ Single step run provided by microprocessors





Types of System Calls (Cont.)

■ Protection

- Control access to resources
- Get and set permissions
 - ▶ `get permission()`
 - ▶ `set permission()`
- Allow and deny user access





Examples of Windows and Unix System Calls

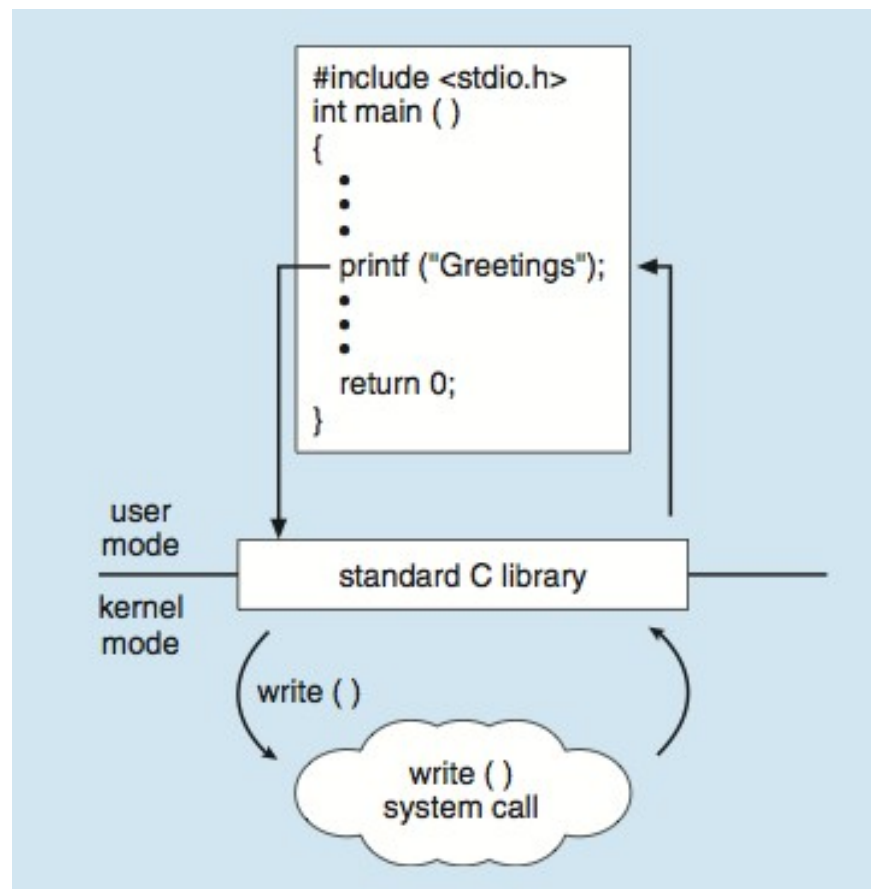
	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()





Standard C Library Example

- C Program Invoking printf() Library Call, which Calls write() System Call





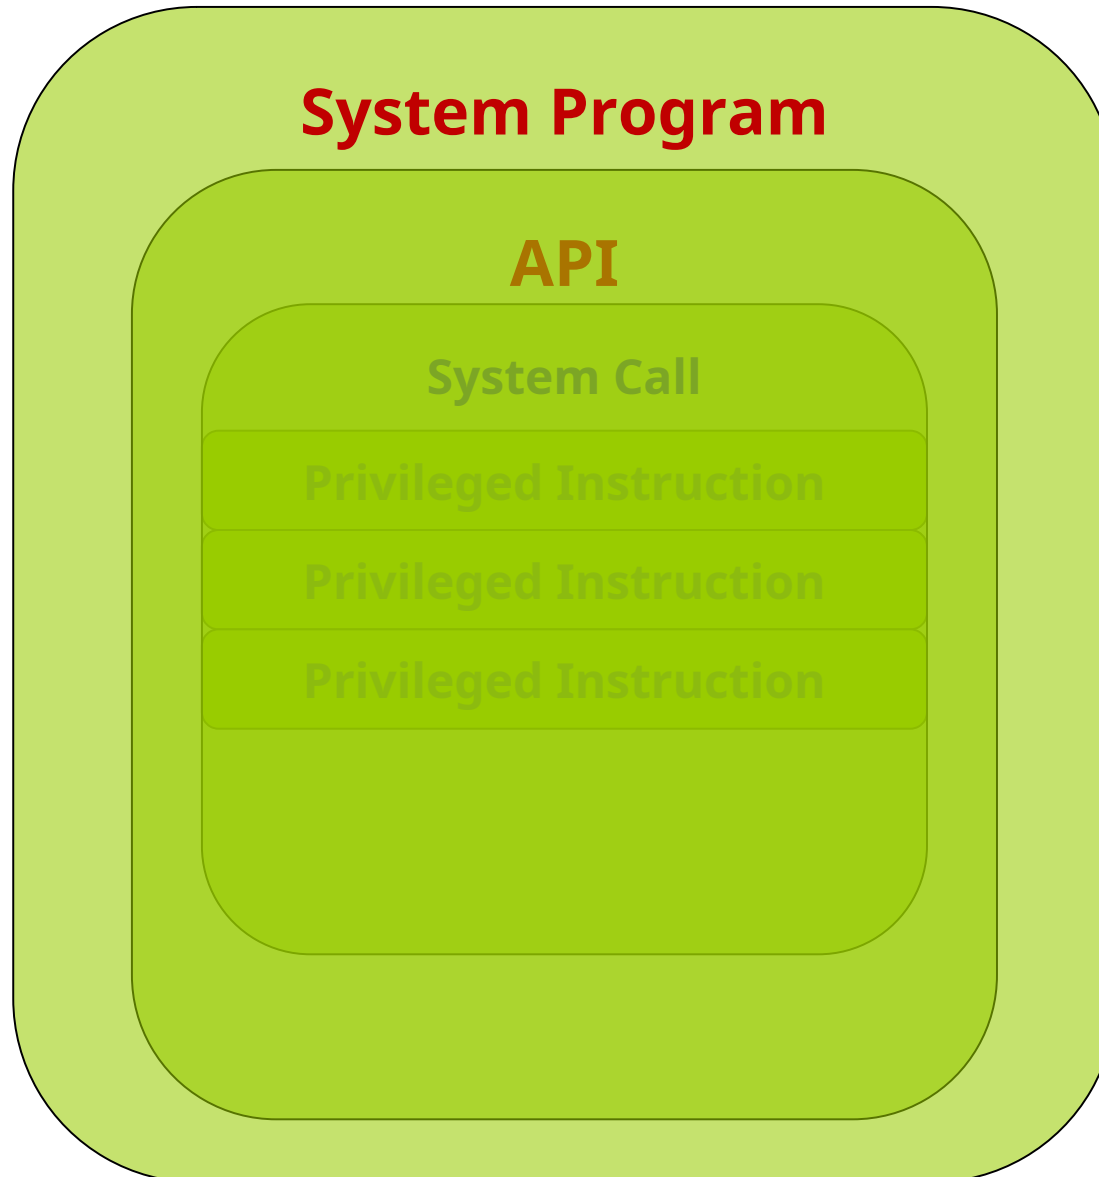
System Programs

- Provide a Convenient Environment for Program Development and Execution
 - Aka, **system utilities**
- Most Users' View of OS defined by System Programs, **not Actual System Calls**
- Some are Simply User Interfaces to System Calls
 - Others are considerably more complex





Privileged Instruction to System Program





System Programs (cont.)

- Can be Divided into:
 - File manipulation
 - Status information
 - File modification
 - Programming language support
 - Program loading and execution
 - Communications
 - Background services
 - Application programs





System Programs (cont.)

■ File Management

- Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories

■ Status Information

- Some ask system for info - date, time, amount of available memory, disk space, number of users
- Others provide detailed performance, logging, and debugging information
- Typically, these programs format and print output to terminal or other output devices
- Some systems implement a **registry** - used to store and retrieve configuration information





System Programs (Cont.)

■ File Modification

- Text editors to create and modify files
- Special commands to search contents of files or perform transformations of the text

■ Programming-Language Support

- Compilers
- Assemblers
- Debuggers and interpreters
- Such system programs provided for C, C++, Java, Visual Basic, Perl, etc.





System Programs (Cont.)

■ Program Loading and Execution

- Absolute loaders, relocatable loaders, linkage editors, and overlay-loaders, debugging systems for higher-level and machine language

■ Communications

- Provide mechanism for creating virtual connections among processes, users, and computer systems
- Allow users to **send messages** to one another's screens, **browse web pages**, send **electronic-mail messages**, **log in remotely**, **transfer files** from one machine to another





System Programs (Cont.)

■ Background Services

- Launch at boot time
 - ▶ Some for system startup, then terminate
 - ▶ Some from system boot to shutdown
- Provide facilities like disk checking, process scheduling, error logging, printing
- Known as **services**, **subsystems**, **daemons**

■ Application Programs (such as web browsers, spreadsheets, database systems, games)

- Don't pertain to system
- Run by users
- Not typically considered part of OS
- Launched by command line, mouse click, finger poke






OS Design and Implementation

■ Design and Implementation of OS

- Internal structure of different OSs can vary widely
- Start design by **defining goals and specifications**
 - ▶ Throughput, energy efficiency, response time, real-time, user-friendly
- Also affected by:
 - ▶ Choice of HW
 - ▶ Type of system: single user, multiuser, real-time, distributed, general purpose
 - ▶ E.g. 1: VxWorks: real-time OS for embedded systems
 - ▶ E.g. 2: MVS: a large multiuser, multi-access OS (IBM)





OS Design and Implementation (cont.)

■ Important Q: **User** goals and **System** goals

- User goals – OS should be convenient to use, easy to learn, reliable, safe, and fast
- System goals – OS should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient





OS Design and Implementation (Cont.)

- Important Principle to Separate: **Policy** vs. **Mechanisms**
 - **Policies**: decide what will be done
 - **Mechanisms** determine how to do sth
- Example: Timer
 - Timer construct: a mechanism used for ensuring CPU protection
 - Policies
 - How long a process can acquire a CPU?
 - Timer is given to user processes only for 10ms
- **Separation** of Policy from Mechanism
 - Allows maximum flexibility if policy decisions are to be changed later





Implementation

■ Much Variation

- Early OSes in **assembly** language
- Then system programming languages like Algol, PL/1
- **Now C, C++**

■ Usually a **Mix** of Languages

- Lowest levels in **assembly** while **main body** in **C**
- Systems programs in C, C++, scripting languages like PERL, Python, shell scripts





Implementation (cont.)

■ High-Level Languages

- Faster code development 😊
- Compact
- Easy to understand and debug
- Improvements in compiler further optimizes code
- Easier to **Port** to other HW
- **Slower** 😞
- **Increased storage/memory** 😞

■ Assembly Languages

- Faster 😊
- **Only can be run on target ISA** 😞





Implementation (cont.)

■ Hybrid Solution

- OS is large
- But small fraction of code is performance bottleneck

■ Critical Sections of OS

- CPU scheduler → Assembly
- Memory manager → Assembly
- Other Parts → C/C++





Operating System Structure

- General-Purpose OS is Very Large Program
- Various Ways to Structure Ones
 - Simple structure – MS-DOS
 - More complex – UNIX
 - Layered – an abstraction
 - Microkernel – Mach





Simple Structure: MS-DOS

■ History

- Originally designed and implemented by very few people

■ Written to Provide Most Functionality in Least space

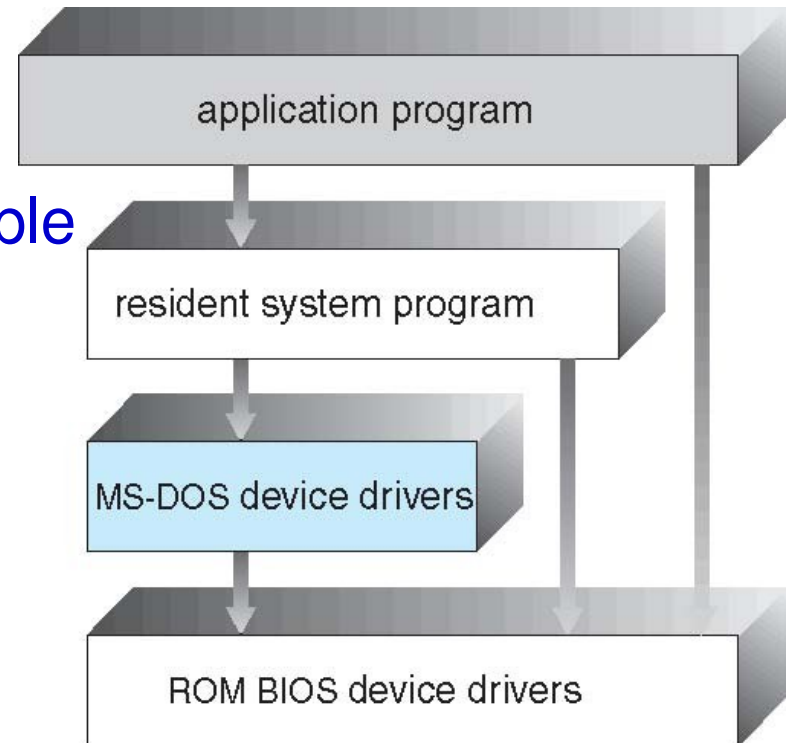
■ Not Well-Defined Structure

- Not divided into modules

■ Interfaces and Levels of Functionality Not Well Separated

- Although it has some structure

■ Very Vulnerable to Errant or Malicious Programs





Non-Simple Structure – UNIX

■ Limited by HW Functionality

- Original UNIX OS had **limited structuring**

■ Consists of Two Separable Parts

- Systems programs
- Kernel

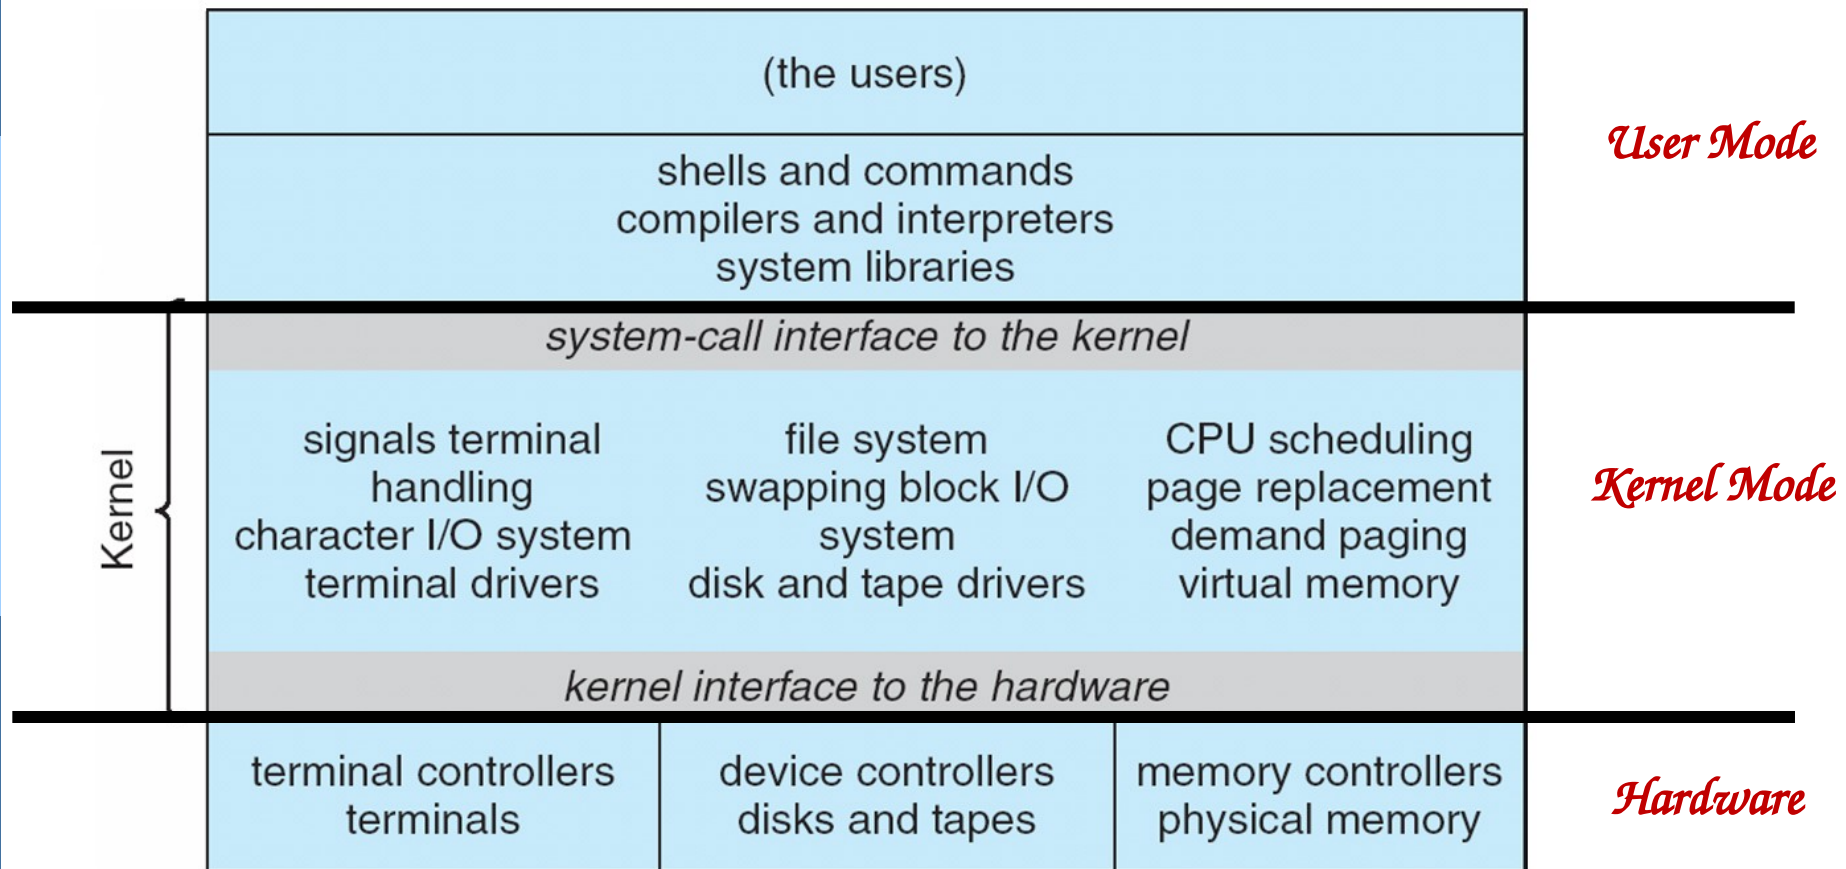
- ▶ Consists of everything **below system-call interface** and **above physical HW**
- ▶ Provides file system, CPU scheduling, memory management, and other OS functions
- ▶ A large number of functions **for one level**





Traditional UNIX System Structure

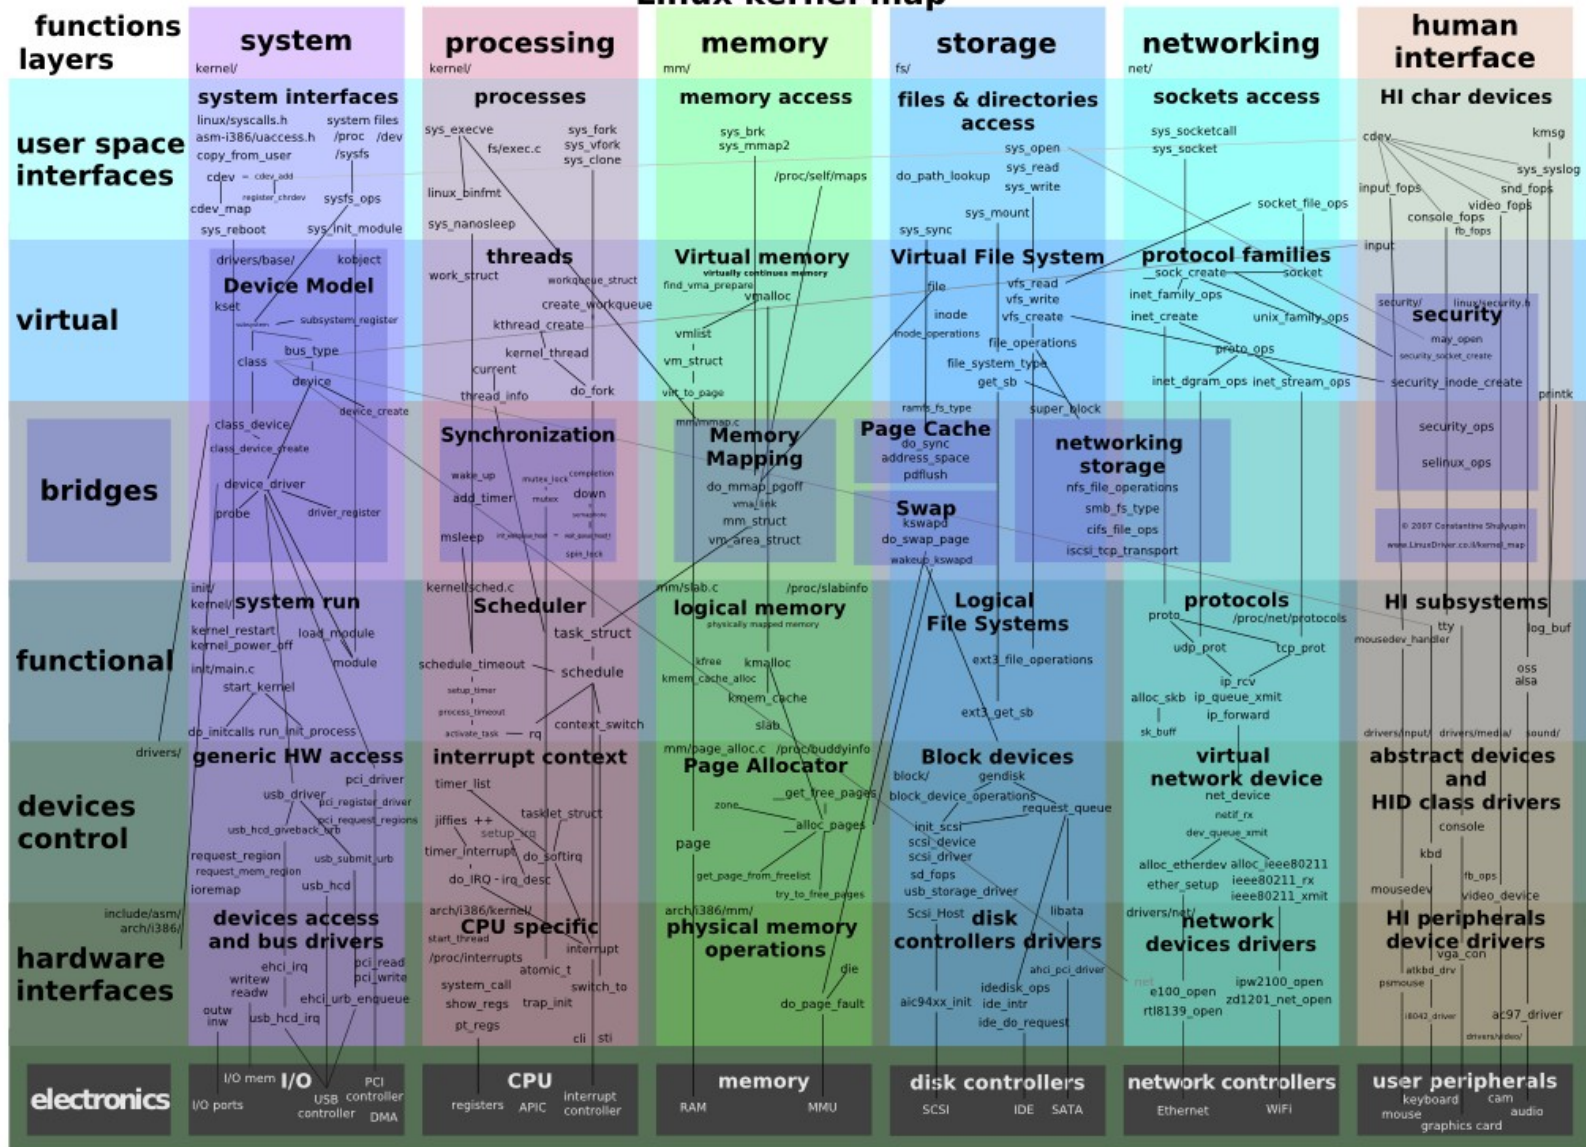
Beyond simple but not fully layered





Linux Structure

Linux kernel map



© https://en.wikipedia.org/wiki/Linux_kernel





version 4.10, 2017-03-10
outlines the Linux storage stack as of Kernel version 4.10






Layered Approach

■ OS Divided into a Number of Layers (Levels)

- Each built on top of lower layers
- Bottom layer (layer 0): HW
- Highest (layer N): user interface



■ Modularity

- Layers are selected such that each uses functions (operations) and services of **only lower-level layers**






Layered Approach (cont.)

■ Pros

- Simplicity of construction
- Simplified verification scheme
- Ease of debugging
- Can easily update or extend each level



■ Cons

- Difficulty to define layers
- Layers may need to interact with various layers
- Too many layers → performance degradation





Microkernel System Structure

- Moves as Much from Kernel into User Space
 - Non-essential components implemented as system and user-level programs
 - → Results in **much smaller kernel** called microkernel
- Microkernel
 - Provide minimal process and memory management and communication facility
- **Mach** Examples of **Microkernel**
 - Mac OS X kernel (**Darwin**) partly based on Mach
 - QNX: a real-time OS
 - True64 UNIX (or formerly Digital UNIX)





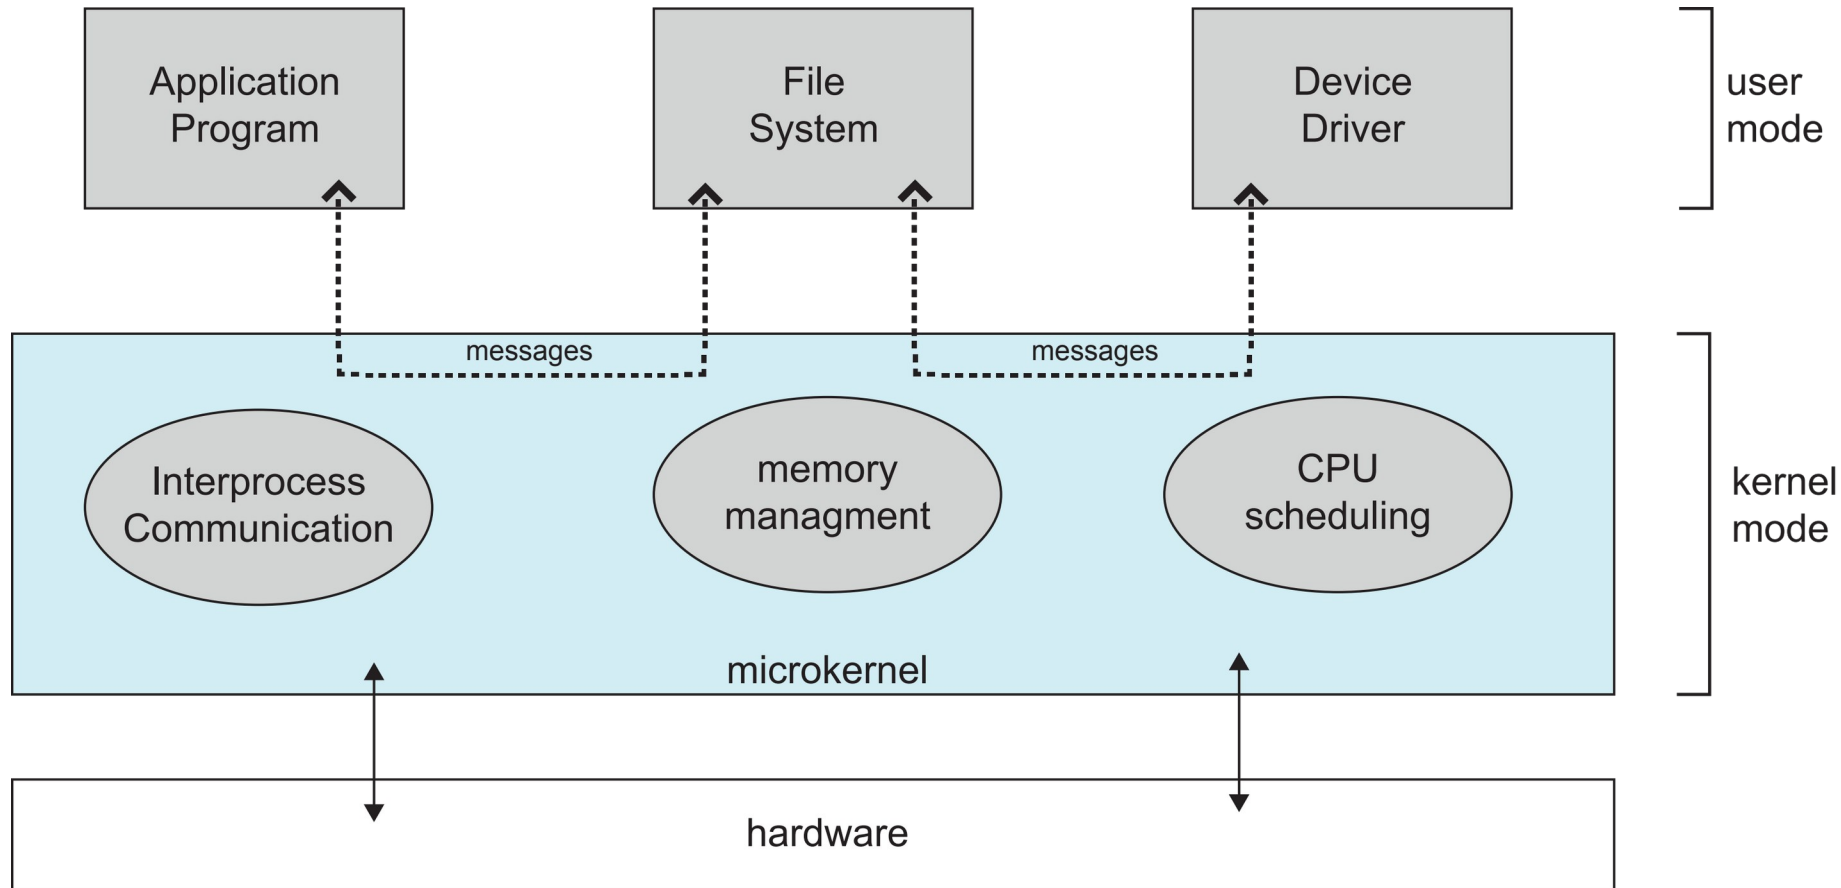
Microkernel System Structure (cont.)

- Communication Takes Place between User Modules using **Message Passing**
- **Benefits**
 - Easier to extend a microkernel
 - Easier to port OS to new architectures
 - ▶ Fewer modification needed
 - More reliable (less code running in kernel mode)
 - More secure
- **Detriments**
 - Performance overhead of user space to kernel space communication (e.g., WinNT vs. Win 95)





Microkernel System Structure (cont.)





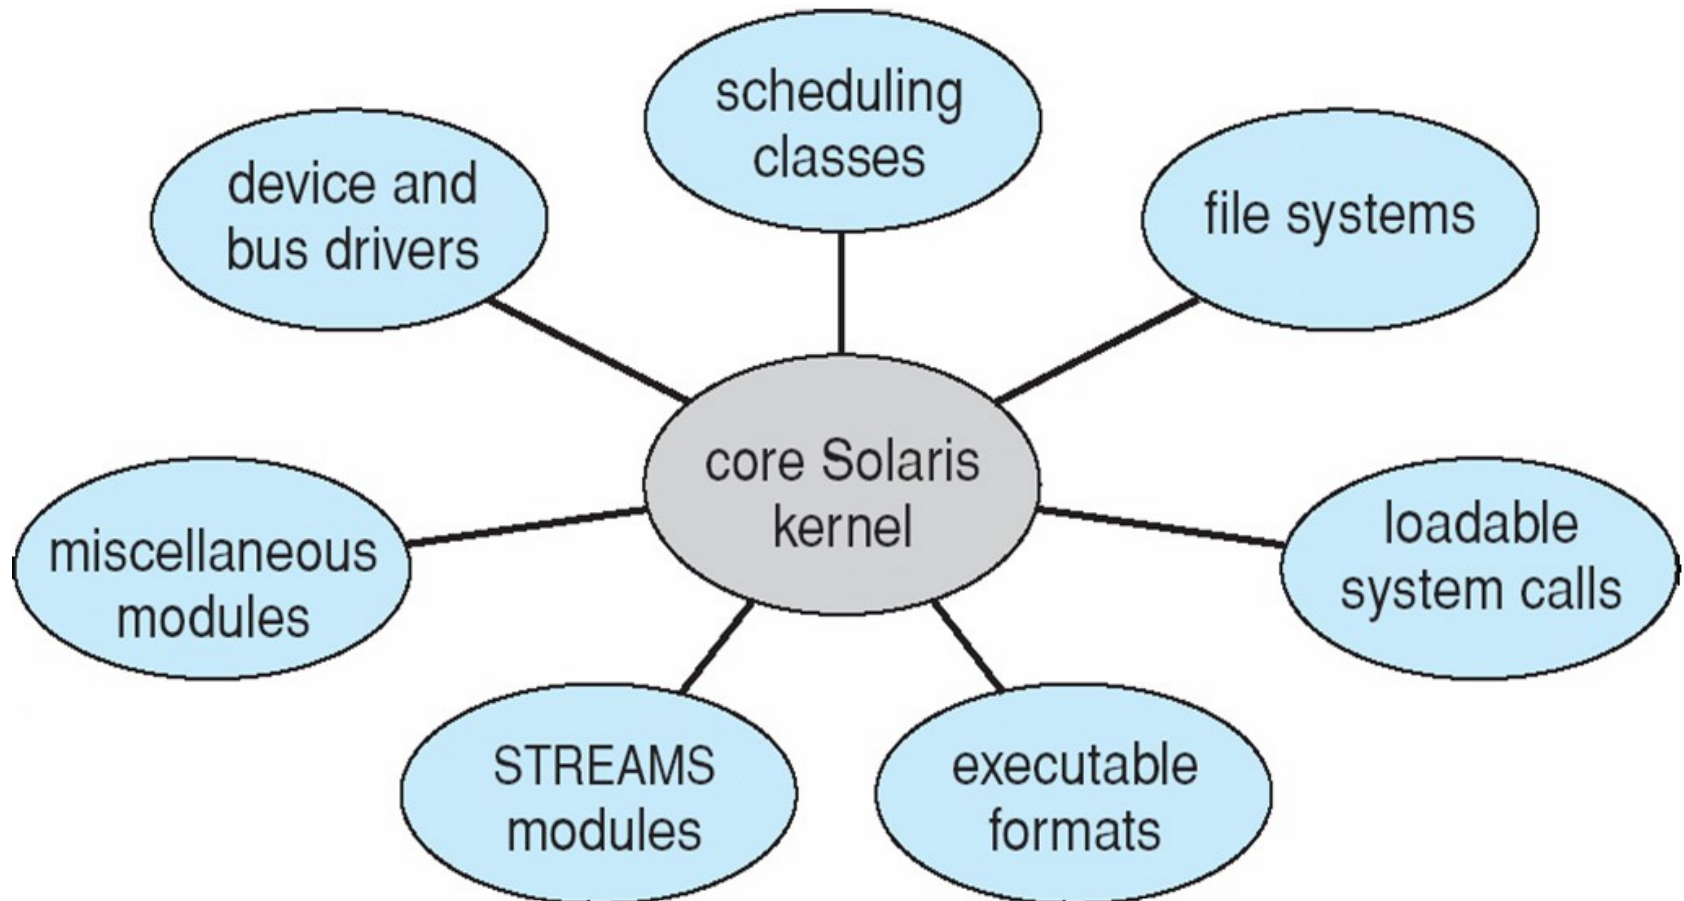
Modules

- **Loadable Kernel Modules** (Used in Adv. OS)
 - Components can be loaded either at **boot time** or during **run time**
 - Uses **object-oriented** approach
 - Each **core** component is **separate**
 - Each talks to others over **known interfaces**
 - Each is **loadable as needed** within kernel
- **Similar to Layers** but with More Flexibility
 - Linux, Solaris, etc
- **Similar to Microkernels** but More Efficient
 - Modules do not need to invoke message passing to communicate





Solaris Modular Approach





Hybrid Systems

- **Most Modern OSes actually not One Pure Model**
 - Hybrid combines multiple approaches to address performance, security, usability needs
 - **Linux** and **Solaris** kernels in kernel address space, so **monolithic**, plus modular for **dynamic loading** of functionality
 - **Windows** mostly **monolithic**, plus microkernel for different subsystem ***personalities***
- **Apple Mac OS X hybrid, Layered**
 - Below is kernel consisting of **Mach** microkernel and BSD Unix parts, plus I/O kit and dynamically loadable modules (called **kernel extensions**)





Mac OS X Structure

graphical user interface

Aqua

application environments and services

Java

Cocoa

Quicktime

BSD

kernel environment

- **Memory management**
- **Remote procedural calls**
- **Interprocess communications**
- **Thread scheduling**

Mach

- **Command line interface**
- **Networking**
- **File systems**
- **POSIX APIs**

BSD

I/O kit

kernel extensions





Virtual Machines


■ Features

- To **abstract** HW of a single computer into several different execution environments
- Creates an **illusion** that each environment running its own **private** computer
- Allows OSes to run applications within other OSes
 - Vast and growing industry
- **Host OS** runs several **guest OSes** as processes
- First appeared commercially on IBM mainframes
 - As VM operating system in 1972

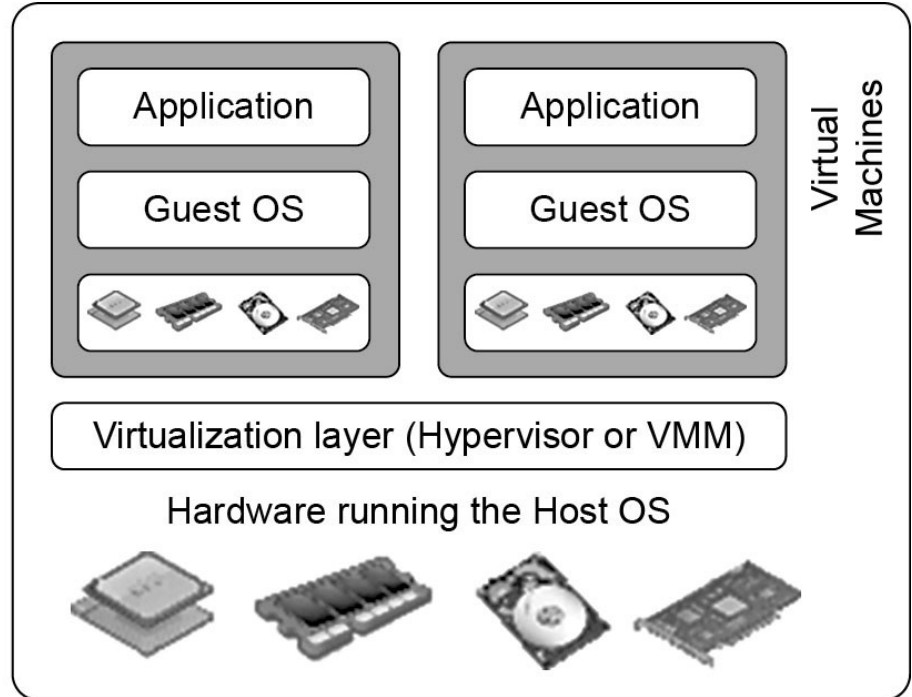




Virtual Machines (cont.)



(a) Traditional computer



(b) After virtualization

(Courtesy of VMWare, 2008)





Virtual Machines (cont.)

■ Hypervisor (HV): Virtual Machine Monitor


- SW/HW that creates and runs virtual machines

■ Type 1

- HV runs directly on host's HW to control HW and to manage guest OSes
- A guest OS thus runs on another level above HV

■ Type 2

- HV runs within a conventional OS environment
- With HV layer as a distinct second SW level, guest OS run at third level above HW



Hypervisor Design:

Two approaches

Type 2 Hypervisor



Examples:

Virtual PC & Virtual Server
VMware Workstation
KVM

Type 1 Hypervisor





Examples:

Hyper-V
Xen
VMware ESX

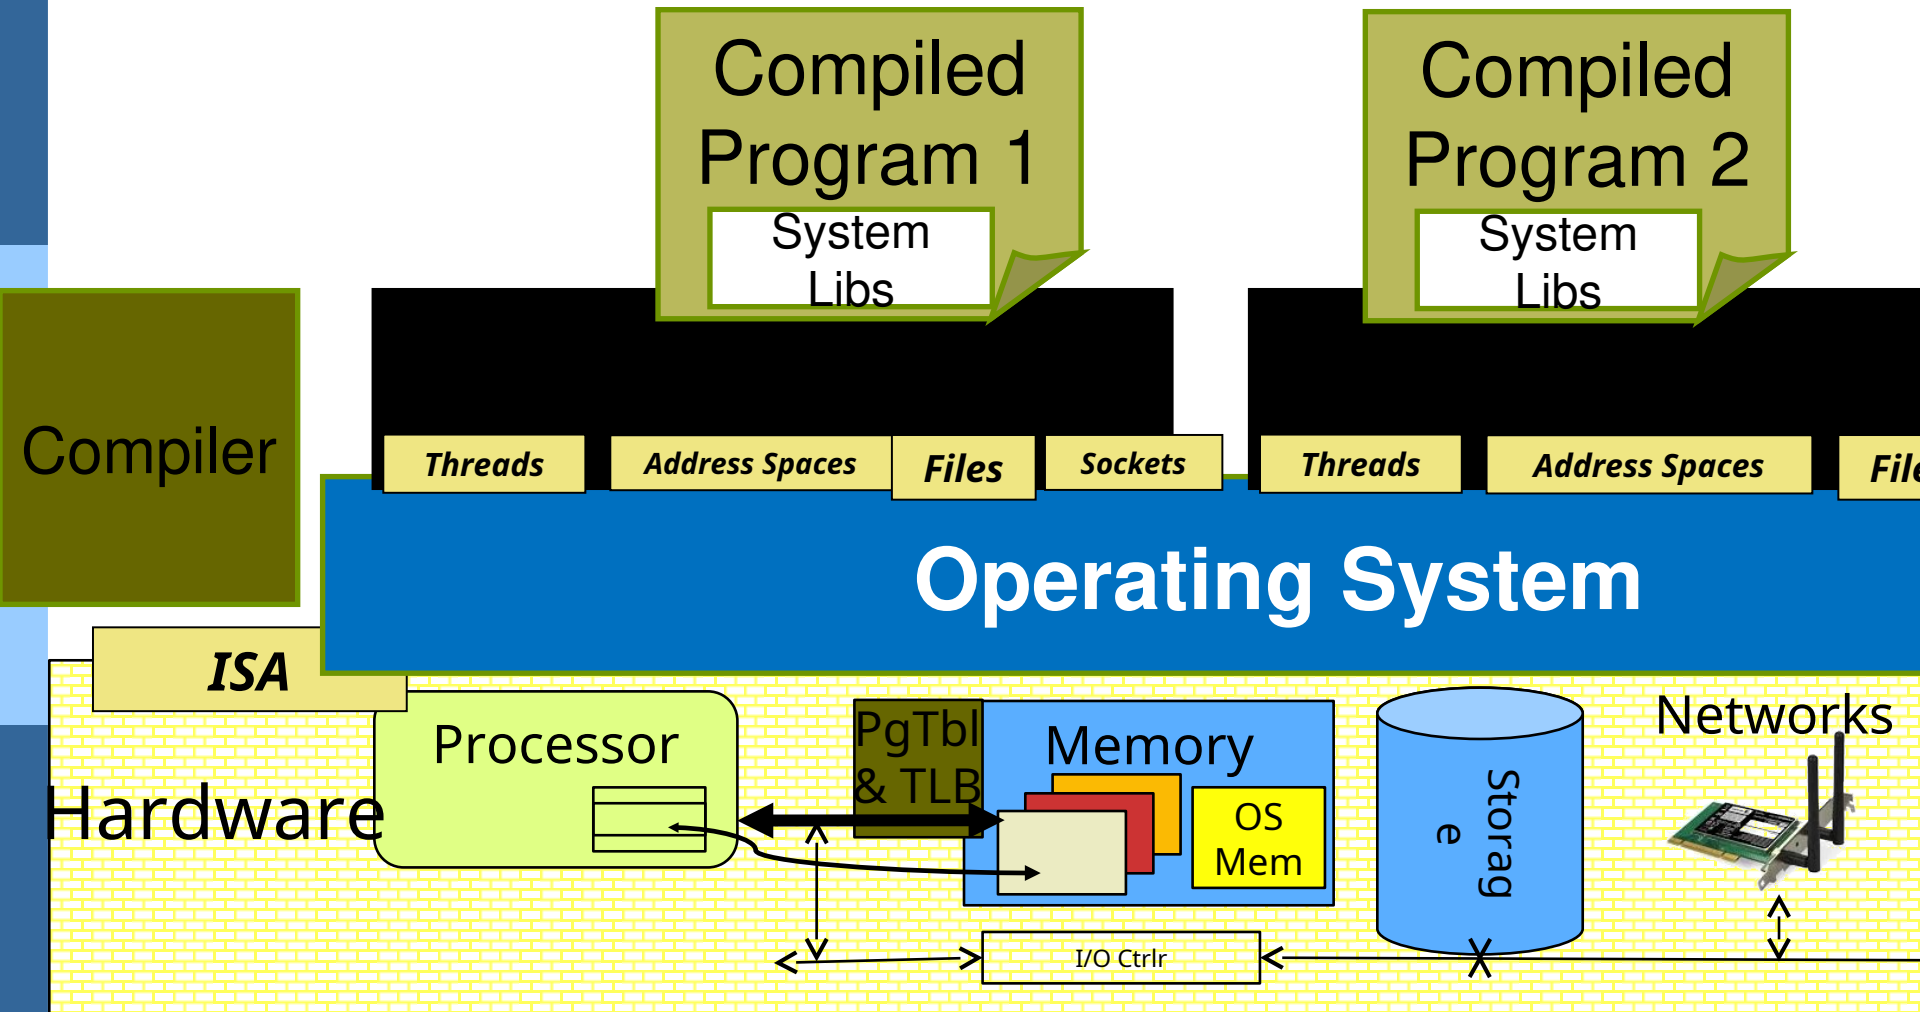


Virtualization (cont.)






Virtualization (cont.)

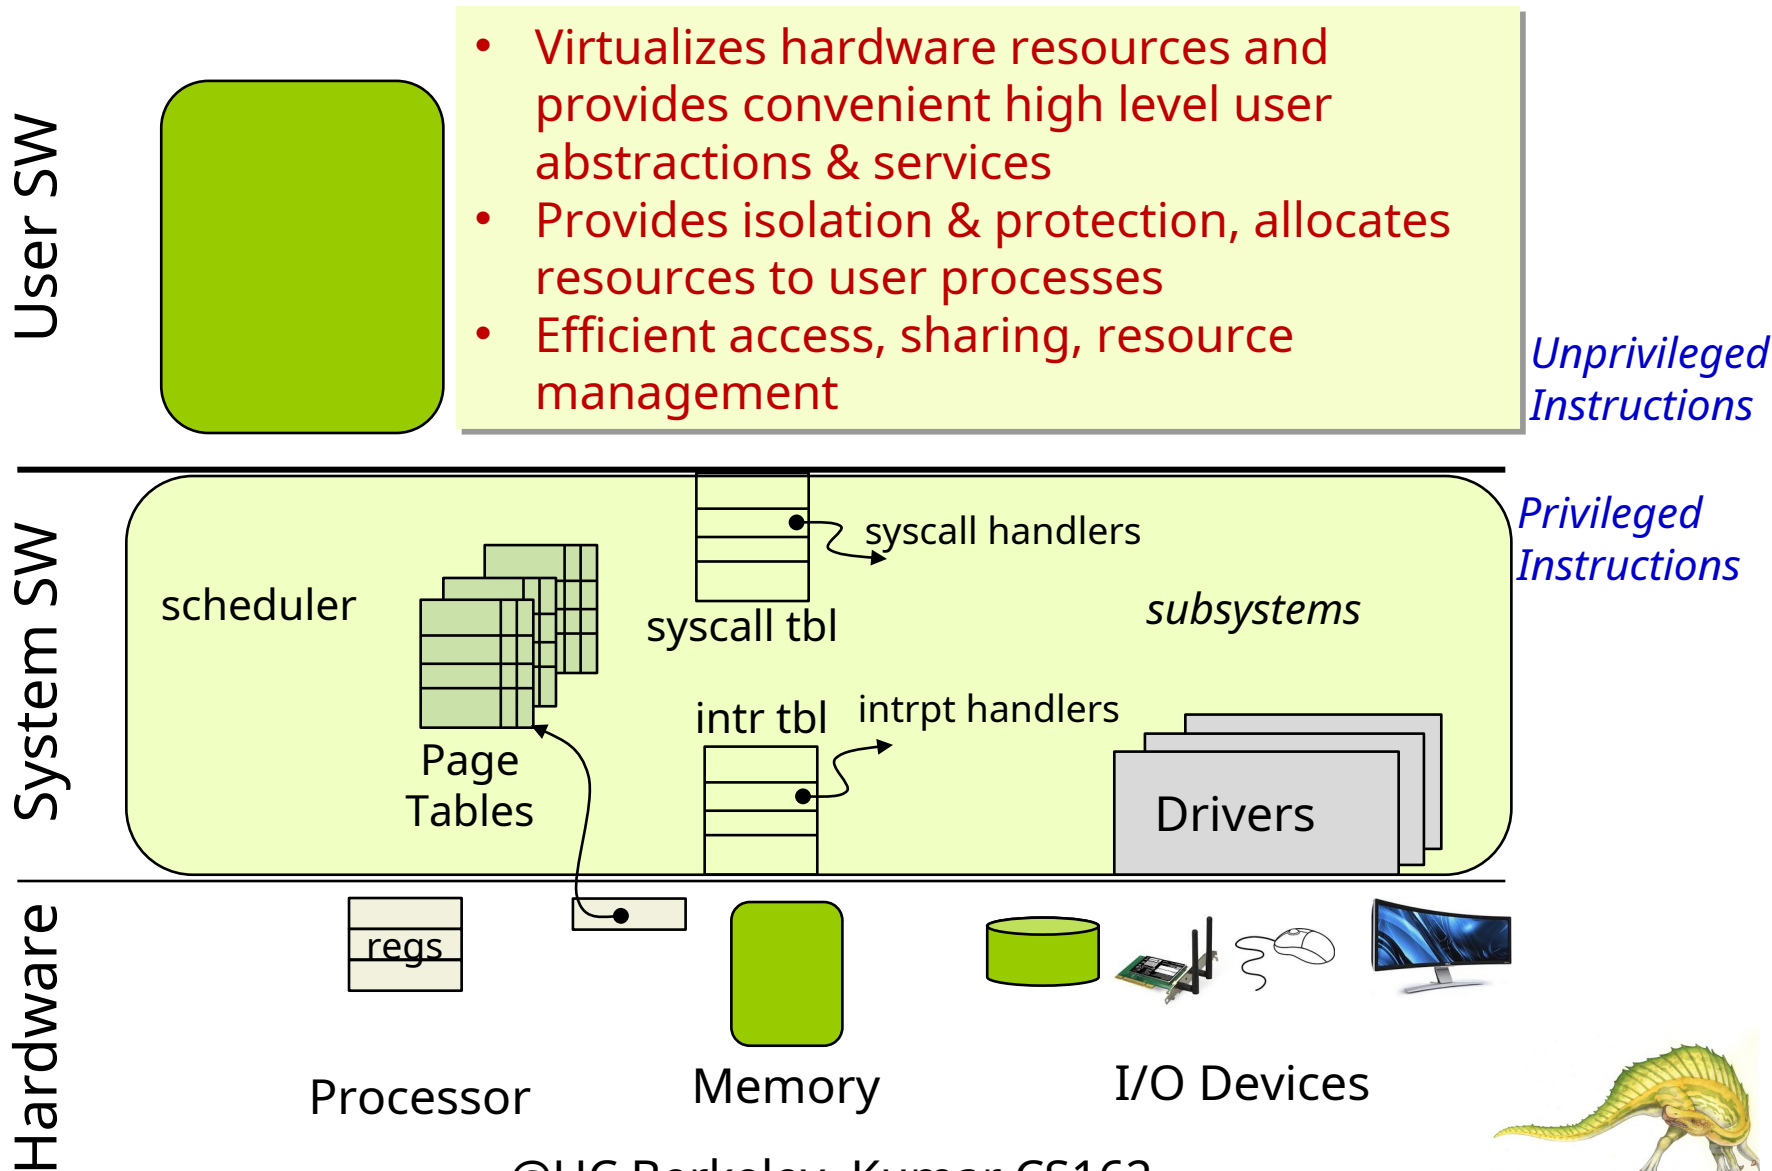


@UC Berkeley, Kumar CS162





OS Software Supports User-Level SW



@UC Berkeley, Kumar CS162





Virtual Machines (cont.)

■ Main Benefits

- Guest OSes do not have to be compliant with HW
 - ▶ Making it possible to run different OSes on same computer
- Host system protected from other VMs
 - ▶ A virus in a guest OS damage OS itself but is unlikely to affect the host or other guest OSes
- Alleviates system development
 - ▶ System development is very dangerous on non-virtualized OS
- A VM can easily be relocated from one physical machine to another as needed





Virtual Machines (cont.)

■ Main Benefits

- Rapid porting and testing of program
 - ▶ Quality assurance engineers can test their programs on multiple platforms on a single HW
- Can consolidate multiple systems in a single HW
 - ➔ Significant power efficiency in data centers
- Can offer pre-installed applications on VMs
 - ▶ Easy to install and config new applications on systems
 - ▶ Technical support of applications becomes straightforward (do not need to tune or config for different platforms)





Virtual Machines (cont.)

■ Virtualization

- Guest OS & applications run at **almost full speed**
- Only system's resources need to be virtualized
 - ▶ VMware and Java Virtual Machine (JVM)
 - ▶ ISA is not virtualized

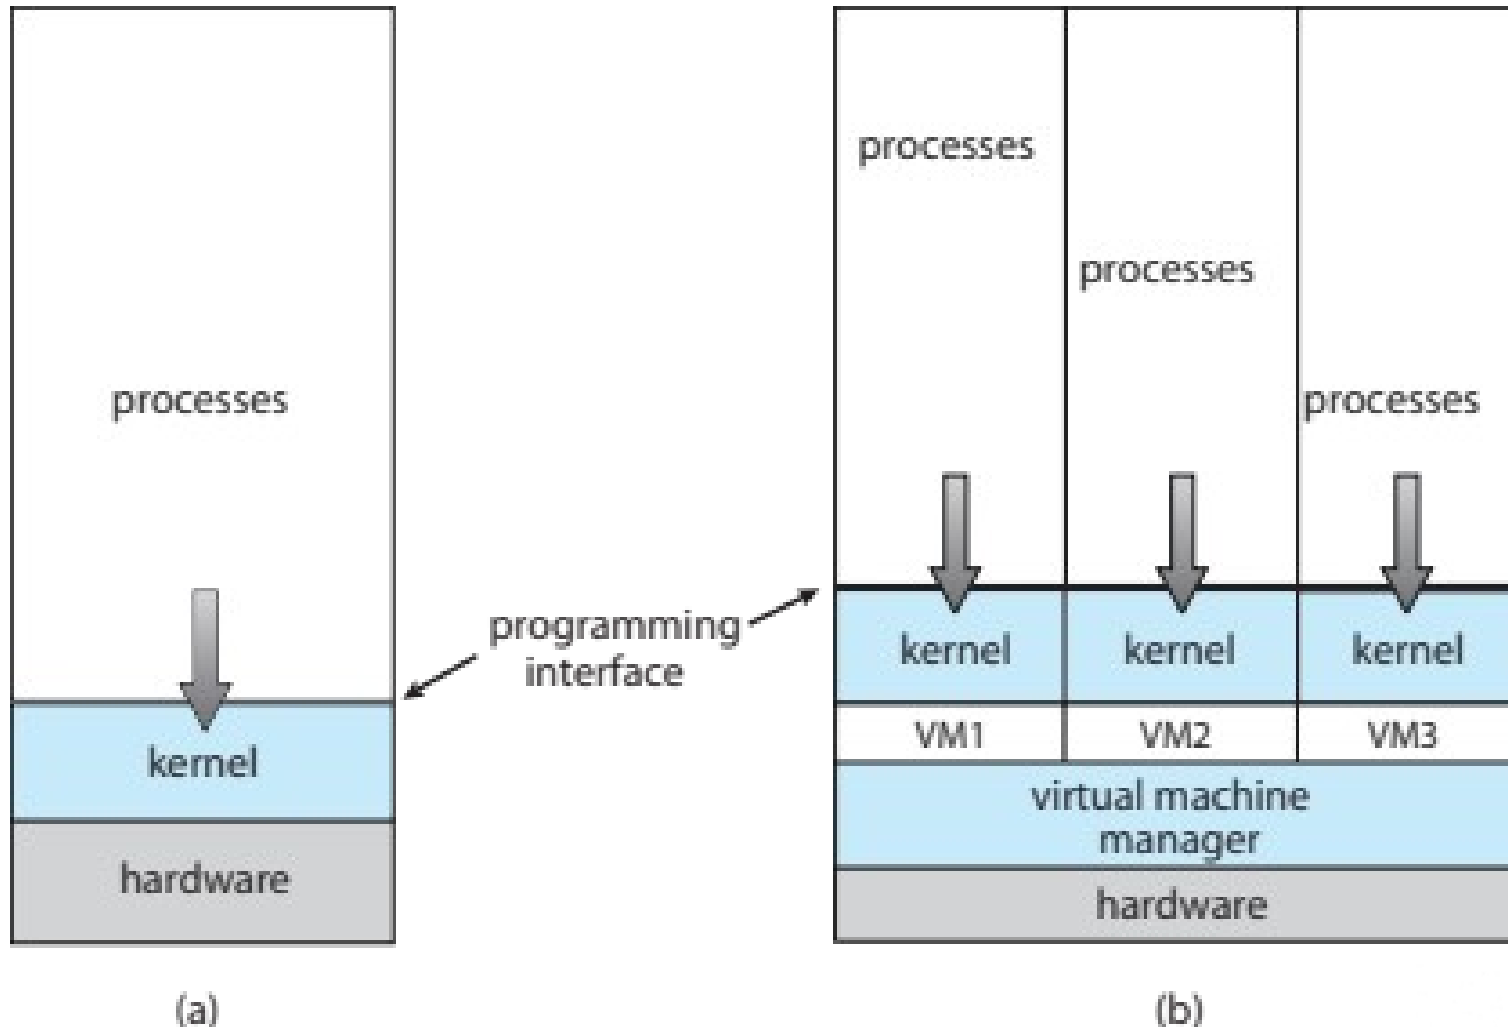
■ Emulation

- ISA is virtually implemented (or emulated)
- Used when source CPU type different from target type (i.e. PowerPC to Intel x86)
- Generally slowest method (by 10X)





Virtual Machines (cont.)






Virtual Machines (cont.)

■ Implementation

- Underlying machine: user and kernel mode
- HV can run in **kernel** mode
- VMs can execute only in **user** mode
- VM: virtual user mode and virtual kernel mode
 - ▶ Both runs in physical user mode
- ➔ This could adversely **affect performance**
- This requires some sort of HW mode support for virtualization: called **host mode** or **HV mode**
 - ▶ Example: AMD supports host mode vs. guest mode

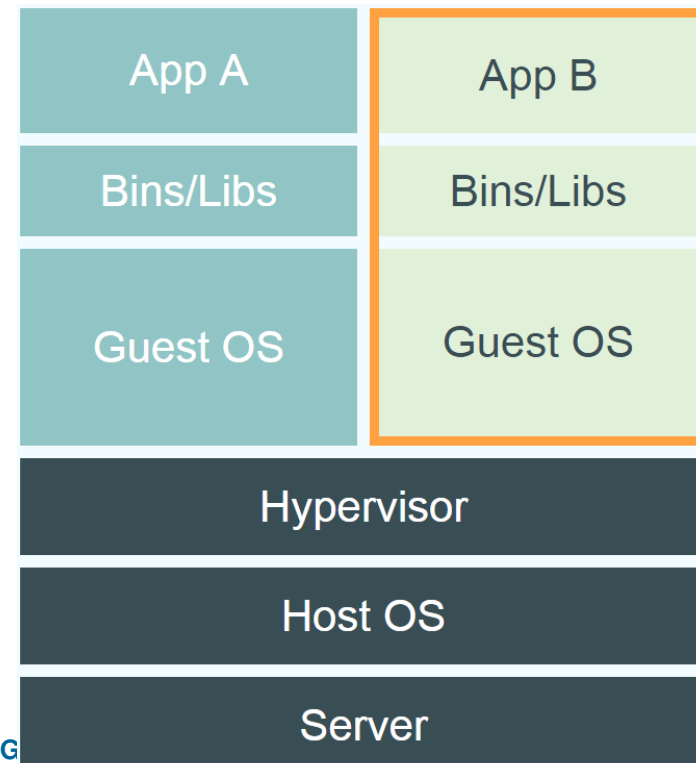
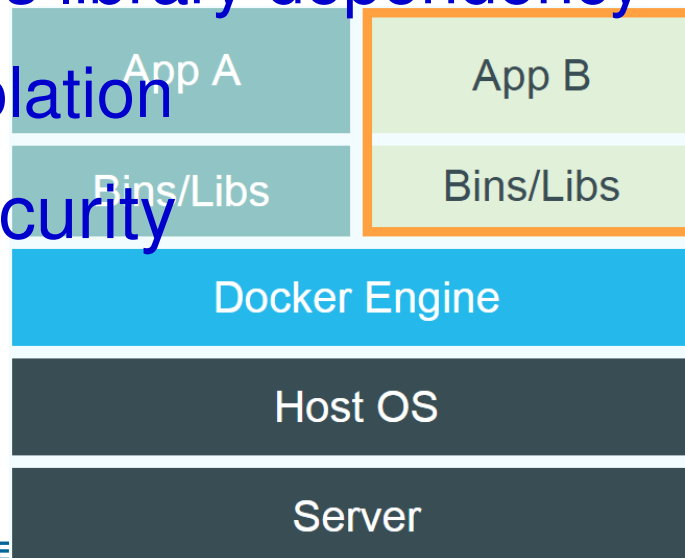




Docker/Container vs. Virtualization

■ Container/Docker

- It does OS virtualization (not HW virtualization)
- OS-Level virtualization to deliver SW in packages
- Share one OS (no multiple guest OS)
- Higher performance
- Resolves library dependency
- Less isolation
- Less security

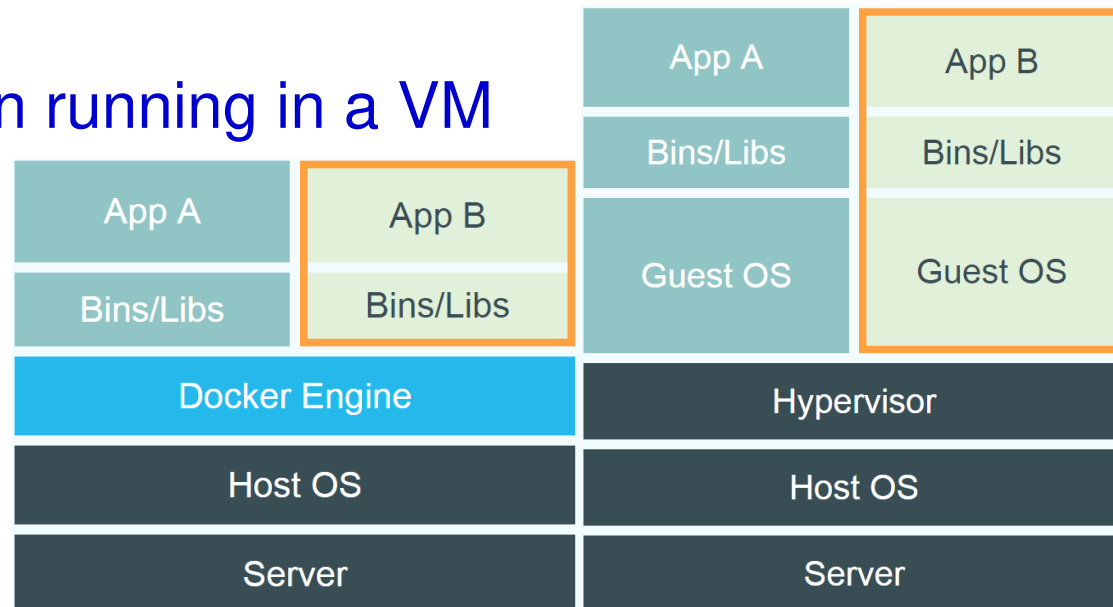




Docker/Container vs. Virtualization

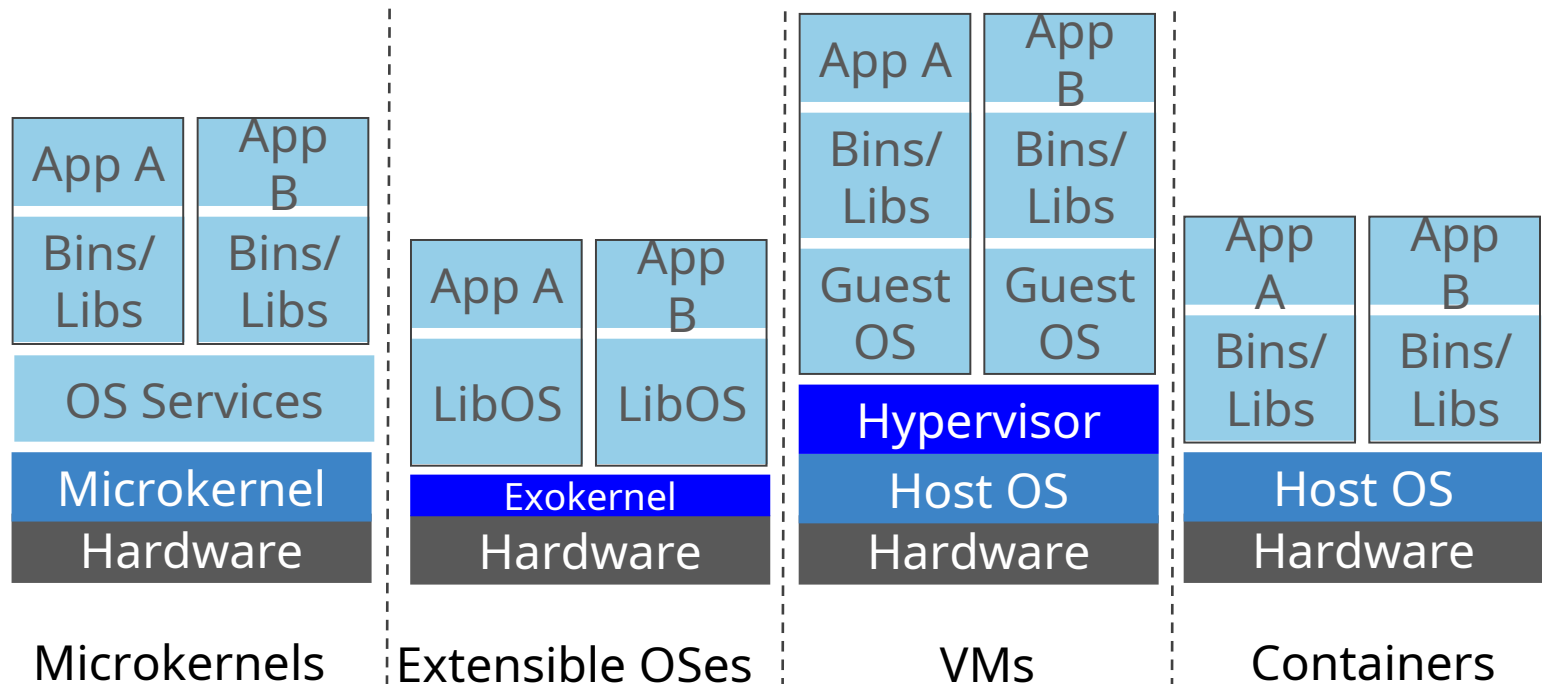
■ Container/Docker

- Each container contains *only* application & its libraries and dependencies
- Small, fast, and portable
- A containerized application has more direct access to HW
 - ▶ Vs. an application running in a VM





VMs vs. Containers vs. Microkernels



© <http://www.eecs.berkeley.edu/~kubitron/cs262>

By John Kubiawicz





Reading Assignment





Reading Assignments

■ OS Debugging

- **Debugging** is finding and fixing errors, or **Bugs**
- OS generate **log files** containing error Info

■ Performance Monitoring and Tuning

- Top command

■ Dtrace

■ Operating System Generation

■ Para-Virtualization

■ Containers





OS Debugging

- **Debugging** is Finding and Fixing errors, or **Bugs**
- OS Generate **Log Files** Containing Error Information
- Failure of an Application can Generate **Core Dump** File Capturing Memory of Process
- OS Failure can Generate **Crash Dump** File Containing Kernel Memory
- Beyond Crashes, Performance Tuning can Optimize System Performance
 - Sometimes using **trace listings** of activities, recorded for analysis
 - **Profiling** is periodic sampling of instruction pointer to look for statistical trends

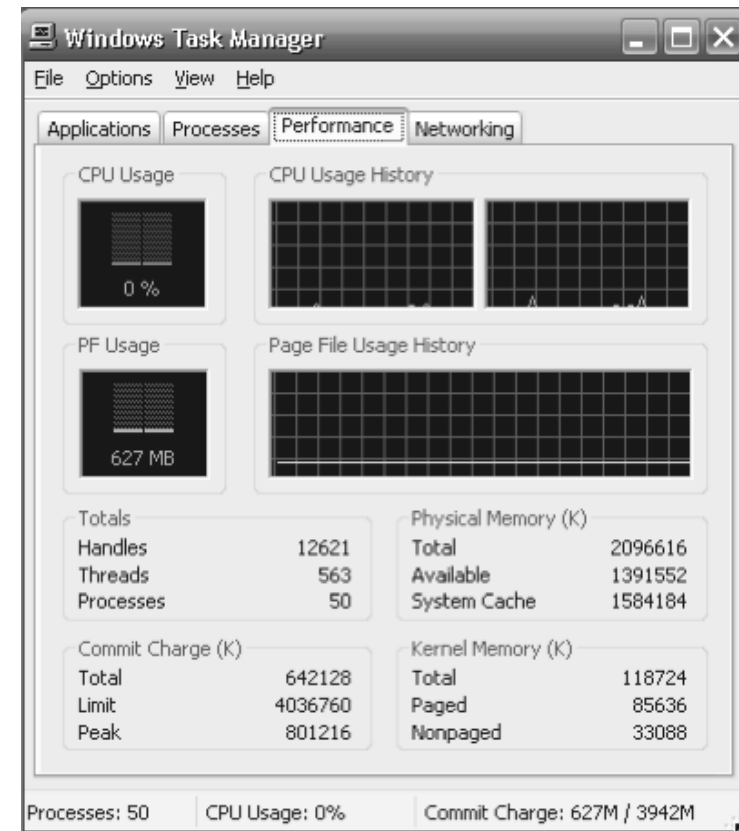
Kernighan's Law: "**Debugging is twice as hard as writing** the code in the first place. Therefore, if you write the code as cleverly as possible, you are, by definition, **not smart enough to debug it.**"





Performance Tuning

- Improve Performance by Removing Bottlenecks
- OS must Provide Means of Computing and Displaying Measures of System Behavior
- For example, “top” program or Windows Task Manager






DTrace

- DTrace tool in Solaris, FreeBSD, Mac OS X allows **live instrumentation** on production systems
- **Probes** fire when code is executed within a **provider**, capturing state data and sending it to **consumers** of those probes
- Example of following XEventsQueued system call move from libc library to kernel and back

```
# ./all.d 'pgrep xclock' XEventsQueued
dtrace: script './all.d' matched 52377 probes
CPU FUNCTION
0 -> XEventsQueued U
0 -> _XEventsQueued U
0 -> _X11TransBytesReadable U
0 <- _X11TransBytesReadable U
0 -> _X11TransSocketBytesReadable U
0 <- _X11TransSocketBytesreadable U
0 -> ioctl U
0 -> ioctl K
0 -> getf K
0 -> set_active_fd K
0 <- set_active_fd K
0 <- getf K
0 -> get_udatamodel K
0 <- get_udatamodel K
...
0 -> releasef K
0 -> clear_active_fd K
0 <- clear_active_fd K
0 -> cv_broadcast K
0 <- cv_broadcast K
0 <- releasef K
0 <- ioctl K
0 <- ioctl U
0 <- _XEventsQueued U
0 <- XEventsQueued U
```





Dtrace (Cont.)

- DTrace code to record amount of time each process with UserID 101 is in running mode (on CPU) in nanoseconds

```
sched:::on-cpu
uid == 101
{
    self->ts = timestamp;
}

sched:::off-cpu
self->ts
{
    @time[execname] = sum(timestamp - self->ts);
    self->ts = 0;
}
```

```
# dtrace -s sched.d
dtrace: script 'sched.d' matched 6 probes
^C
gnome-settings-d      142354
gnome-vfs-daemon      158243
dsdm                  189804
wnck-applet           200030
gnome-panel           277864
clock-applet          374916
mapping-daemon        385475
xscreensaver          514177
metacity              539281
Xorg                  2579646
gnome-terminal         5007269
mixer_applet2         7388447
java                  10769137
```

Figure 2.21 Output of the D code.





Operating System Generation

- OSes Designed to Run on any of a Class of Machines
 - System must be configured for each specific computer site
- **SYSGEN** Program Obtains Info Concerning Specific Configuration of HW system
 - Used to build system-specific compiled kernel or system-tuned
 - Can generate more efficient code than one general kernel



End of Lecture 2

